

Project 3 - Walmart

May 20, 2022

0.1 About Walmart

Walmart is an American multinational retail corporation that operates a chain of supercenters, discount departmental stores, and grocery stores from the United States. Walmart has more than 100 million customers worldwide.

0.1.1 Business Problem

The Management team at Walmart Inc. wants to analyze the customer purchase behavior (specifically, purchase amount) against the customer's gender and the various other factors to help the business make better decisions. They want to understand if the spending habits differ between male and female customers: Do women spend more on Black Friday than men? (Assume 50 million customers are male and 50 million are female).

0.1.2 Here the assumption is that 50 million customers are male and 50 million customers are female.

The company collected the transactional data of customers who purchased products from the Walmart Stores during Black Friday. The dataset has the following features:

- User_ID: User ID
- Product_ID: Product ID
- Gender: Sex of User
- Age: Age in bins
- Occupation: Occupation(Masked)
- City_Category: Category of the City (A,B,C)
- StayInCurrentCityYears: Number of years stay in current city
- Marital_Status: Marital Status
- ProductCategory: Product Category (Masked)
- Purchase: Purchase Amount

Black Friday is a term for the Friday after Thanksgiving in the United States. It traditionally marks the start of the Christmas shopping season in the United States.

[]:

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
from scipy import stats
```

```
[2]: #Importing the dataset
df=pd.read_csv('Walmart.csv')
df.head()
```

```
[2]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
0	1000001	P00069042	F	0-17	10	A	
1	1000001	P00248942	F	0-17	10	A	
2	1000001	P00087842	F	0-17	10	A	
3	1000001	P00085442	F	0-17	10	A	
4	1000002	P00285442	M	55+	16	C	

	Stay_In_Current_City_Years	Marital_Status	Product_Category	Purchase
0	2	0	3	8370
1	2	0	1	15200
2	2	0	12	1422
3	2	0	12	1057
4	4+	0	8	7969

```
[3]: #No. of rows and columns of the dataset
df.shape
```

```
[3]: (550068, 10)
```

```
[4]: #No. of null values
df.isna().sum()/len(df)
#There are no null values in any of the columns.
```

```
[4]:
```

User_ID	0.0
Product_ID	0.0
Gender	0.0
Age	0.0
Occupation	0.0
City_Category	0.0
Stay_In_Current_City_Years	0.0
Marital_Status	0.0
Product_Category	0.0
Purchase	0.0

dtype: float64

```
[5]: #Information about the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	User_ID	550068 non-null	int64
1	Product_ID	550068 non-null	object
2	Gender	550068 non-null	object
3	Age	550068 non-null	object
4	Occupation	550068 non-null	int64
5	City_Category	550068 non-null	object
6	Stay_In_Current_City_Years	550068 non-null	object
7	Marital_Status	550068 non-null	int64
8	Product_Category	550068 non-null	int64
9	Purchase	550068 non-null	int64

dtypes: int64(5), object(5)

memory usage: 42.0+ MB

```
[6]: print(df['Product_Category'].nunique())
      # There are 20 unique product categories.
      print()
      print(df['Product_Category'].unique().tolist())
      print()
      print(df['Product_Category'].value_counts())
```

20

[3, 1, 12, 8, 5, 4, 2, 6, 14, 11, 13, 15, 7, 16, 18, 10, 17, 9, 20, 19]

5	150933
1	140378
8	113925
11	24287
2	23864
6	20466
3	20213
4	11753
16	9828
15	6290
13	5549
10	5125
12	3947
7	3721
18	3125
20	2550
19	1603
14	1523
17	578
9	410

Name: Product_Category, dtype: int64

```
[7]: print(df['Product_ID'].nunique())
print()
# There are 3631 unique product IDs.
print(df['Product_ID'].value_counts()[:10]) #Top 10 product IDs
```

3631

P00265242	1880
P00025442	1615
P00110742	1612
P00112142	1562
P00057642	1470
P00184942	1440
P00046742	1438
P00058042	1422
P00059442	1406
P00145042	1406

Name: Product_ID, dtype: int64

```
[8]: print(df['Age'].nunique())
print()
# There are 7 unique age categories.
print(df['Age'].unique().tolist())
print()
print(df['Age'].value_counts())
```

7

['0-17', '55+', '26-35', '46-50', '51-55', '36-45', '18-25']

26-35	219587
36-45	110013
18-25	99660
46-50	45701
51-55	38501
55+	21504
0-17	15102

Name: Age, dtype: int64

```
[9]: print(df['Occupation'].nunique())
print()
# There are 20 unique Occupation years.
print(df['Occupation'].unique().tolist())
print()
print(df['Occupation'].value_counts())
```

21

[10, 16, 15, 7, 20, 9, 1, 12, 17, 0, 3, 4, 11, 8, 19, 2, 18, 5, 14, 13, 6]

4	72308
0	69638
7	59133
1	47426
17	40043
20	33562
12	31179
14	27309
2	26588
16	25371
6	20355
3	17650
10	12930
5	12177
15	12165
11	11586
19	8461
13	7728
18	6622
9	6291
8	1546

Name: Occupation, dtype: int64

```
[10]: print(df['City_Category'].nunique())
print()
# There are 3 unique City categories.
print(df['City_Category'].unique().tolist())
print()
print(df['City_Category'].value_counts())
```

3

['A', 'C', 'B']

B	231173
C	171175
A	147720

Name: City_Category, dtype: int64

```
[11]: print(df['Stay_In_Current_City_Years'].nunique())
print()
# There are 5 unique Current City Stay Years.
print(df['Stay_In_Current_City_Years'].unique().tolist())
print()
print(df['Stay_In_Current_City_Years'].value_counts())
```

5

['2', '4+', '3', '1', '0']

1 193821
2 101838
3 95285
4+ 84726
0 74398

Name: Stay_In_Current_City_Years, dtype: int64

```
[12]: print(df['Gender'].nunique())  
print()  
# There are 2 unique Gender categories.  
print(df['Gender'].unique().tolist())  
print()  
print(df['Gender'].value_counts())
```

2

['F', 'M']

M 414259
F 135809

Name: Gender, dtype: int64

```
[13]: print(df['Marital_Status'].nunique())  
print()  
# There are 2 unique Marital Status categories.  
print(df['Marital_Status'].unique().tolist())  
print()  
print(df['Marital_Status'].value_counts())
```

2

[0, 1]

0 324731
1 225337


Name: Marital_Status, dtype: int64

```
[14]: print(df['User_ID'].nunique())  
# There are 5891 unique users.
```

5891


[]:

```
[15]: #Using Boxplot on Purchase
plt.figure(figsize=(15,5))
sns.boxplot(data=df,x='Purchase',color='orange')
plt.xticks(np.arange(0,26000,500),rotation=90)
plt.show()
#We see that there are few outliers after which are above 21500.
#The median is around 8000.
```



```
[16]: #Removing the rows for which outliers are present in the Purchase column
q75,q25 = np.percentile(df['Purchase'],[75,25])
intr_qr = q75-q25
maximum = q75+(1.5*intr_qr)
minimum = q25-(1.5*intr_qr)
df.loc[df['Purchase'] < minimum,'Purchase'] = np.nan
df.loc[df['Purchase'] > maximum,'Purchase'] = np.nan
df.dropna(inplace=True)
```


```
[17]: #Using Boxplot on Purchase again. We see that the outliers have been removed.
plt.figure(figsize=(15,5))
sns.boxplot(data=df,x='Purchase',color='orange')
plt.xticks(np.arange(0,26000,500),rotation=90)
plt.show()
#We see that there are few outliers after which are above 21500.
#The median is around 8000.
```



```
[18]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 547391 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User_ID               547391 non-null  int64
1   Product_ID            547391 non-null  object
2   Gender                547391 non-null  object
3   Age                  547391 non-null  object
4   Occupation            547391 non-null  int64
5   City_Category         547391 non-null  object
6   Stay_In_Current_City_Years  547391 non-null  object
7   Marital_Status        547391 non-null  int64
8   Product_Category      547391 non-null  int64
9   Purchase              547391 non-null  float64
dtypes: float64(1), int64(4), object(5)
memory usage: 45.9+ MB
```

```
[19]: #Distribution Of Purchase
plt.figure(figsize=(10,6))
sns.histplot(data=df,x='Purchase',bins=20,kde=True)
plt.xticks(np.arange(0,28000,2000))
plt.show()
#We see that the distribution is right skewed and most of the purchase prices
→are in the range 4000-13000.
#There are also a few purchases between 15000-16000 and between 19000-21000.
#Since the distribution is right skewed, therefore the mean would be greater
→than the median.
```


```
[20]: df.describe()
# Here also we see that the mean Purchase price is greater than the median,
# Purchase price, since it is right skewed.
#The average purchase price is 9195.63 and the median purchase price is 8038.
```

```
[20]:
```

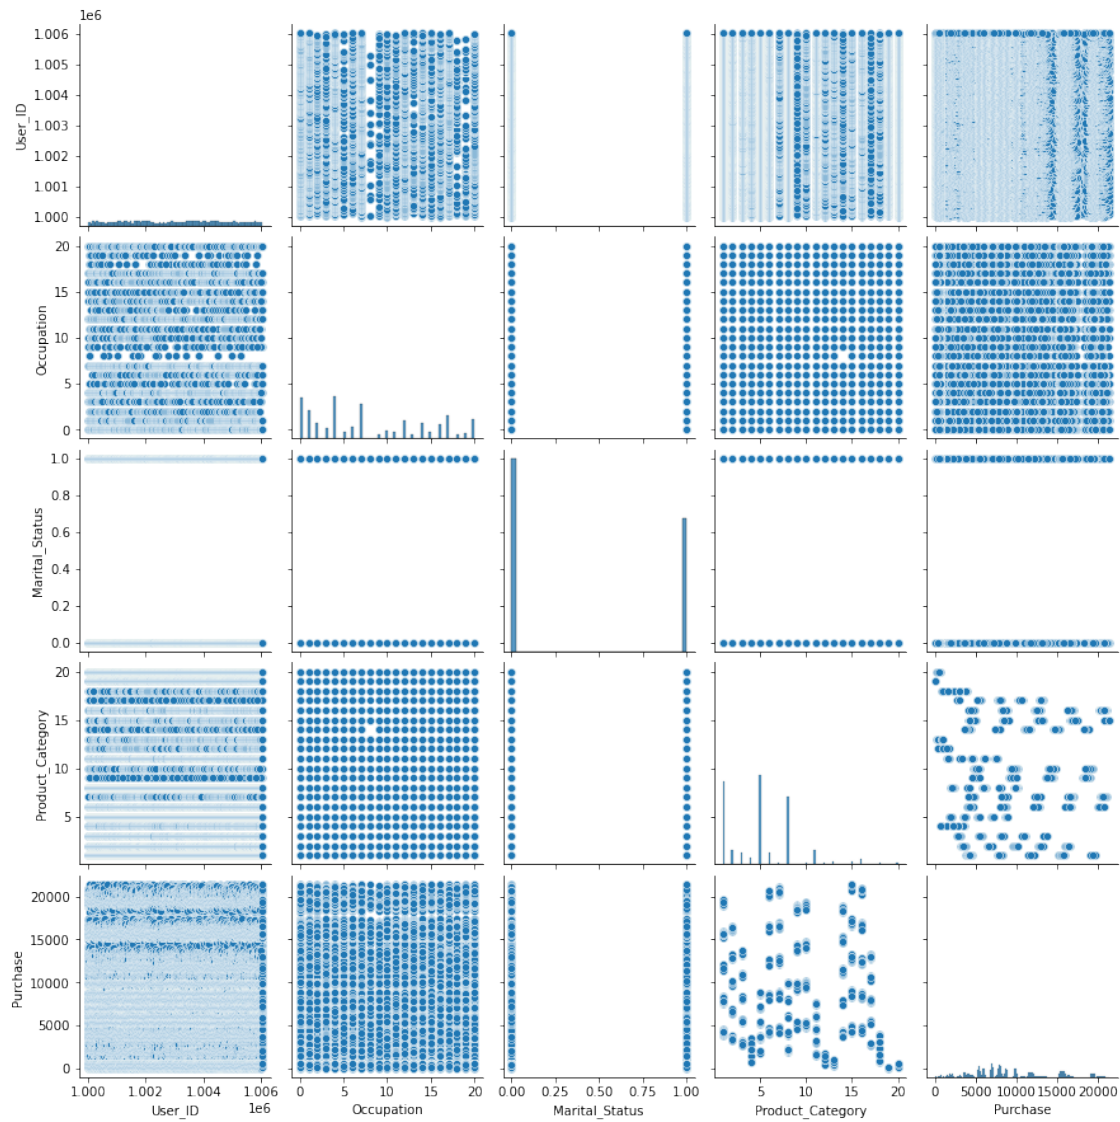
	User_ID	Occupation	Marital_Status	Product_Category \
count	5.473910e+05	547391.000000	547391.000000	547391.000000
mean	1.003028e+06	8.074627	0.409486	5.378945
std	1.727357e+03	6.521586	0.491739	3.927383
min	1.000001e+06	0.000000	0.000000	1.000000
25%	1.001516e+06	2.000000	0.000000	1.000000
50%	1.003075e+06	7.000000	0.000000	5.000000
75%	1.004478e+06	14.000000	1.000000	8.000000
max	1.006040e+06	20.000000	1.000000	20.000000

	Purchase
count	547391.000000
mean	9195.627195
std	4938.872953
min	12.000000
25%	5721.000000
50%	8038.000000
75%	12019.000000
max	21399.000000

```
[21]: plt.figure(figsize=(7,6))
sns.heatmap(df.corr(),cmap='Blues',annot=True)
plt.show()
#From the heatmap, we observe that there isn't a strong relationship between
→ any of the 2 columns.
#The strongest correlation we see is between Product_Category and Purchase.
```




```
[22]: sns.pairplot(df)
plt.show()
#From the pairplot, we get to see that we do not see a relationship between
→ product category and purchase.
#We therefore cannot rely on the correlation coefficient value of -0.37.
```



[]:

0.1.3 Question 1 - Purchase price of males and females?


```
[23]: plt.figure()
sns.countplot(data=df,x='Gender')
plt.show()
```




```
[24]: print(df.groupby(by='Gender').count()['Purchase'])
```

```
Gender
F      135220
M      412171
Name: Purchase, dtype: int64
```

```
[25]: plt.figure(figsize=(7,5))
sns.boxplot(data=df,x='Gender',y='Purchase')
plt.show()
```



```
[26]: plt.figure(figsize=(7,5))  
sns.kdeplot(data=df,hue='Gender',x='Purchase')  
plt.show()
```



```
[27]: print(df.groupby(by='Gender').mean()['Purchase'])
```


```
Gender
F    8671.049039
M    9367.724355
Name: Purchase, dtype: float64
```

Inference - We see that the average amount spent by males is more than females. In the dataset provided, the no of males making a purchase are 3 times of females making a purchase.

```
[ ]:
```

0.1.4 Question 2 - Purchase price for each of the age categories?


```
[28]: plt.figure(figsize=(7,5))
sns.countplot(data=df,x='Age')
plt.show()
```



```
[29]: print(df.groupby(by='Age').count()['Purchase'])
```

```
Age
0-17      15032
18-25     99334
26-35    218661
36-45    109409
46-50     45442
51-55     38191
55+       21322
Name: Purchase, dtype: int64
```

```
[30]: plt.figure(figsize=(7,5))
sns.boxplot(data=df,x='Age',y='Purchase')
plt.show()
```



```
[31]: print(df.groupby(by='Age').mean()['Purchase'].sort_values(ascending=False))
```


```
Age
51-55    9423.121704
36-45    9254.202214
55+      9216.650220
26-35    9193.469924
46-50    9128.985080
18-25    9124.031731
0-17     8867.447046
Name: Purchase, dtype: float64
```

Inference : We see that the average purchase price of all the age categories is almost same, with the age category 51-55 having the highest purchase price and the age group 0-17 having the lowest average purchase price. Most of the customers are in the age category 18-25, 36-45 and 26-35.

```
[ ]:
```


0.1.5 Question 3 - Purchase price for each of the marital status categories?


```
[32]: plt.figure()  
sns.countplot(data=df,x='Marital_Status')  
plt.show()
```



```
[33]: print(df.groupby(by='Marital_Status').count()['Purchase'])
```

```
Marital_Status  
0    323242  
1    224149  
Name: Purchase, dtype: int64
```

```
[34]: plt.figure(figsize=(7,5))  
sns.boxplot(data=df,x='Marital_Status',y='Purchase')  
plt.show()
```



```
[35]: print(df.groupby(by='Marital_Status').mean()['Purchase'] .
      ↪ sort_values(ascending=False))
```


```
Marital_Status
0    9201.581849
1    9187.040076
Name: Purchase, dtype: float64
```

Inference - We see that the average purchase price of both the marital status categories is almost same. The number of unmarried people buying the products is almost 1.5x of that of who are married.

```
[ ]:
```

0.1.6 Question 4 - Average purchase price for each of the city categories ?


```
[36]: plt.figure()
      sns.countplot(data=df,x='City_Category')
      plt.show()
```



```
[37]: print(df.groupby(by='City_Category').count()['Purchase'])
```

```
City_Category
A      147036
B      230114
C      170241
Name: Purchase, dtype: int64
```

```
[38]: plt.figure(figsize=(7,5))
sns.boxplot(data=df,x='City_Category',y='Purchase')
plt.show()
```



```
[39]: print(df.groupby(by='City_Category').mean()['Purchase'].
        ↪sort_values(ascending=False))
```


```
City_Category
C    9645.647300
B    9086.502707
A    8845.367393
Name: Purchase, dtype: float64
```

Inference - We observe that the average purchase price is highest in City C and least in city A. The number of people buying from City B is the highest and lowest for city A among all cities.

```
[ ]:
```

0.1.7 Question 5 - Purchase price for each of the stay_in_the_current_year categories ?


```
[40]: plt.figure()
        sns.countplot(data=df,x='Stay_In_Current_City_Years')
        plt.show()
```



```
[41]: print(df.groupby(by='Stay_In_Current_City_Years').count()['Purchase'].  
        ↪sort_values(ascending=False))
```

```
Stay_In_Current_City_Years  
1      192845  
2      101384  
3       94804  
4+       84322  
0        74036  
Name: Purchase, dtype: int64
```

```
[42]: plt.figure(figsize=(7,5))  
       sns.boxplot(data=df,x='Stay_In_Current_City_Years',y='Purchase')  
       plt.show()
```



```
[43]: print(df.groupby(by='Stay_In_Current_City_Years').mean()['Purchase'] .
      ↪sort_values(ascending=False))
```


```
Stay_In_Current_City_Years
2      9258.292028
3      9215.953451
4+     9208.837895
1      9179.275916
0      9111.331555
Name: Purchase, dtype: float64
```

Inference - We observe that the average purchase price is higher for people whose current city staying years are greater than or equal to 2 years. The number of people who make the most purchases are those that have lived in the current city for 1 year.

```
[ ]:
```

0.1.8 Question 6 - Purchase price for each of the Product categories ?


```
[44]: plt.figure()
      sns.countplot(data=df,x='Product_Category')
      plt.show()
```



```
[45]: print(df.groupby(by='Product_Category').count()['Purchase'].
      ↪sort_values(ascending=False))
```

```
Product_Category
5      150933
1      140378
8      113925
11     24287
2      23864
6      20466
3      20213
4      11753
16      9828
15      5963
13      5549
12      3947
7       3721
18      3125
10      2850
20      2550
19      1603
14      1523
17       578
9        335
Name: Purchase, dtype: int64
```

```
[46]: plt.figure(figsize=(9,7))
sns.boxplot(data=df,x='Product_Category',y='Purchase')
plt.show()
```



```
[47]: print(df.groupby(by='Product_Category').mean()['Purchase'].
↪sort_values(ascending=False))
```

```
Product_Category
10    16626.385965
7     16365.689600
6     15838.478550
16    14766.037037
15    14412.504109
9     13852.325373
1     13606.218596
14    13141.625739
2     11251.935384
17    10170.759516
3     10096.705734
8      7498.958078
5      6240.088178
```



```

11      4685.268456
18      2972.864320
4       2329.659491
12      1350.859894
13       722.400613
20      370.481176
19       37.041797
Name: Purchase, dtype: float64


```

Inference - We see that the highest average purchase price is for product category 10, and the least average purchase price is for product category 19. The most sold product categories are 5,1 and 8.

```
[ ]:
```

0.1.9 Question 7 - Purchase price for each of the occupation categories ?

```
[48]: plt.figure()
sns.countplot(data=df,x='Occupation')
plt.show()
```



```
[49]: print(df.groupby(by='Occupation').count()['Purchase'].
        ↳sort_values(ascending=False))
```

```


Occupation
4      72040

```

```
0    69310
7    58875
1    47174
17   39855
20   33355
12   30995
14   27173
2    26435
16   25251
6    20261
3    17568
10   12888
5    12133
15   12086
11   11500
19    8412
13    7667
18    6595
9     6278
8     1540
```

Name: Purchase, dtype: int64

```
[50]: plt.figure(figsize=(9,7))
      sns.boxplot(data=df,x='Occupation',y='Purchase')
      plt.show()
```



```
[51]: print(df.groupby(by='Occupation').mean()['Purchase'].
      ↪sort_values(ascending=False))
```

```
Occupation
17    9758.679087
12    9717.192386
15    9691.443157
8      9479.513636
14    9431.785228
7      9365.188025
16    9328.979090
5      9283.268854
13    9194.099387
6      9191.133261
4      9161.759189
18    9113.718423
3      9112.929019
11    9108.370957
0      9057.506291
10     8912.886639
1      8876.457053
```

```

2      8869.590505
20     8746.967501
19     8627.760818
9       8607.359828
Name: Purchase, dtype: float64

```

Inference - We see that the highest average purchase price is for 17 years of occupation and lowest average purchase price is for 9 years of occupation. People having 4,0 and 7 years of occupation are the most frequent buyers.


```
[ ]:
```

0.1.10 Question 8 - What product categories are mostly bought by each of the gender?

```

[52]: plt.figure(figsize=(15,4))
      sns.heatmap(pd.
      ↪crosstab(index=df['Gender'],columns=df['Product_Category']),cmap='Blues')
      plt.show()

```



```

[53]: pd.crosstab(index=df['Gender'],columns=df['Product_Category'],normalize='index')

```

```

[53]: Product_Category      1      2      3      4      5      6  \
Gender
F      0.183634  0.041843  0.044417  0.026912  0.310317  0.033715
M      0.280338  0.044171  0.034469  0.019686  0.264385  0.038593

Product_Category      7      8      9      10     11     12  \
Gender
F      0.006974  0.248173  0.000414  0.004718  0.035047  0.011330
M      0.006740  0.194985  0.000677  0.005367  0.047427  0.005859

Product_Category      13     14     15     16     17     18  \
Gender
F      0.010812  0.004607  0.007358  0.017764  0.000459  0.002825

```


M	0.009916	0.002184	0.012053	0.018017	0.001252	0.006655
Product_Category	19	20				
Gender						
F	0.003335	0.005347				
M	0.002795	0.004433				

Inference - We observe that irrespective of the gender, people mostly buy the products 1, 5 and 8.

[]:

0.1.11 Question 9 - What product categories are mostly bought by each of the age categories?

```
[54]: plt.figure(figsize=(15,5))
sns.heatmap(pd.
↳ crosstab(index=df['Age'],columns=df['Product_Category'],cmap='Oranges')
plt.show())
```



```
[55]: pd.crosstab(index=df['Age'],columns=df['Product_Category'],normalize='index')
```

```
[55]: Product_Category    1         2         3         4         5         6  \
Age
0-17      0.238491  0.053552  0.079830  0.050426  0.288052  0.026543
18-25      0.271428  0.044577  0.047416  0.024795  0.287132  0.037741
26-35      0.266390  0.040830  0.035041  0.019171  0.281134  0.038804
36-45      0.252703  0.044896  0.035226  0.021516  0.268506  0.035637
46-50      0.230492  0.046323  0.030280  0.021786  0.263435  0.035694
51-55      0.236941  0.046634  0.024194  0.017753  0.259040  0.037967
55+        0.206876  0.042444  0.022840  0.014914  0.251712  0.040428
```

Product_Category	7	8	9	10	11	12	\
Age							
0-17	0.003526	0.150213	0.000998	0.003725	0.049228	0.008316	
18-25	0.004842	0.180311	0.000483	0.003584	0.046278	0.004419	
26-35	0.007551	0.202395	0.000563	0.004582	0.045157	0.005012	
36-45	0.007394	0.212926	0.000859	0.006416	0.045270	0.009085	
46-50	0.007196	0.234497	0.000572	0.006668	0.046301	0.011443	
51-55	0.006965	0.244560	0.000655	0.006415	0.038177	0.011338	
55+	0.006285	0.291155	0.000188	0.008723	0.026311	0.015946	

Product_Category	13	14	15	16	17	18	\
Age							
0-17	0.007451	0.002594	0.009713	0.015234	0.000399	0.001796	
18-25	0.007611	0.002315	0.009664	0.016087	0.000413	0.003413	
26-35	0.009586	0.002579	0.010345	0.018833	0.000581	0.004765	
36-45	0.011425	0.002852	0.012220	0.017869	0.001234	0.006416	
46-50	0.012125	0.003279	0.012477	0.019343	0.002091	0.007724	
51-55	0.012647	0.004032	0.012464	0.017596	0.002802	0.011076	
55+	0.014117	0.003517	0.010083	0.017681	0.003142	0.011303	


Product_Category	19	20
Age		
0-17	0.003925	0.005987
18-25	0.002768	0.004721
26-35	0.002575	0.004107
36-45	0.002925	0.004625
46-50	0.003279	0.004995
51-55	0.003509	0.005237
55+	0.004831	0.007504

Inference - We observe that irrespective of the age category, people mostly buy the products 1, 5 and 8.

[]:

0.1.12 Question 10 - What product categories are mostly bought by each of the city categories?

```
[56]: plt.figure(figsize=(15,5))
sns.heatmap(pd.
    ↳ crosstab(index=df['City_Category'], columns=df['Product_Category']), cmap='Reds')
plt.show()
```



```
[57]: pd.
      ↪ crosstab(index=df['City_Category'], columns=df['Product_Category'], normalize='index')
```

```
[57]: Product_Category      1      2      3      4      5      6  \
City_Category
A      0.238588  0.041765  0.033618  0.020743  0.287079  0.037453
B      0.253148  0.045386  0.037316  0.022710  0.278723  0.037051
C      0.276338  0.042757  0.039256  0.020424  0.261888  0.037788

Product_Category      7      8      9     10     11     12  \
City_Category
A      0.008338  0.218851  0.000667  0.005121  0.044894  0.007230
B      0.006949  0.206650  0.000626  0.005067  0.045564  0.007279
C      0.005263  0.200851  0.000546  0.005469  0.042299  0.007102

Product_Category     13     14     15     16     17     18  \
City_Category
A      0.010977  0.003271  0.011052  0.019369  0.000823  0.005121
B      0.009869  0.002746  0.010890  0.017548  0.001160  0.006036
C      0.009774  0.002408  0.010761  0.017281  0.001116  0.005774


Product_Category     19     20
City_Category
A      0.001857  0.003183
B      0.002008  0.003272
C      0.005099  0.007807
```

Inference - We observe that irrespective of the city category, people mostly buy the products 1, 5 and 8.

```
[ ]:
```

0.1.13 Question 11 - What product categories are mostly bought by each of the marital status categories?

```
[58]: plt.figure(figsize=(15,5))
sns.heatmap(pd.
↳ crosstab(index=df['Marital_Status'],columns=df['Product_Category']),cmap='Greens')
plt.show()
```



```
[59]: pd.
↳ crosstab(index=df['Marital_Status'],columns=df['Product_Category'],normalize='index')
```

```
[59]: Product_Category      1      2      3      4      5      6  \
Marital_Status
0      0.261027  0.043738  0.038235  0.022203  0.277365  0.037554
1      0.249847  0.043391  0.035039  0.020415  0.273376  0.037149

Product_Category      7      8      9     10     11     12  \
Marital_Status
0      0.006311  0.202359  0.000603  0.004724  0.045378  0.006292
1      0.007499  0.216436  0.000625  0.005902  0.042913  0.008535

Product_Category     13     14     15     16     17     18  \
Marital_Status
0      0.009782  0.002617  0.010633  0.017674  0.000922  0.005077
1      0.010649  0.003020  0.011269  0.018358  0.001249  0.006621


Product_Category     19     20
Marital_Status
0      0.002927  0.004579
1      0.002931  0.004774
```


Inference - We observe that irrespective of the marital status, people mostly buy the products 1, 5 and 8.

[]:

0.1.14 Question 12 - What product categories are mostly bought by each of the occupation categories?

```
[60]: plt.figure(figsize=(15,5))
sns.heatmap(pd.
↪crosstab(index=df['Occupation'],columns=df['Product_Category']),cmap='pink_r')
plt.show()
```



```
[61]: pd.
↪crosstab(index=df['Occupation'],columns=df['Product_Category'],normalize='index')
```

```
[61]: Product_Category      1      2      3      4      5      6  \
Occupation
0      0.254552  0.042865  0.038018  0.021368  0.273914  0.037051
1      0.218659  0.040743  0.033048  0.020626  0.279158  0.037457
2      0.217212  0.041990  0.036240  0.019368  0.287989  0.039304
3      0.224499  0.037910  0.034950  0.022655  0.300888  0.037056
4      0.267046  0.042185  0.045891  0.023640  0.283995  0.036882
5      0.300503  0.048545  0.036842  0.020770  0.270007  0.031237
6      0.227087  0.040521  0.037017  0.021963  0.280884  0.036277
7      0.272391  0.042650  0.027244  0.018667  0.265291  0.039117
8      0.332468  0.064286  0.043506  0.025325  0.240909  0.029221
9      0.180471  0.044600  0.057343  0.039822  0.343103  0.030742
10     0.253492  0.055866  0.082014  0.051831  0.281502  0.022346
11     0.230000  0.041391  0.029913  0.019043  0.294870  0.044435
12     0.298661  0.045878  0.038296  0.018487  0.259171  0.035393
13     0.201513  0.047215  0.019695  0.015912  0.257206  0.034303
```

14	0.276230	0.040629	0.028999	0.017407	0.281014	0.043462
15	0.279331	0.047907	0.036323	0.019030	0.263363	0.035744
16	0.257455	0.044751	0.030256	0.017623	0.269098	0.037385
17	0.320336	0.049454	0.037862	0.020374	0.244286	0.034826
18	0.265807	0.046247	0.024261	0.018954	0.264898	0.041698
19	0.246909	0.038635	0.031859	0.027104	0.281740	0.040181
20	0.206056	0.043832	0.037356	0.021166	0.294379	0.042332

Product_Category Occupation	7	8	9	10	11	12 \
0	0.008065	0.203246	0.000606	0.005107	0.054134	0.006420
1	0.007928	0.254335	0.000339	0.006275	0.035634	0.009200
2	0.008663	0.237034	0.000378	0.005599	0.043541	0.009192
3	0.007798	0.218750	0.000740	0.007001	0.039447	0.010189
4	0.006233	0.191977	0.000486	0.003526	0.044309	0.005164
5	0.003050	0.158988	0.000495	0.004286	0.065194	0.006429
6	0.007206	0.236168	0.000642	0.005182	0.042199	0.007946
7	0.005554	0.214352	0.000594	0.005146	0.040968	0.007898
8	0.003247	0.170779	0.000649	0.005195	0.040909	0.001948
9	0.002389	0.208187	0.000159	0.003504	0.025645	0.013061
10	0.003259	0.141217	0.000776	0.002871	0.048805	0.008070
11	0.007304	0.224087	0.000783	0.007217	0.034087	0.009130
12	0.005936	0.184933	0.000903	0.006033	0.040878	0.005646
13	0.006391	0.290726	0.000130	0.009260	0.032738	0.014478
14	0.005741	0.202959	0.000699	0.004269	0.037059	0.006035
15	0.007033	0.206437	0.000993	0.004882	0.033262	0.005213
16	0.007762	0.211239	0.000792	0.004911	0.049265	0.007485
17	0.004316	0.170920	0.000979	0.005194	0.045114	0.005796
18	0.005155	0.169219	0.000303	0.006217	0.091281	0.004852
19	0.007370	0.196624	0.000594	0.004280	0.063956	0.006895
20	0.011363	0.231240	0.000540	0.006716	0.042153	0.007585

Product_Category Occupation	13	14	15	16	17	18 \
0	0.010388	0.002828	0.009797	0.017674	0.000909	0.006146
1	0.011744	0.003582	0.010726	0.016619	0.001293	0.005045
2	0.009344	0.003064	0.010895	0.016947	0.000530	0.005712
3	0.012523	0.003871	0.009164	0.019467	0.000854	0.005749
4	0.008204	0.002499	0.009883	0.016546	0.000639	0.003720
5	0.007912	0.001896	0.011209	0.017555	0.000742	0.007830
6	0.010118	0.003554	0.009328	0.020532	0.000888	0.004491
7	0.010327	0.002718	0.010191	0.019839	0.001410	0.007830
8	0.005195	0.000000	0.014935	0.010390	0.000649	0.004545
9	0.009716	0.003664	0.005416	0.017999	0.000637	0.003664
10	0.007061	0.002250	0.010320	0.015518	0.000466	0.002017
11	0.010174	0.002609	0.009826	0.015739	0.001217	0.009217
12	0.009969	0.002323	0.015648	0.017551	0.000839	0.004291

13	0.014347	0.003522	0.010826	0.017738	0.002869	0.009000
14	0.010636	0.002024	0.009605	0.018953	0.000773	0.005594
15	0.010343	0.002234	0.013404	0.017127	0.001820	0.007364
16	0.010257	0.003129	0.011643	0.021029	0.001901	0.006970
17	0.011341	0.001907	0.013800	0.017940	0.001179	0.005796
18	0.010766	0.000910	0.009401	0.023199	0.001213	0.008036
19	0.008916	0.003447	0.012601	0.017237	0.000476	0.005825
20	0.010223	0.003628	0.011572	0.017509	0.001379	0.005396

Product_Category	19	20
Occupation		
0	0.002712	0.004199
1	0.002947	0.004642
2	0.002345	0.004653
3	0.002789	0.003700
4	0.002471	0.004706
5	0.001978	0.004533
6	0.003109	0.004886
7	0.003261	0.004552
8	0.003247	0.002597
9	0.003504	0.006371
10	0.004268	0.006052
11	0.003391	0.005565
12	0.003839	0.005323
13	0.003782	0.008347
14	0.003349	0.004563
15	0.002979	0.005213
16	0.002535	0.004515
17	0.003638	0.004943
18	0.003791	0.003791
19	0.001545	0.003804
20	0.001949	0.003628


Inference - We observe that irrespective of the occupation years, people mostly buy the products 1, 5 and 8.

[]:

0.1.15 Question 13 - What product categories are mostly bought by each of the Current City Stay Years categories?

[62]:

```
plt.figure(figsize=(15,5))
sns.heatmap(pd.
    ↳crosstab(index=df['Stay_In_Current_City_Years'],columns=df['Product_Category']),cmap='Green')
plt.show()
```



```
[63]: pd.
      ↪ crosstab(index=df['Stay_In_Current_City_Years'], columns=df['Product_Category'], normalize='i
```

```
[63]: Product_Category      1      2      3      4      5  \
Stay_In_Current_City_Years
0      0.250621  0.043222  0.037671  0.021436  0.278243
1      0.253151  0.043449  0.036340  0.021255  0.275361
2      0.261511  0.045471  0.038586  0.021049  0.277223
3      0.263502  0.044618  0.037815  0.022172  0.278174
4+     0.255094  0.040855  0.034617  0.021714  0.269835
```

```
Product_Category      6      7      8      9     10  \
Stay_In_Current_City_Years
0      0.036969  0.007537  0.207210  0.000648  0.004822
1      0.037678  0.006694  0.213405  0.000565  0.005263
2      0.036820  0.005928  0.199617  0.000671  0.005346
3      0.036296  0.007278  0.201194  0.000686  0.005200
4+     0.039005  0.006890  0.214867  0.000534  0.005254
```

```
Product_Category     11     12     13     14     15  \
Stay_In_Current_City_Years
0      0.047072  0.007267  0.010468  0.002917  0.011832
1      0.042236  0.007301  0.010402  0.002982  0.010760
2      0.045431  0.007230  0.009656  0.002456  0.010722
3      0.042235  0.007162  0.009947  0.002489  0.010580
4+     0.047995  0.006985  0.010033  0.002929  0.010934
```

```
Product_Category     16     17     18     19     20
Stay_In_Current_City_Years
0      0.017829  0.001162  0.005754  0.002688  0.004633
1      0.018305  0.000871  0.006264  0.003148  0.004568
2      0.017794  0.001134  0.005188  0.002920  0.005247
```


3	0.017130	0.001213	0.005158	0.002679	0.004472
4+	0.018382	0.001115	0.005645	0.002929	0.004388

Inference - We observe that irrespective of the current years stay, people mostly buy the products 1, 5 and 8.

[]:

0.1.16 Question 14 - Top 10 Product IDs and which product category they belong to?

```
[64]: plt.figure(figsize=(9,7))
sns.barplot(y=df['Product_ID'].value_counts()[:10],x=df['Product_ID'].
↪value_counts()[:10].index)
plt.xticks(rotation=30)
plt.show()
```



```
[65]: df.loc[df['Product_ID'].isin(df['Product_ID'].value_counts()[:10].
↪index), 'Product_Category'].unique()
```

```
[65]: array([1, 5, 6, 8], dtype=int64)
```

Inference - The top 10 product IDs belong to the product categories 1,5,6 and 8.

```
[ ]:
```


```
[ ]:
```

0.2 Confidence Interval Using The Central Limit Theorem For The Entire Population


- Let's assume that the population mean and standard deviation are finite.
- Let us draw 1 lakh samples from the given dataset, with each sample having 10000 data points.

```
[66]: data=df['Purchase']  
sampling_distribution=[]  
for i in range(100000):  
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[67]: #Plotting the sampling distribution  
sns.histplot(data=sampling_distribution,kde=True)  
plt.show()  
#It looks like this is a normal distribution. But we need to confirm if it is  
↪ so.  
#We can confirm using QQ Plot.
```



```
[68]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[69]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[70]: #68% Confidence Interval :
print('Population Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Population Mean Purchase 68% Confidence Interval : 9146.522418690758 to 9245.009750147245

```
[71]: #95% Confidence Interval :
print('Population Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```

Population Mean Purchase 95% Confidence Interval : 9097.278752962513 to 9294.25341587549


```
[ ]:
```

0.3 Confidence Interval Using The Central Limit Theorem For Males


- Let's assume that the population mean and standard deviation are finite.
- Let us draw 1 lakh samples from the given dataset, with each sample having 10000 data points.

```
[72]: data=df.loc[df['Gender']=='M', 'Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[73]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution, kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[74]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```

From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[75]: mean_sampling_distribution = np.mean(sampling_distribution)
      std_sampling_distribution=np.std(sampling_distribution)
```

```
[76]: #68% Confidence Interval :
      print('Male Mean Purchase 68% Confidence Interval :
            ↳',mean_sampling_distribution-std_sampling_distribution,
              'to',mean_sampling_distribution+std_sampling_distribution)
```

Male Mean Purchase 68% Confidence Interval : 9317.892020424133 to 9417.956468955868

```
[77]: #95% Confidence Interval :
      print('Male Mean Purchase 95% Confidence Interval :
            ↳',mean_sampling_distribution-2*std_sampling_distribution,
              'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Male Mean Purchase 95% Confidence Interval : 9267.859796158267 to 9467.988693221734

```
[ ]:
```


0.4 Confidence Interval Using The Central Limit Theorem For Females

```
[78]: data=df.loc[df['Gender']=='F', 'Purchase']  
sampling_distribution=[]  
for i in range(100000):  
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[79]: #Plotting the sampling distribution  
sns.histplot(data=sampling_distribution, kde=True)  
plt.show()  
#It looks like this is a normal distribution. But we need to confirm it is so.  
#We can confirm using QQ Plot.
```



```
[80]: #QQ Plot  
fig, ax1 = plt.subplots()  
plt.grid()  
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[81]: mean_sampling_distribution = np.mean(sampling_distribution)
      std_sampling_distribution=np.std(sampling_distribution)
```

```
[82]: #68% Confidence Interval :
      print('Female Mean Purchase 68% Confidence Interval :
            ↳',mean_sampling_distribution-std_sampling_distribution,
              'to',mean_sampling_distribution+std_sampling_distribution)
```

Female Mean Purchase 68% Confidence Interval : 8624.236932917702 to 8718.085025300297

```
[83]: #95% Confidence Interval :
      print('Female Mean Purchase 95% Confidence Interval :
            ↳',mean_sampling_distribution-2*std_sampling_distribution,
              'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Female Mean Purchase 95% Confidence Interval : 8577.312886726402 to 8765.009071491597

```
[ ]:
```


0.5 Confidence Interval Using The Central Limit Theorem For City A

```
[84]: data=df.loc[df['City_Category']=='A','Purchase']
      sampling_distribution=[]
      for i in range(100000):
          sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[85]: #Plotting the sampling distribution
      sns.histplot(data=sampling_distribution,kde=True)
      plt.show()
      #It looks like this is a normal distribution. But we need to confirm it is so.
      #We can confirm using QQ Plot.
```



```
[86]: #QQ Plot
      fig, ax1 = plt.subplots()
      plt.grid()
      prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[87]: mean_sampling_distribution = np.mean(sampling_distribution)
      std_sampling_distribution=np.std(sampling_distribution)
```

```
[88]: #68% Confidence Interval :
      print('City A Mean Purchase 68% Confidence Interval :
            ↳',mean_sampling_distribution-std_sampling_distribution,
              'to',mean_sampling_distribution+std_sampling_distribution)
```

City A Mean Purchase 68% Confidence Interval : 8797.471169466544 to 8893.498509299456

```
[89]: #95% Confidence Interval :
      print('City A Mean Purchase 95% Confidence Interval :
            ↳',mean_sampling_distribution-2*std_sampling_distribution,
              'to',mean_sampling_distribution+2*std_sampling_distribution)
```


City A Mean Purchase 95% Confidence Interval : 8749.457499550088 to 8941.512179215912

```
[ ]:
```


0.6 Confidence Interval Using The Central Limit Theorem For City B

```
[90]: data=df.loc[df['City_Category']=='B','Purchase']  
data  
sampling_distribution=[]  
for i in range(100000):  
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[91]: #Plotting the sampling distribution  
sns.histplot(data=sampling_distribution,kde=True)  
plt.show()  
#It looks like this is a normal distribution. But we need to confirm it is so.  
#We can confirm using QQ Plot.
```



```
[92]: #QQ Plot  
fig, ax1 = plt.subplots()  
plt.grid()  
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[93]: mean_sampling_distribution = np.mean(sampling_distribution)
      std_sampling_distribution=np.std(sampling_distribution)
```

```
[94]: #68% Confidence Interval :
      print('City B Mean Purchase 68% Confidence Interval :
            ↳',mean_sampling_distribution-std_sampling_distribution,
              'to',mean_sampling_distribution+std_sampling_distribution)
```

City B Mean Purchase 68% Confidence Interval : 9037.7509426006 to 9135.225770363399

```
[95]: #95% Confidence Interval :
      print('City B Mean Purchase 95% Confidence Interval :
            ↳',mean_sampling_distribution-2*std_sampling_distribution,
              'to',mean_sampling_distribution+2*std_sampling_distribution)
```


City B Mean Purchase 95% Confidence Interval : 8989.0135287192 to 9183.963184244798

```
[ ]:
```


0.7 Confidence Interval Using The Central Limit Theorem For City C

```
[96]: data=df.loc[df['City_Category']=='C','Purchase']
      sampling_distribution=[]
      for i in range(100000):
          sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[97]: #Plotting the sampling distribution
      sns.histplot(data=sampling_distribution,kde=True)
      plt.show()
      #It looks like this is a normal distribution. But we need to confirm it is so.
      #We can confirm using QQ Plot.
```



```
[98]: #QQ Plot
      fig, ax1 = plt.subplots()
      plt.grid()
      prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```

From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[99]: mean_sampling_distribution = np.mean(sampling_distribution)
      std_sampling_distribution=np.std(sampling_distribution)
```

```
[100]: #68% Confidence Interval :
print('City C Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

City C Mean Purchase 68% Confidence Interval : 9594.768033726903 to 9696.675063373097

```
[101]: #95% Confidence Interval :
print('City C Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```

City C Mean Purchase 95% Confidence Interval : 9543.814518903804 to 9747.628578196196

```
[ ]:
```

```
[ ]:
```

0.8 Confidence Interval Using The Central Limit Theorem For Age Group 0-17


```
[102]: print(df['Age'].nunique())  
# There are 7 unique age categories.  
print(df['Age'].unique().tolist())
```

7

```
['0-17', '55+', '26-35', '46-50', '51-55', '36-45', '18-25']
```


```
[103]: data=df.loc[df['Age']=='0-17','Purchase']  
sampling_distribution=[]  
for i in range(100000):  
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[104]: #Plotting the sampling distribution  
sns.histplot(data=sampling_distribution,kde=True)  
plt.show()  
#It looks like this is a normal distribution. But we need to confirm it is so.  
#We can confirm using QQ Plot.
```



```
[105]: #QQ Plot  
fig, ax1 = plt.subplots()  
plt.grid()
```

```
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[106]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[107]: #68% Confidence Interval :
print('Age Group 0-17 Mean Purchase 68% Confidence Interval :
      ↳',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Age Group 0-17 Mean Purchase 68% Confidence Interval : 8817.16969308802 to 8917.755087865979

```
[108]: #95% Confidence Interval :
print('Age Group 0-17 Mean Purchase 95% Confidence Interval :
      ↳',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Age Group 0-17 Mean Purchase 95% Confidence Interval : 8766.876995699042 to 8968.047785254957

```
[ ]:
```


0.9 Confidence Interval Using The Central Limit Theorem For Age Group 18-25

```
[109]: data=df.loc[df['Age']=='18-25','Purchase']
data
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[110]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution,kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[111]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[112]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[113]: #68% Confidence Interval :
print('Age Group 18-25 Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Age Group 18-25 Mean Purchase 68% Confidence Interval : 9074.407415798321 to 9174.163115891677

```
[114]: #95% Confidence Interval :
print('Age Group 18-25 Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Age Group 18-25 Mean Purchase 95% Confidence Interval : 9024.529565751642 to 9224.040965938357

```
[ ]:
```


0.10 Confidence Interval Using The Central Limit Theorem For Age Group 26-35

```
[115]: data=df.loc[df['Age']=='26-35','Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[116]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution,kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[117]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[118]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[119]: #68% Confidence Interval :
print('Age Group 26-35 Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Age Group 26-35 Mean Purchase 68% Confidence Interval : 9144.238069132618 to 9242.814961225387

```
[120]: #95% Confidence Interval :
print('Age Group 26-35 Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Age Group 26-35 Mean Purchase 95% Confidence Interval : 9094.949623086235 to 9292.10340727177

```
[ ]:
```


0.11 Confidence Interval Using The Central Limit Theorem For The Age Group 36-45

```
[121]: data=df.loc[df['Age']=='36-45','Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[122]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution,kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[123]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```

From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[124]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[125]: #68% Confidence Interval :
print('Age Group 36-45 Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Age Group 36-45 Mean Purchase 68% Confidence Interval : 9204.905608025932 to 9303.751895302066

```
[126]: #95% Confidence Interval :
print('Age Group 36-45 Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Age Group 36-45 Mean Purchase 95% Confidence Interval : 9155.482464387866 to 9353.175038940133

```
[ ]:
```


0.12 Confidence Interval Using The Central Limit Theorem For Age Group 46-50

```
[127]: data=df.loc[df['Age']=='46-50','Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[128]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution,kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[129]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[130]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[131]: #68% Confidence Interval :
print('Age Group 46-50 Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Age Group 46-50 Mean Purchase 68% Confidence Interval : 9080.259436324037 to 9177.665785283962

```
[132]: #95% Confidence Interval :
print('Age Group 46-50 Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Age Group 46-50 Mean Purchase 95% Confidence Interval : 9031.556261844073 to 9226.368959763926

```
[ ]:
```


0.13 Confidence Interval Using The Central Limit Theorem For Age Group 51-55

```
[133]: data=df.loc[df['Age']=='51-55','Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[134]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution,kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[135]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[136]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[137]: #68% Confidence Interval :
print('Age Group 51-55 Mean Purchase 68% Confidence Interval :
      ↳',mean_sampling_distribution-std_sampling_distribution,
        'to',mean_sampling_distribution+std_sampling_distribution)
```

Age Group 51-55 Mean Purchase 68% Confidence Interval : 9373.080889980605 to 9472.571365837395

```
[138]: #95% Confidence Interval :
print('Age Group 51-55 Mean Purchase 95% Confidence Interval :
      ↳',mean_sampling_distribution-2*std_sampling_distribution,
        'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Age Group 51-55 Mean Purchase 95% Confidence Interval : 9323.33565205221 to 9522.31660376579

```
[ ]:
```


0.14 Confidence Interval Using The Central Limit Theorem For Age Group 55+

```
[139]: data=df.loc[df['Age']=='55+', 'Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[140]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution, kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[141]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[142]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[143]: #68% Confidence Interval :
print('Age Group 55+ Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Age Group 55+ Mean Purchase 68% Confidence Interval : 9167.893286845856 to 9265.333105306143

```
[144]: #95% Confidence Interval :
print('Age Group 55+ Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```

Age Group 55+ Mean Purchase 95% Confidence Interval : 9119.173377615713 to 9314.053014536286

```
[ ]:
```


0.15 Confidence Interval Using The Central Limit Theorem For Stay In Current City 0 Year

```
[145]: print(df['Stay_In_Current_City_Years'].nunique())  
# There are 5 unique Stay In Current City Years.  
print(df['Stay_In_Current_City_Years'].unique().tolist())
```


```
5  
['2', '4+', '3', '1', '0']
```

```
[146]: data=df.loc[df['Stay_In_Current_City_Years']=='0','Purchase']  
sampling_distribution=[]  
for i in range(100000):  
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[147]: #Plotting the sampling distribution  
sns.histplot(data=sampling_distribution,kde=True)  
plt.show()  
#It looks like this is a normal distribution. But we need to confirm it is so.  
#We can confirm using QQ Plot.
```



```
[148]: #QQ Plot  
fig, ax1 = plt.subplots()  
plt.grid()  
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```

From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[149]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[150]: #68% Confidence Interval :
print('Current City Stay 0 Year Mean Purchase 68% Confidence Interval :
      ↳',mean_sampling_distribution-std_sampling_distribution,
        'to',mean_sampling_distribution+std_sampling_distribution)
```

Current City Stay 0 Year Mean Purchase 68% Confidence Interval :
9062.30545395957 to 9160.26637463843

```
[151]: #95% Confidence Interval :
print('Current City Stay 0 Year Mean Purchase 95% Confidence Interval :
      ↳',mean_sampling_distribution-2*std_sampling_distribution,
        'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Current City Stay 0 Year Mean Purchase 95% Confidence Interval :
9013.324993620141 to 9209.24683497786

```
[ ]:
```


0.16 Confidence Interval Using The Central Limit Theorem For Stay In Current City 1 Year

```
[152]: data=df.loc[df['Stay_In_Current_City_Years']=='1','Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[153]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution,kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[154]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[155]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[156]: #68% Confidence Interval :
print('Current City Stay 1 Year Mean Purchase 68% Confidence Interval :
      ↳',mean_sampling_distribution-std_sampling_distribution,
        'to',mean_sampling_distribution+std_sampling_distribution)
```

Current City Stay 1 Year Mean Purchase 68% Confidence Interval :
9129.580985775216 to 9228.605304208782

```
[157]: #95% Confidence Interval :
print('Current City Stay 1 Year Mean Purchase 95% Confidence Interval :
      ↳',mean_sampling_distribution-2*std_sampling_distribution,
        'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Current City Stay 1 Year Mean Purchase 95% Confidence Interval :
9080.068826558432 to 9278.117463425566

```
[ ]:
```


0.17 Confidence Interval Using The Central Limit Theorem For Stay In Current City 2 Years

```
[158]: data=df.loc[df['Stay_In_Current_City_Years']=='2','Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[159]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution,kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[160]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[161]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[162]: #68% Confidence Interval :
print('Current City Stay 2 Years Mean Purchase 68% Confidence Interval :
      ↳',mean_sampling_distribution-std_sampling_distribution,
        'to',mean_sampling_distribution+std_sampling_distribution)
```

Current City Stay 2 Years Mean Purchase 68% Confidence Interval :
9208.49904511685 to 9307.956619385148

```
[163]: #95% Confidence Interval :
print('Current City Stay 2 Years Mean Purchase 95% Confidence Interval :
      ↳',mean_sampling_distribution-2*std_sampling_distribution,
        'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Current City Stay 2 Years Mean Purchase 95% Confidence Interval :
9158.770257982702 to 9357.685406519297

```
[ ]:
```


0.18 Confidence Interval Using The Central Limit Theorem For Stay In Current City 3 Years

```
[164]: data=df.loc[df['Stay_In_Current_City_Years']=='3','Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[165]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution,kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[166]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[167]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[168]: #68% Confidence Interval :
print('Current City Stay 3 Years Mean Purchase 68% Confidence Interval :
      ↳',mean_sampling_distribution-std_sampling_distribution,
        'to',mean_sampling_distribution+std_sampling_distribution)
```

Current City Stay 3 Years Mean Purchase 68% Confidence Interval :
9166.689156775368 to 9265.08956621463

```
[169]: #95% Confidence Interval :
print('Current City Stay 3 Years Mean Purchase 95% Confidence Interval :
      ↳',mean_sampling_distribution-2*std_sampling_distribution,
        'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Current City Stay 3 Years Mean Purchase 95% Confidence Interval :
9117.488952055739 to 9314.28977093426

```
[ ]:
```


0.19 Confidence Interval Using The Central Limit Theorem For Stay In Current City 4+ Years

```
[170]: data=df.loc[df['Stay_In_Current_City_Years']=='4+', 'Purchase']
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[171]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution, kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[172]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```

From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[173]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[174]: #68% Confidence Interval :
print('Current City Stay 4+ Years Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Current City Stay 4+ Years Mean Purchase 68% Confidence Interval :
9159.369842172933 to 9257.65787423307

```
[175]: #95% Confidence Interval :
print('Current City Stay 4+ Years Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```

Current City Stay 4+ Years Mean Purchase 95% Confidence Interval :
9110.225826142865 to 9306.801890263137

```
[ ]:
```


0.20 Confidence Interval Using The Central Limit Theorem For Married People

```
[176]: print(df['Marital_Status'].unique())  
# There are 2 unique Marital Status categories.  
print(df['Marital_Status'].unique().tolist())
```


```
2  
[0, 1]
```

```
[177]: data=df.loc[df['Marital_Status']==1, 'Purchase']  
sampling_distribution=[]  
for i in range(100000):  
    sampling_distribution.append(np.mean(np.random.choice(data, size=10000)))
```

```
[178]: #Plotting the sampling distribution  
sns.histplot(data=sampling_distribution, kde=True)  
plt.show()  
#It looks like this is a normal distribution. But we need to confirm it is so.  
#We can confirm using QQ Plot.
```



```
[179]: #QQ Plot  
fig, ax1 = plt.subplots()  
plt.grid()  
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[180]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[181]: #68% Confidence Interval :
print('Married Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Married Mean Purchase 68% Confidence Interval : 9137.843988230801 to 9236.235351369203

```
[182]: #95% Confidence Interval :
print('Married Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```


Married Mean Purchase 95% Confidence Interval : 9088.6483066616 to 9285.431032938404

```
[ ]:
```


0.21 Confidence Interval Using The Central Limit Theorem For Unmarried People

```
[183]: Unmarried_data=df.loc[df['Marital_Status']==0,'Purchase']
Unmarried_data
sampling_distribution=[]
for i in range(100000):
    sampling_distribution.append(np.mean(np.random.choice(Unmarried_data,
↪size=10000)))

[184]: #Plotting the sampling distribution
sns.histplot(data=sampling_distribution,kde=True)
plt.show()
#It looks like this is a normal distribution. But we need to confirm it is so.
#We can confirm using QQ Plot.
```



```
[185]: #QQ Plot
fig, ax1 = plt.subplots()
plt.grid()
prob = stats.probplot(sampling_distribution, dist=stats.norm, plot=ax1)
```



From the above QQ PLOT, we can confirm that the sampling distribution follows a normal distribution, since almost all the points are lying on the 45 degree line.

```
[186]: mean_sampling_distribution = np.mean(sampling_distribution)
std_sampling_distribution=np.std(sampling_distribution)
```

```
[187]: #68% Confidence Interval :
print('Unmarried Mean Purchase 68% Confidence Interval :
      ↪',mean_sampling_distribution-std_sampling_distribution,
      'to',mean_sampling_distribution+std_sampling_distribution)
```

Unmarried Mean Purchase 68% Confidence Interval : 9152.00225431859 to 9251.066608847408

```
[188]: #95% Confidence Interval :
print('Unmarried Mean Purchase 95% Confidence Interval :
      ↪',mean_sampling_distribution-2*std_sampling_distribution,
      'to',mean_sampling_distribution+2*std_sampling_distribution)
```

Unmarried Mean Purchase 95% Confidence Interval : 9102.47007705418 to 9300.598786111817

```
[ ]:
```

1 Recommendations

1)

From the analysis done in this notebook, we observe that not only the most sold product categories are 5,1 and 8, but also irrespective of the gender, age category, city category, marital status, occupation years and current city stay years, people mostly buy the product categories 1, 5 and 8. Therefore Walmart should always try to keep abundant stock of these 3 product categories at all times.

2)

The top 10 product IDs belong to the product categories 1,5,6 and 8, which indicate that these product IDs are the favourites from their product categories. Therefore Walmart should always try to keep abundant stock of these product IDs from their categories at all times.

3) The Male Mean Purchase 95% Confidence Interval is from 9267.859796158267 to 9467.988693221734 and the Female Mean Purchase 95% Confidence Interval is from 8577.312886726402 to 8765.009071491597. ##### We see that the 95% confidence interval for the average amount spent by males is greater than females, and there is a clear distinction between the two. Therefore, we can expect that men will buy more on average than women. Also, the no of males making a purchase are 3 times of females making a purchase and the average spend by males is more than females, therefore Walmart should try to keep abundant stock of products which men usually buy and prefer.

4)

- Age Group 0-17 Mean Purchase 95% Confidence Interval : 8766.876995699042 to 8968.047785254957
- Age Group 18-25 Mean Purchase 95% Confidence Interval : 9024.529565751642 to 9224.040965938357
- Age Group 26-35 Mean Purchase 95% Confidence Interval : 9094.949623086235 to 9292.10340727177
- Age Group 36-45 Mean Purchase 95% Confidence Interval : 9155.482464387866 to 9353.175038940133
- Age Group 46-50 Mean Purchase 95% Confidence Interval : 9031.556261844073 to 9226.368959763926
- Age Group 51-55 Mean Purchase 95% Confidence Interval : 9323.33565205221 to 9522.31660376579
- Age Group 55+ Mean Purchase 95% Confidence Interval : 9119.173377615713 to 9314.053014536286 ##### We see that the 95% confidence interval for mean purchase price for the age categories 51-55 are greater than all other categories, and the age category 0-17 has the lowest 95% mean purchase price confidence interval. We can therefore assume that people who are 51-55 years old have their mean purchase price on the higher side and their confidence interval can be easily distinguished from the other age categories. People who are 0-17 years old have the lowest average purchase price, and their confidence interval can also be easily distinguished from the other age categories. Since most of the customers are in the

age category 18-25, 36-45 and 26-35, therefore Walmart should try to keep abundant stock of products which they usually buy and prefer.

5)

- Married Mean Purchase 95% Confidence Interval : 9088.6483066616 to 9285.431032938404
- Unmarried Mean Purchase 95% Confidence Interval : 9102.47007705418 to 9300.598786111817 ##### We see that the 95% confidence intervals for the mean purchase price for both married and unmarried are kind of similar, so we cannot find a distinction between their mean purchase prices. However, since the the number of unmarried people buying the products is almost 1.5x of that of who are married, therefore Walmart should try to keep abundant stock of products which unmarried people usually buy and prefer.

6)

- Current City Stay 0 Year Mean Purchase 95% Confidence Interval : 9013.324993620141 to 9209.24683497786
- Current City Stay 1 Year Mean Purchase 95% Confidence Interval : 9080.068826558432 to 9278.117463425566
- Current City Stay 2 Years Mean Purchase 95% Confidence Interval : 9158.770257982702 to 9357.685406519297
- Current City Stay 3 Years Mean Purchase 95% Confidence Interval : 9117.488952055739 to 9314.28977093426
- Current City Stay 4+ Years Mean Purchase 95% Confidence Interval : 9110.225826142865 to 9306.801890263137 ##### We observe that the 95% confidence intervals of average purchase price for different current city staying years are overlapping and therefore we cannot conclude that there is a distinction among all the categories. However, the number of people who make the most purchases are those that have lived in the current city for 1 year, and thus Walmart can keep this factor in account and stock its inventory accordingly.

7)

- City A Mean Purchase 95% Confidence Interval : 8749.457499550088 to 8941.512179215912
- City B Mean Purchase 95% Confidence Interval : 8989.0135287192 to 9183.963184244798
- City C Mean Purchase 95% Confidence Interval : 9543.814518903804 to 9747.628578196196 ##### We observe that the 95% confidence interval for the mean purchase price for city C is greater than city B which is in turn greater than city A. There is a clear distinction between their average purchase prices, therefore we can conclude that city C on average makes higher purchases than city B which in turn is higher than city A. Also, we observe that the number of people buying from City B is the highest, followed by C and A. Therefore Walmart should try to keep abundant stock of products firstly in city B, then C and then A.

8)

We see that the highest average purchase price is for 17 years of occupation, which might be because people having more occupation years might be earning more. Therefore, Walmart can expect that people having higher occupation years might have higher purchase price on average. Also, people having 0,4 and 7 years of occupation

are the most frequent buyers. Therefore Walmart should try to keep abundant stock of products which these people usually buy and prefer.

[]: