

Project_9_Ola

September 29, 2022

0.1 Problem Statement

Recruiting and retaining drivers is seen by industry watchers as a tough battle for Ola. Churn among drivers is high and it's very easy for drivers to stop working for the service on the fly or jump to Uber depending on the rates.

As the companies get bigger, the high churn could become a bigger problem. To find new drivers, Ola is casting a wide net, including people who don't have cars for jobs. But this acquisition is really costly. Losing drivers frequently impacts the morale of the organization and acquiring new drivers is more expensive than retaining existing ones.

You are working as a data scientist with the Analytics Department of Ola, focused on driver team attrition. You are provided with the monthly information for a segment of drivers for 2019 and 2020 and tasked to predict whether a driver will be leaving the company or not based on their attributes like

- Demographics (city, age, gender etc.)
- Tenure information (joining date, Last Date)
- Historical data regarding the performance of the driver (Quarterly rating, Monthly business acquired, grade, Income)

0.1.1 Column Profiling:

- MMMM-YY : Reporting Date (Monthly)
- Driver_ID : Unique id for drivers
- Age : Age of the driver
- Gender : Gender of the driver – Male : 0, Female: 1
- City : City Code of the driver
- Education_Level : Education level – 0 for 10+ ,1 for 12+ ,2 for graduate
- Income : Monthly average Income of the driver
- Date Of Joining : Joining date for the driver
- LastWorkingDate : Last date of working for the driver
- Joining Designation : Designation of the driver at the time of joining
- Grade : Grade of the driver at the time of reporting
- Total Business Value : The total business value acquired by the driver in a month (negative business indicates cancellation/refund or car EMI adjustments)
- Quarterly Rating : Quarterly rating of the driver: 1,2,3,4,5 (higher is better)

```
[1]: #Importing the required libraries
import numpy as np
```

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import ConfusionMatrixDisplay
from xgboost import XGBClassifier
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import PowerTransformer
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.metrics import roc_auc_score, roc_curve
import warnings
warnings.filterwarnings('ignore')

```

```

[2]: df=pd.read_csv('https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/
↳002/492/original/ola_driver_scaler.csv')
df.head()

```

```

[2]:      Unnamed: 0      MMM-YY  Driver_ID  Age  Gender  City  Education_Level  \
0              0  01/01/19           1  28.0    0.0  C23              2
1              1  02/01/19           1  28.0    0.0  C23              2
2              2  03/01/19           1  28.0    0.0  C23              2
3              3  11/01/20           2  31.0    0.0   C7              2
4              4  12/01/20           2  31.0    0.0   C7              2

      Income  Dateofjoining  LastWorkingDate  Joining  Designation  Grade  \
0    57387      24/12/18              NaN              1          1
1    57387      24/12/18              NaN              1          1
2    57387      24/12/18      03/11/19              1          1
3    67016      11/06/20              NaN              2          2
4    67016      11/06/20              NaN              2          2

      Total Business Value  Quarterly Rating
0              2381060              2
1             -665480              2
2              0              2
3              0              1
4              0              1

```

```

[3]: #Dropping the first column and checking the first 15 rows.
df=df.iloc[:,1:]

```

```
df.head(15)
```

```
[3]:
```

	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	\
0	01/01/19	1	28.0	0.0	C23	2	57387	
1	02/01/19	1	28.0	0.0	C23	2	57387	
2	03/01/19	1	28.0	0.0	C23	2	57387	
3	11/01/20	2	31.0	0.0	C7	2	67016	
4	12/01/20	2	31.0	0.0	C7	2	67016	
5	12/01/19	4	43.0	0.0	C13	2	65603	
6	01/01/20	4	43.0	0.0	C13	2	65603	
7	02/01/20	4	43.0	0.0	C13	2	65603	
8	03/01/20	4	43.0	0.0	C13	2	65603	
9	04/01/20	4	43.0	0.0	C13	2	65603	
10	01/01/19	5	29.0	0.0	C9	0	46368	
11	02/01/19	5	29.0	0.0	C9	0	46368	
12	03/01/19	5	29.0	0.0	C9	0	46368	
13	08/01/20	6	31.0	1.0	C11	1	78728	
14	09/01/20	6	31.0	1.0	C11	1	78728	

	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
0	24/12/18	NaN	1		1	
1	24/12/18	NaN	1		1	
2	24/12/18	03/11/19	1		1	
3	11/06/20	NaN	2		2	
4	11/06/20	NaN	2		2	
5	12/07/19	NaN	2		2	
6	12/07/19	NaN	2		2	
7	12/07/19	NaN	2		2	
8	12/07/19	NaN	2		2	
9	12/07/19	27/04/20	2		2	
10	01/09/19	NaN	1		1	
11	01/09/19	NaN	1		1	
12	01/09/19	03/07/19	1		1	
13	31/07/20	NaN	3		3	
14	31/07/20	NaN	3		3	

	Total Business Value	Quarterly Rating
0	2381060	2
1	-665480	2
2	0	2
3	0	1
4	0	1
5	0	1
6	0	1
7	0	1
8	350000	1
9	0	1

```

10          0          1
11      120360          1
12          0          1
13          0          1
14          0          1

```

```
[4]: #Checking the shape of the given dataset
df.shape
```

```
[4]: (19104, 13)
```

```
[5]: #Checking for percentage of missing values
np.round(df.isna().mean()*100,2)
```

```
[5]: MMM-YY          0.00
Driver_ID          0.00
Age                0.32
Gender             0.27
City               0.00
Education_Level    0.00
Income             0.00
Dateofjoining      0.00
LastWorkingDate    91.54
Joining Designation 0.00
Grade              0.00
Total Business Value 0.00
Quarterly Rating   0.00
dtype: float64
```

```
[6]: #Checking the data_type of each of the columns
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19104 entries, 0 to 19103
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MMM-YY                 19104 non-null  object
1   Driver_ID              19104 non-null  int64
2   Age                    19043 non-null  float64
3   Gender                 19052 non-null  float64
4   City                   19104 non-null  object
5   Education_Level        19104 non-null  int64
6   Income                 19104 non-null  int64
7   Dateofjoining          19104 non-null  object
8   LastWorkingDate        1616 non-null   object
9   Joining Designation    19104 non-null  int64
10  Grade                  19104 non-null  int64

```

```

11 Total Business Value 19104 non-null int64
12 Quarterly Rating      19104 non-null int64
dtypes: float64(2), int64(7), object(4)
memory usage: 1.9+ MB

```

```

[7]: #Checking for duplicate observations.
df.duplicated().sum()

```

```
[7]: 0
```

0.1.2 Preliminary Observations

- There are a few missing values in columns - “Age” and “Gender”. We will take care of missing values later.
- There are a lot of missing values in column - “LastWorkingDate”. From this column, we can get our target variable. The drivers who have churned have a LastWorkingDate.
- “MMM-YY”, “Dateofjoining” and “LastWorkingDate” columns are in object data type. We might need to change to date_type format.
- Gender is a categorical variables as given in the problem statement, but is represented in float. We can convert to integer later.
- There are no duplicate observations.
- The report is generated on the 1st of every month.

```
[ ]:
```

```

[8]: #Creating Another DataFrame to store Aggregate Level Information for each driver
dataset=pd.DataFrame(data=df['Driver_ID'].unique(),columns=['Driver_ID'])
print(dataset.shape) #There are total 2381 drivers.
dataset.head()

```

```
(2381, 1)
```

```

[8]: Driver_ID
0         1
1         2
2         4
3         5
4         6

```

```

[9]: #Getting the list of drivers who have churned and who have not_churned
churned_drivers=set(df.loc[~df['LastWorkingDate'].isna(),'Driver_ID'])
not_churned_drivers=set(df['Driver_ID']).difference(churned_drivers)

```

```
[ ]:
```

0.1.3 Getting the Target Column - “Churn”

```
[10]: def get_churn(x):  
        if x in churned_drivers:  
            return 1  
        else:  
            return 0  
dataset['Churn']=dataset['Driver_ID'].apply(get_churn)  
dataset.head()
```

```
[10]:   Driver_ID  Churn  
0         1      1  
1         2      0  
2         4      1  
3         5      1  
4         6      0
```

```
[11]: print("Number Of Churned Drivers",(dataset['Churn']==1).sum())  
print("Number Of Not_Churned Drivers",(dataset['Churn']==0).sum())
```

Number Of Churned Drivers 1616
Number Of Not_Churned Drivers 765

```
[12]: df.loc[~df['LastWorkingDate'].isna(),'Churn']=1  
df.loc[df['LastWorkingDate'].isna(),'Churn']=0  
df.head()
```

```
[12]:   MMM-YY  Driver_ID  Age  Gender  City  Education_Level  Income  \  
0  01/01/19         1  28.0    0.0  C23                2   57387  
1  02/01/19         1  28.0    0.0  C23                2   57387  
2  03/01/19         1  28.0    0.0  C23                2   57387  
3  11/01/20         2  31.0    0.0   C7                2   67016  
4  12/01/20         2  31.0    0.0   C7                2   67016
```

```
   Dateofjoining  LastWorkingDate  Joining Designation  Grade  \  
0    24/12/18             NaN                1      1  
1    24/12/18             NaN                1      1  
2    24/12/18    03/11/19                1      1  
3    11/06/20             NaN                2      2  
4    11/06/20             NaN                2      2
```

```
   Total Business Value  Quarterly Rating  Churn  
0          2381060                2    0.0  
1         -665480                2    0.0  
2              0                2    1.0  
3              0                1    0.0  
4              0                1    0.0
```

[]:

0.1.4 Getting The Duration For Which Each Driver Has Been Driving Vehicles

- For drivers who have churned, we can get the difference in no of days between DateOfJoining and LastWorkingDate
- For drivers who have not churned, we can get the difference in no of days between DateOfJoining and the Last_Monthly_Report_Date
- My expectation is that the last month report date is same for all not_churned drivers. Need to check.

```
[13]: #Converting "MMM-YY", "Dateofjoining" and "LastWorkingDate" columns to date_time
      ↪format
df[["MMM-YY", "Dateofjoining", "LastWorkingDate"]] = df[["MMM-YY", "Dateofjoining", "LastWorkingDate"]
      ↪astype('datetime64')
```

```
[14]: #Finding out the last_working_date for churned_drivers.
temp=df.loc[df['Churn']==1, ['Driver_ID', 'LastWorkingDate']]

#Merging with the new Dataset.
dataset=pd.merge(left=dataset, right=temp, on='Driver_ID', how='left')
dataset.head()
```

```
[14]:   Driver_ID  Churn  LastWorkingDate
0         1      1      2019-03-11
1         2      0              NaT
2         4      1      2020-04-27
3         5      1      2019-03-07
4         6      0              NaT
```

```
[15]: #Finding out the latest date for monthly report for not_churned_drivers
last_report_dates=df[df['Driver_ID'].isin(not_churned_drivers)].
      ↪groupby(by='Driver_ID').max()['MMM-YY'].reset_index()

#Merging with the new Dataset.
dataset=pd.merge(left=dataset, right=last_report_dates, on='Driver_ID', how='left')
dataset.head()
```

```
[15]:   Driver_ID  Churn  LastWorkingDate  MMM-YY
0         1      1      2019-03-11      NaT
1         2      0              NaT  2020-12-01
2         4      1      2020-04-27      NaT
3         5      1      2019-03-07      NaT
4         6      0              NaT  2020-12-01
```

```
[16]: # Lets check if the last date for monthly report for all not_churned_drivers is
      ↪the same.
dataset['MMM-YY'].value_counts()
```

```
# This is strange. The last date should be the same for all drivers who have,
↳ not churned. Lets inspect further.
# There are few drivers for whom the last report date is not "2020-12-01",
↳ which is the most occurring value.
```

```
[16]: 2020-12-01    741
      2020-02-01     6
      2020-05-01     6
      2019-12-01     4
      2020-01-01     3
      2020-03-01     2
      2020-06-01     2
      2019-09-01     1
      Name: MMM-YY, dtype: int64
```

```
[17]: # Lets find out the report for the drivers whose last month report date is,
      ↳ "2020-02-01". Expecting 6 drivers.
      temp_driver_ids=last_report_dates.
      ↳ loc[last_report_dates['MMM-YY']=='2020-02-01','Driver_ID'].tolist()
      print(len(temp_driver_ids))

      df[df['Driver_ID'].isin(temp_driver_ids)]
      # From the below table it is confirmed that the 6 drivers haven't churned,
      ↳ yet, but still don't have records till "2020-12-01"
      # We can consider dropping such driver_ids.
```

6

```
[17]:
```

	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	\
433	2020-01-01	66	27.0	1.0	C4	2	104286	
434	2020-02-01	66	27.0	1.0	C4	2	104286	
4209	2019-11-01	612	31.0	0.0	C17	0	29685	
4210	2019-12-01	612	31.0	0.0	C17	0	29685	
4211	2020-01-01	612	31.0	0.0	C17	0	29685	
4212	2020-02-01	612	31.0	0.0	C17	0	29685	
5054	2020-01-01	755	36.0	0.0	C8	1	25894	
5055	2020-02-01	755	36.0	0.0	C8	1	25894	
8226	2019-09-01	1224	33.0	0.0	C28	2	78442	
8227	2019-10-01	1224	33.0	0.0	C28	2	78442	
8228	2019-11-01	1224	33.0	0.0	C28	2	78442	
8229	2019-12-01	1224	33.0	0.0	C28	2	78442	
8230	2020-01-01	1224	33.0	0.0	C28	2	78442	
8231	2020-02-01	1224	33.0	0.0	C28	2	78442	
12788	2019-01-01	1893	38.0	0.0	C26	1	132819	
12789	2019-02-01	1893	38.0	0.0	C26	1	132819	
12790	2019-03-01	1893	38.0	0.0	C26	1	132819	

12791	2019-04-01	1893	38.0	0.0	C26	1	132819
12792	2019-05-01	1893	38.0	0.0	C26	1	132819
12793	2019-06-01	1893	38.0	0.0	C26	1	132819
12794	2019-07-01	1893	38.0	0.0	C26	1	132819
12795	2019-08-01	1893	38.0	0.0	C26	1	132819
12796	2019-09-01	1893	38.0	0.0	C26	1	132819
12797	2019-10-01	1893	38.0	0.0	C26	1	132819
12798	2019-11-01	1893	38.0	0.0	C26	1	132819
12799	2019-12-01	1893	38.0	0.0	C26	1	132819
12800	2020-01-01	1893	39.0	0.0	C26	1	132819
12801	2020-02-01	1893	39.0	0.0	C26	1	132819
12802	2020-02-01	1894	33.0	1.0	C7	2	51264

	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
433	2020-01-21	NaT	3		3	
434	2020-01-21	NaT	3		3	
4209	2019-11-03	NaT	1		1	
4210	2019-11-03	NaT	1		1	
4211	2019-11-03	NaT	1		1	
4212	2019-11-03	NaT	1		1	
5054	2019-12-30	NaT	3		3	
5055	2019-12-30	NaT	3		3	
8226	2019-09-03	NaT	3		3	
8227	2019-09-03	NaT	3		3	
8228	2019-09-03	NaT	3		3	
8229	2019-09-03	NaT	3		3	
8230	2019-09-03	NaT	3		3	
8231	2019-09-03	NaT	3		3	
12788	2017-01-20	NaT	5		5	
12789	2017-01-20	NaT	5		5	
12790	2017-01-20	NaT	5		5	
12791	2017-01-20	NaT	5		5	
12792	2017-01-20	NaT	5		5	
12793	2017-01-20	NaT	5		5	
12794	2017-01-20	NaT	5		5	
12795	2017-01-20	NaT	5		5	
12796	2017-01-20	NaT	5		5	
12797	2017-01-20	NaT	5		5	
12798	2017-01-20	NaT	5		5	
12799	2017-01-20	NaT	5		5	
12800	2017-01-20	NaT	5		5	
12801	2017-01-20	NaT	5		5	
12802	2020-02-09	NaT	3		3	

	Total Business Value	Quarterly Rating	Churn
433	0	1	0.0
434	0	1	0.0

4209	0	1	0.0
4210	200000	1	0.0
4211	0	1	0.0
4212	0	1	0.0
5054	0	1	0.0
5055	0	1	0.0
8226	0	1	0.0
8227	0	2	0.0
8228	664800	2	0.0
8229	374910	2	0.0
8230	3318790	3	0.0
8231	0	3	0.0
12788	0	1	0.0
12789	0	1	0.0
12790	0	1	0.0
12791	0	1	0.0
12792	0	1	0.0
12793	0	1	0.0
12794	830060	2	0.0
12795	32730	2	0.0
12796	812880	2	0.0
12797	0	1	0.0
12798	0	1	0.0
12799	0	1	0.0
12800	0	1	0.0
12801	0	1	0.0
12802	0	1	0.0

```
[18]: driver_ids_to_be_dropped=list(last_report_dates.loc[last_report_dates['MMM-YY'].
↳isin(last_report_dates['MMM-YY'].value_counts()[1:].index),'Driver_ID'])

#Total Drivers to be dropped
len(driver_ids_to_be_dropped)

print('As % of total drivers -',(100*len(driver_ids_to_be_dropped))/
↳len(dataset['Driver_ID']))
# There are about 1% not_churned drivers for whom we do not have complete_
↳monthly date. We can remove these drivers.
```

As % of total drivers - 1.0079798404031919

```
[19]: #Dropping observations of the 24 strange drivers from both df and dataset.
dataset=dataset[~dataset['Driver_ID'].isin(driver_ids_to_be_dropped)].
↳reset_index(drop=True)
df=df[~df['Driver_ID'].isin(driver_ids_to_be_dropped)].reset_index(drop=True)
```

```
[ ]:
```

[20]: *#Combining columns "LastWorkingDate" and "MMM-YY" into 1 column.*

```
def function(dataset):
    arr=[]
    Churn=dataset['Churn']
    LastWorkingDate=dataset['LastWorkingDate']
    MMM_YY=dataset['MMM-YY']
    for i,value in enumerate(Churn):
        if value==1:
            arr.append(LastWorkingDate[i])
        else:
            arr.append(MMM_YY[i])
    return arr
dataset['Last_Date']=function(dataset)
dataset.head(8)
```

[20]:

	Driver_ID	Churn	LastWorkingDate	MMM-YY	Last_Date
0	1	1	2019-03-11	NaT	2019-03-11
1	2	0	NaT	2020-12-01	2020-12-01
2	4	1	2020-04-27	NaT	2020-04-27
3	5	1	2019-03-07	NaT	2019-03-07
4	6	0	NaT	2020-12-01	2020-12-01
5	8	1	2020-11-15	NaT	2020-11-15
6	11	0	NaT	2020-12-01	2020-12-01
7	12	1	2019-12-21	NaT	2019-12-21

[]:

[21]: *#Getting "DateOfJoining" column for every driver*

```
temp=df.loc[~df[['Driver_ID','Dateofjoining']].
↳duplicated(),['Driver_ID','Dateofjoining']]
temp.head()

#Merging with the new Dataset.
dataset=pd.merge(left=dataset,right=temp,on='Driver_ID')
dataset.head()
```

[21]:

	Driver_ID	Churn	LastWorkingDate	MMM-YY	Last_Date	Dateofjoining
0	1	1	2019-03-11	NaT	2019-03-11	2018-12-24
1	2	0	NaT	2020-12-01	2020-12-01	2020-11-06
2	4	1	2020-04-27	NaT	2020-04-27	2019-12-07
3	5	1	2019-03-07	NaT	2019-03-07	2019-01-09
4	6	0	NaT	2020-12-01	2020-12-01	2020-07-31

[22]: *#Calculating the total duration of driving for each driver in days by*

```
↳subtracting "Dateofjoining" from "Last_Date"
dataset['Duration']=(dataset['Last_Date']-dataset['Dateofjoining']).dt.days
dataset.head()
```


```
[22]: Driver_ID  Churn  LastWorkingDate      MMM-YY  Last_Date  Dateofjoining  \
0         1      1      2019-03-11          NaT  2019-03-11    2018-12-24
1         2      0           NaT  2020-12-01  2020-12-01    2020-11-06
2         4      1      2020-04-27          NaT  2020-04-27    2019-12-07
3         5      1      2019-03-07          NaT  2019-03-07    2019-01-09
4         6      0           NaT  2020-12-01  2020-12-01    2020-07-31
```

```
Duration
0      77
1      25
2     142
3      57
4     123
```

```
[23]: # Lets check if the "Duration" Column has correct data.
print(dataset[['Duration']].describe().T)
sns.histplot(dataset['Duration'])
plt.show()

# There are some very high values for "Duration" which can be dealt with later.
# But "Duration" feature also has negative values. But time can never be
→negative. Lets inspect further.
```

```
count      mean      std  min  25%  50%  75%  max
Duration  2357.0  425.196012  564.592189 -27.0  92.0  182.0  467.0  2801.0
```



```
[24]: #Checking for observations which have negative values.
dataset[dataset['Duration']<0]
# The below records are of drivers who have not churned yet, but have joined OLA
↳ in Dec_2020, but the last_monthly_report date is "2020-12-01"
# This is is strange, because the report would get generated on 1st of every
↳ month, which is before their joining date.
# We can therefore drop these drivers because of this incorrect data
```

```
[24]:
```

	Driver_ID	Churn	LastWorkingDate	MMM-YY	Last_Date	Dateofjoining	\
6	11	0	NaT	2020-12-01	2020-12-01	2020-12-07	
124	148	0	NaT	2020-12-01	2020-12-01	2020-12-11	
175	205	0	NaT	2020-12-01	2020-12-01	2020-12-20	
228	274	0	NaT	2020-12-01	2020-12-01	2020-12-20	
240	288	0	NaT	2020-12-01	2020-12-01	2020-12-18	
249	297	0	NaT	2020-12-01	2020-12-01	2020-12-19	
257	309	0	NaT	2020-12-01	2020-12-01	2020-12-11	
258	310	0	NaT	2020-12-01	2020-12-01	2020-12-11	
319	378	0	NaT	2020-12-01	2020-12-01	2020-12-13	
333	398	0	NaT	2020-12-01	2020-12-01	2020-12-07	
431	508	0	NaT	2020-12-01	2020-12-01	2020-12-14	
434	511	0	NaT	2020-12-01	2020-12-01	2020-12-15	
461	540	0	NaT	2020-12-01	2020-12-01	2020-12-05	
480	561	0	NaT	2020-12-01	2020-12-01	2020-12-21	
493	575	0	NaT	2020-12-01	2020-12-01	2020-12-21	
514	599	0	NaT	2020-12-01	2020-12-01	2020-12-14	
535	622	0	NaT	2020-12-01	2020-12-01	2020-12-25	
569	664	0	NaT	2020-12-01	2020-12-01	2020-12-11	
607	711	0	NaT	2020-12-01	2020-12-01	2020-12-28	
611	715	0	NaT	2020-12-01	2020-12-01	2020-12-19	
623	733	0	NaT	2020-12-01	2020-12-01	2020-12-14	
678	802	0	NaT	2020-12-01	2020-12-01	2020-12-12	
685	810	0	NaT	2020-12-01	2020-12-01	2020-12-04	
690	816	0	NaT	2020-12-01	2020-12-01	2020-12-12	
696	824	0	NaT	2020-12-01	2020-12-01	2020-12-14	
697	825	0	NaT	2020-12-01	2020-12-01	2020-12-11	
705	834	0	NaT	2020-12-01	2020-12-01	2020-12-07	
770	913	0	NaT	2020-12-01	2020-12-01	2020-12-15	
868	1029	0	NaT	2020-12-01	2020-12-01	2020-12-18	
975	1157	0	NaT	2020-12-01	2020-12-01	2020-12-07	
980	1162	0	NaT	2020-12-01	2020-12-01	2020-12-14	
1005	1194	0	NaT	2020-12-01	2020-12-01	2020-12-13	
1006	1195	0	NaT	2020-12-01	2020-12-01	2020-12-18	
1021	1210	0	NaT	2020-12-01	2020-12-01	2020-12-19	
1073	1273	0	NaT	2020-12-01	2020-12-01	2020-12-07	
1143	1358	0	NaT	2020-12-01	2020-12-01	2020-12-05	
1281	1521	0	NaT	2020-12-01	2020-12-01	2020-12-07	
1317	1568	0	NaT	2020-12-01	2020-12-01	2020-12-11	

1363	1619	0	NaT	2020-12-01	2020-12-01	2020-12-08
1374	1632	0	NaT	2020-12-01	2020-12-01	2020-12-06
1389	1650	0	NaT	2020-12-01	2020-12-01	2020-12-11
1437	1708	0	NaT	2020-12-01	2020-12-01	2020-12-15
1467	1740	0	NaT	2020-12-01	2020-12-01	2020-12-05
1522	1812	0	NaT	2020-12-01	2020-12-01	2020-12-13
1584	1880	0	NaT	2020-12-01	2020-12-01	2020-12-11
1618	1925	0	NaT	2020-12-01	2020-12-01	2020-12-15
1642	1955	0	NaT	2020-12-01	2020-12-01	2020-12-11
1654	1971	0	NaT	2020-12-01	2020-12-01	2020-12-20
1706	2030	0	NaT	2020-12-01	2020-12-01	2020-12-05
1715	2039	0	NaT	2020-12-01	2020-12-01	2020-12-15
1732	2060	0	NaT	2020-12-01	2020-12-01	2020-12-18
1754	2088	0	NaT	2020-12-01	2020-12-01	2020-12-27
1883	2233	0	NaT	2020-12-01	2020-12-01	2020-12-19
1907	2261	0	NaT	2020-12-01	2020-12-01	2020-12-21
1937	2298	0	NaT	2020-12-01	2020-12-01	2020-12-08
1986	2360	0	NaT	2020-12-01	2020-12-01	2020-12-14
2130	2521	0	NaT	2020-12-01	2020-12-01	2020-12-18
2223	2631	0	NaT	2020-12-01	2020-12-01	2020-12-07

Duration	
6	-6
124	-10
175	-19
228	-19
240	-17
249	-18
257	-10
258	-10
319	-12
333	-6
431	-13
434	-14
461	-4
480	-20
493	-20
514	-13
535	-24
569	-10
607	-27
611	-18
623	-13
678	-11
685	-3
690	-11
696	-13

697	-10
705	-6
770	-14
868	-17
975	-6
980	-13
1005	-12
1006	-17
1021	-18
1073	-6
1143	-4
1281	-6
1317	-10
1363	-7
1374	-5
1389	-10
1437	-14
1467	-4
1522	-12
1584	-10
1618	-14
1642	-10
1654	-19
1706	-4
1715	-14
1732	-17
1754	-26
1883	-18
1907	-20
1937	-7
1986	-13
2130	-17
2223	-6

```
[25]: # Getting the list of above Drivers.
drivers_to_drop=list(set(dataset.loc[dataset['Duration']<0,'Driver_ID']))

# Lets have a look at few of the records of these drivers.
df[df['Driver_ID'].isin(drivers_to_drop)][:10]
# The same issue can be seen from below records as well. The joining date is_
→after the date of report generation.
```

```
[25]:
```

	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	\
21	2020-12-01	11	28.0	1.0	C19	2	42172	
994	2020-12-01	148	41.0	0.0	C15	0	45288	
1396	2020-12-01	205	39.0	0.0	C12	0	15826	
1774	2020-12-01	274	39.0	0.0	C7	1	53224	

1852	2020-12-01	288	41.0	1.0	C17	0	62901
1907	2020-12-01	297	38.0	0.0	C8	2	109473
1974	2020-12-01	309	30.0	1.0	C5	1	77597
1975	2020-12-01	310	27.0	1.0	C3	2	58439
2466	2020-12-01	378	33.0	0.0	C9	2	35057
2584	2020-12-01	398	30.0	0.0	C5	2	63280

	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
21	2020-12-07	NaT	1		1	
994	2020-12-11	NaT	2		2	
1396	2020-12-20	NaT	1		1	
1774	2020-12-20	NaT	2		2	
1852	2020-12-18	NaT	2		2	
1907	2020-12-19	NaT	3		3	
1974	2020-12-11	NaT	3		3	
1975	2020-12-11	NaT	1		1	
2466	2020-12-13	NaT	3		3	
2584	2020-12-07	NaT	2		2	

	Total Business Value	Quarterly Rating	Churn
21	0	1	0.0
994	0	1	0.0
1396	0	1	0.0
1774	0	1	0.0
1852	0	1	0.0
1907	0	1	0.0
1974	0	1	0.0
1975	0	1	0.0
2466	0	1	0.0
2584	0	1	0.0

```
[26]: # Lets drop the above drivers from both our datasets.
dataset=dataset[~dataset['Driver_ID'].isin(drivers_to_drop)]
df=df[~df['Driver_ID'].isin(drivers_to_drop)]
dataset.reset_index(drop=True,inplace=True)
df.reset_index(drop=True,inplace=True)
```

```
[27]: # Dropping the columns - 'LastWorkingDate', 'MMM-YY', 'Last_Date'
↳and'Dateofjoining'
dataset.
↳drop(columns=['LastWorkingDate','MMM-YY','Last_Date','Dateofjoining'],inplace=True)
```

```
[28]: # A look at our updated dataset.
dataset.head(3)
```

```
[28]: Driver_ID  Churn  Duration
0          1      1         77
```


1	2	0	25
2	4	1	142

[]:

0.1.5 Inspecting and Aggregating “City” for each Driver in the new Dataset

```
[29]: #Lets do a quick check if there there are multiple cities assigned to each
      ↪ driver
      df.groupby(by='Driver_ID').nunique()['City'].value_counts()

      #We observe that all drivers have only 1 city mapped against them. Lets create
      ↪ a new feature in our aggregated
      # dataset and store the city for each driver
```

```
[29]: 1    2299
      Name: City, dtype: int64
```

```
[30]: #Getting the city for each driver
      cities=df.loc[~df[['Driver_ID','City']].duplicated(),['Driver_ID','City']]
      cities.head()
```

```
[30]:   Driver_ID City
      0         1  C23
      3         2   C7
      5         4  C13
     10         5   C9
     13         6  C11
```

```
[31]: dataset = pd.merge(left=dataset,right=cities,on='Driver_ID')
      dataset.head()
      #Sucessfully created a new feature which stores the city for each driver.
```

```
[31]:   Driver_ID  Churn  Duration City
      0         1      1         77  C23
      1         2      0         25   C7
      2         4      1        142  C13
      3         5      1         57   C9
      4         6      0        123  C11
```

[]:

0.1.6 Inspecting and Aggregating “Education_Level” for each Driver in the new Dataset

```
[32]: #Lets do a quick check if there are multiple cities assigned to each
      ↪ driver
df.groupby(by='Driver_ID').nunique()['Education_Level'].value_counts()

#We observe that all drivers have only 1 Education_Level against them. Lets
      ↪ create a new feature in our aggregated
# dataset and store the Education_Level for each driver
```

```
[32]: 1    2299
      Name: Education_Level, dtype: int64
```

```
[33]: # Lets look at the different values of Education_Level
df['Education_Level'].unique()
```

```
[33]: array([2, 0, 1], dtype=int64)
```

```
[34]: #Getting the Education_Level for each driver
education_level=df.loc[~df[['Driver_ID','Education_Level']].
      ↪ duplicated(),['Driver_ID','Education_Level']]
education_level.head()
```

```
[34]:   Driver_ID  Education_Level
0         1             2
3         2             2
5         4             2
10        5             0
13        6             1
```

```
[35]: dataset = pd.merge(left=dataset,right=education_level,on='Driver_ID')
dataset.head()
#Successfully created a new feature which stores the city for each driver.
```

```
[35]:   Driver_ID  Churn  Duration  City  Education_Level
0         1      1         77  C23             2
1         2      0         25  C7              2
2         4      1        142  C13             2
3         5      1         57  C9              0
4         6      0        123  C11             1
```

```
[ ]:
```

0.1.7 Inspecting and Aggregating “Income” for each Driver in the new Dataset

```
[36]: # Lets have a look at few observation to understand how average income for each
      ↪ driver is stored in the dataset.
df.head(15)

#It seems that the average income is repeated for every observation of each
↪ driver.
```

```
[36]:      MMM-YY  Driver_ID  Age  Gender  City  Education_Level  Income  \
0  2019-01-01         1  28.0    0.0  C23             2  57387
1  2019-02-01         1  28.0    0.0  C23             2  57387
2  2019-03-01         1  28.0    0.0  C23             2  57387
3  2020-11-01         2  31.0    0.0   C7             2  67016
4  2020-12-01         2  31.0    0.0   C7             2  67016
5  2019-12-01         4  43.0    0.0  C13             2  65603
6  2020-01-01         4  43.0    0.0  C13             2  65603
7  2020-02-01         4  43.0    0.0  C13             2  65603
8  2020-03-01         4  43.0    0.0  C13             2  65603
9  2020-04-01         4  43.0    0.0  C13             2  65603
10 2019-01-01         5  29.0    0.0   C9             0  46368
11 2019-02-01         5  29.0    0.0   C9             0  46368
12 2019-03-01         5  29.0    0.0   C9             0  46368
13 2020-08-01         6  31.0    1.0  C11             1  78728
14 2020-09-01         6  31.0    1.0  C11             1  78728
```

```
      Dateofjoining  LastWorkingDate  Joining  Designation  Grade  \
0    2018-12-24             NaT             1             1
1    2018-12-24             NaT             1             1
2    2018-12-24    2019-03-11             1             1
3    2020-11-06             NaT             2             2
4    2020-11-06             NaT             2             2
5    2019-12-07             NaT             2             2
6    2019-12-07             NaT             2             2
7    2019-12-07             NaT             2             2
8    2019-12-07             NaT             2             2
9    2019-12-07    2020-04-27             2             2
10   2019-01-09             NaT             1             1
11   2019-01-09             NaT             1             1
12   2019-01-09    2019-03-07             1             1
13   2020-07-31             NaT             3             3
14   2020-07-31             NaT             3             3
```

```
      Total Business Value  Quarterly Rating  Churn
0             2381060             2    0.0
1             -665480             2    0.0
2              0             2    1.0
```

3	0	1	0.0
4	0	1	0.0
5	0	1	0.0
6	0	1	0.0
7	0	1	0.0
8	350000	1	0.0
9	0	1	1.0
10	0	1	0.0
11	120360	1	0.0
12	0	1	1.0
13	0	1	0.0
14	0	1	0.0

```
[37]: #Lets do a quick check if there there are multiple incomes assigned to each
      ↪ driver
df.groupby(by='Driver_ID').nunique()['Income'].value_counts()

#We observe that most drivers have only 1 Income against them.
#However there are few drivers who have 2 different average incomes. Lets
      ↪ inspect.
```

```
[37]: 1    2255
      2     44
      Name: Income, dtype: int64
```

```
[38]: income_count=(df.groupby(by='Driver_ID').nunique()['Income'])
      income_count[income_count==2].index

# These drivers have 2 average incomes in the dataset. Lets look at the
      ↪ observations for driver_id 26.
```

```
[38]: Int64Index([ 26,   54,   60,   98,  275,  307,  320,  368,  434,  537,  568,
                  580,  582,  638,  716,  789,  888, 1031, 1050, 1161, 1165, 1206,
                  1249, 1274, 1316, 1327, 1770, 1783, 1817, 1840, 1852, 1877, 1918,
                  2008, 2070, 2087, 2198, 2272, 2390, 2407, 2543, 2567, 2625, 2690],
                  dtype='int64', name='Driver_ID')
```

```
[39]: df[df['Driver_ID']==26]

# We can confirm from the below observations that certain drivers have 2
      ↪ different incomes.
```

```
[39]:      MMM-YY  Driver_ID  Age  Gender  City  Education_Level  Income  \
137  2019-01-01         26  41.0    0.0   C14                2  121529
138  2019-02-01         26  41.0    0.0   C14                2  121529
139  2019-03-01         26  41.0    0.0   C14                2  121529
140  2019-04-01         26  41.0    0.0   C14                2  121529
```

141	2019-05-01	26	41.0	0.0	C14	2	121529
142	2019-06-01	26	41.0	0.0	C14	2	121529
143	2019-07-01	26	41.0	0.0	C14	2	121529
144	2019-08-01	26	41.0	0.0	C14	2	121529
145	2019-09-01	26	42.0	0.0	C14	2	121529
146	2019-10-01	26	42.0	0.0	C14	2	121529
147	2019-11-01	26	42.0	0.0	C14	2	121529
148	2019-12-01	26	42.0	0.0	C14	2	121529
149	2020-01-01	26	42.0	0.0	C14	2	121529
150	2020-02-01	26	42.0	0.0	C14	2	121529
151	2020-03-01	26	42.0	0.0	C14	2	132577
152	2020-04-01	26	42.0	0.0	C14	2	132577
153	2020-05-01	26	42.0	0.0	C14	2	132577
154	2020-06-01	26	42.0	0.0	C14	2	132577
155	2020-07-01	26	42.0	0.0	C14	2	132577
156	2020-08-01	26	42.0	0.0	C14	2	132577
157	2020-09-01	26	43.0	0.0	C14	2	132577
158	2020-10-01	26	43.0	0.0	C14	2	132577
159	2020-11-01	26	43.0	0.0	C14	2	132577
160	2020-12-01	26	43.0	0.0	C14	2	132577

	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
137	2018-05-07	NaT	1		3	
138	2018-05-07	NaT	1		3	
139	2018-05-07	NaT	1		3	
140	2018-05-07	NaT	1		3	
141	2018-05-07	NaT	1		3	
142	2018-05-07	NaT	1		3	
143	2018-05-07	NaT	1		3	
144	2018-05-07	NaT	1		3	
145	2018-05-07	NaT	1		3	
146	2018-05-07	NaT	1		3	
147	2018-05-07	NaT	1		3	
148	2018-05-07	NaT	1		3	
149	2018-05-07	NaT	1		3	
150	2018-05-07	NaT	1		3	
151	2018-05-07	NaT	1		4	
152	2018-05-07	NaT	1		4	
153	2018-05-07	NaT	1		4	
154	2018-05-07	NaT	1		4	
155	2018-05-07	NaT	1		4	
156	2018-05-07	NaT	1		4	
157	2018-05-07	NaT	1		4	
158	2018-05-07	NaT	1		4	
159	2018-05-07	NaT	1		4	
160	2018-05-07	NaT	1		4	

	Total Business Value	Quarterly Rating	Churn
137	243110	4	0.0
138	646330	4	0.0
139	17651940	4	0.0
140	242510	4	0.0
141	1098080	4	0.0
142	1212720	4	0.0
143	2695910	4	0.0
144	494710	4	0.0
145	986540	4	0.0
146	9950710	4	0.0
147	13097320	4	0.0
148	1086370	4	0.0
149	732410	2	0.0
150	1511840	2	0.0
151	9226690	4	0.0
152	1940050	2	0.0
153	970030	2	0.0
154	0	2	0.0
155	153590	2	0.0
156	1710410	2	0.0
157	440550	2	0.0
158	744590	2	0.0
159	1578270	2	0.0
160	1453220	2	0.0

```
[40]: #Lets find out the average income of each driver and store it in the aggregated
      ↪ dataset.
      income_avg=df.groupby(by='Driver_ID').mean()['Income'].reset_index()
      income_avg[:5]
```

```
[40]:   Driver_ID  Income
0         1  57387.0
1         2  67016.0
2         4  65603.0
3         5  46368.0
4         6  78728.0
```

```
[41]: dataset=pd.merge(left=dataset,right=income_avg,on='Driver_ID')
      dataset.head()
      #Sucessfully created a new feature which stores the average_income for each
      ↪ driver.
```

```
[41]:   Driver_ID  Churn  Duration  City  Education_Level  Income
0         1      1         77  C23                2  57387.0
1         2      0         25  C7                 2  67016.0
2         4      1        142  C13                2  65603.0
```

3	5	1	57	C9	0	46368.0
4	6	0	123	C11	1	78728.0

[]:

```
[42]: # Lets do another Feature Engineering step for Income
# If the income has increased for a particular driver, then we can flag him as 1, or else 0.
# To achieve this, we can compare the first month income and last month income for every driver.
```

```
# Getting the list of driver_ids
all_driver_ids=dataset['Driver_ID']

#Getting the indexes for every driver, and storing the key-value pair in a dict
dict1={}
for id in all_driver_ids:
    dict1[id]=list(df[df['Driver_ID']==id].index)
```

```
[43]: # Creating a Flag Column
dataset['Income_Flag']=-1

# Now lets compare the first and last month income for every driver and create the flag feature in our dataset.
```

```
for driver_id in all_driver_ids:
    first_index=min(dict1[driver_id])
    last_index=max(dict1[driver_id])
    first_income=df.iloc[first_index]['Income']
    last_income=df.iloc[last_index]['Income']
    if last_income>first_income:
        dataset.loc[dataset['Driver_ID']==driver_id,'Income_Flag']=1
    else:
        dataset.loc[dataset['Driver_ID']==driver_id,'Income_Flag']=0
```

```
[44]: # Lets have a look at the new feature that we created.
dataset.head()
```

```
[44]:
```

	Driver_ID	Churn	Duration	City	Education_Level	Income	Income_Flag
0	1	1	77	C23	2	57387.0	0
1	2	0	25	C7	2	67016.0	0
2	4	1	142	C13	2	65603.0	0
3	5	1	57	C9	0	46368.0	0
4	6	0	123	C11	1	78728.0	0

[]:

0.1.8 Inspecting and Aggregating “Age” for each Driver in the new Dataset

```
[45]: # From the inspection done earlier, we know "Age" feature has missing values.
df['Age'].isna().sum()
```

```
[45]: 61
```

```
[46]: # Lets inspect the dataset for missing values in age feature
df[df['Age'].isna()]
```

```
[46]:
```

	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	\
71	2020-02-01	20	NaN	1.0	C19	0	40342	
96	2019-10-01	22	NaN	0.0	C10	2	31224	
109	2019-07-01	24	NaN	0.0	C24	2	76308	
211	2019-11-01	40	NaN	0.0	C15	0	59182	
260	2019-05-01	49	NaN	0.0	C20	0	53039	
...	
18180	2020-05-01	2690	NaN	0.0	C11	2	77662	
18507	2020-08-01	2730	NaN	1.0	C16	2	69924	
18565	2019-03-01	2738	NaN	0.0	C17	0	23068	
18628	2019-01-01	2751	NaN	0.0	C17	2	53115	
18809	2019-02-01	2774	NaN	0.0	C15	1	42313	

	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
71	2019-10-25	NaT	3		3	
96	2018-05-25	NaT	1		1	
109	2018-05-25	NaT	1		2	
211	2019-11-08	NaT	2		2	
260	2018-05-25	NaT	1		2	
...	
18180	2018-07-17	NaT	1		2	
18507	2019-07-08	NaT	2		2	
18565	2018-09-08	NaT	1		1	
18628	2015-11-05	NaT	1		1	
18809	2018-07-21	NaT	1		1	

	Total Business Value	Quarterly Rating	Churn
71	0	1	0.0
96	200000	3	0.0
109	203240	3	0.0
211	0	1	0.0
260	124190	1	0.0
...
18180	692600	4	0.0
18507	161860	2	0.0
18565	639780	3	0.0
18628	506550	3	0.0
18809	1141280	4	0.0

[61 rows x 14 columns]

```
[47]: # Lets get the list of driver_ids for whom there are missing values in Age.
missing_age_driver_ids=list(set(df.loc[df['Age'].isna(),'Driver_ID']))
len(missing_age_driver_ids)
#There are total 57 such drivers.
```

[47]: 57

```
[48]: # As this dataset has more than 1 observation for every driver,lets check for
    ↳ valid values of "Age" for these drivers
# Since the records could be more than 12 months long, lets check for the
    ↳ maximum age of each driver, since age may change.
missing_age_drivers_found=df[df['Driver_ID'].isin(missing_age_driver_ids)].
    ↳ groupby('Driver_ID').max()['Age']
missing_age_drivers_found

#We found valid age for each of these drivers. Lets use these values in our
    ↳ dataset.
```

```
[48]: Driver_ID
20      26.0
22      41.0
24      31.0
40      32.0
49      22.0
63      28.0
69      32.0
103     26.0
120     27.0
167     26.0
179     26.0
183     26.0
204     31.0
215     26.0
305     24.0
313     31.0
325     26.0
369     32.0
422     27.0
458     31.0
541     27.0
560     42.0
607     27.0
617     22.0
718     27.0
```

778	31.0
901	32.0
954	31.0
1050	32.0
1072	28.0
1247	32.0
1378	28.0
1421	27.0
1430	31.0
1462	27.0
1588	42.0
1611	31.0
1669	22.0
1720	33.0
1852	55.0
1909	31.0
1932	28.0
1936	26.0
2073	21.0
2168	40.0
2273	32.0
2348	24.0
2351	31.0
2460	26.0
2507	26.0
2569	27.0
2618	28.0
2690	26.0
2730	32.0
2738	24.0
2751	32.0
2774	41.0

Name: Age, dtype: float64

```
[49]: mapping=dict(zip(list(missing_age_drivers_found.
    ↪index),list(missing_age_drivers_found)))
mapping
#Now lets fill the missing age values.

for d_id in missing_age_driver_ids:
    df.loc[df['Driver_ID']==d_id,'Age']=mapping[d_id]
```

```
[50]: #Checking if missing values have been removed
df['Age'].isna().sum()
```

```
[50]: 0
```

```
[51]: # Now lets store the max age of all drivers in the aggregate dataset
max_ages_of_drivers=df.groupby(by='Driver_ID').max()['Age']

dataset=pd.merge(left=dataset,right=max_ages_of_drivers,on='Driver_ID')
dataset.head()
#Sucessfully created a new feature which stores the max Age for each driver.
```

```
[51]:
```

	Driver_ID	Churn	Duration	City	Education_Level	Income	Income_Flag	\
0	1	1	77	C23	2	57387.0	0	
1	2	0	25	C7	2	67016.0	0	
2	4	1	142	C13	2	65603.0	0	
3	5	1	57	C9	0	46368.0	0	
4	6	0	123	C11	1	78728.0	0	

	Age
0	28.0
1	31.0
2	43.0
3	29.0
4	31.0

```
[ ]:
```

0.1.9 Inspecting and Aggregating “Gender” for each Driver in the new Dataset

```
[52]: # There were few missing values in "Gender" Feature found in our preliminary_
↳analysis
df['Gender'].isna().sum()
```

```
[52]: 52
```

```
[53]: # Lets check few the observation where there are missing values in the "Gender"
↳feature
df[df['Gender'].isna()][:20]
```

```
[53]:
```

	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	\
239	2019-02-01	43	27.0	NaN	C15	0	12906	
257	2019-02-01	49	22.0	NaN	C20	0	53039	
263	2019-08-01	49	22.0	NaN	C20	0	53039	
463	2019-08-01	68	31.0	NaN	C29	0	79288	
817	2019-02-01	116	21.0	NaN	C11	0	16477	
856	2019-11-01	119	31.0	NaN	C29	1	71000	
1509	2020-12-01	225	32.0	NaN	C14	0	44792	
1885	2019-08-01	296	31.0	NaN	C20	1	65094	
2267	2019-02-01	354	31.0	NaN	C11	0	60555	
2349	2020-03-01	365	24.0	NaN	C22	0	44740	
2656	2019-06-01	407	40.0	NaN	C13	1	58207	

2905	2020-03-01	439	27.0	NaN	C3	1	60246
2969	2020-07-01	446	31.0	NaN	C22	0	50832
3262	2020-01-01	489	31.0	NaN	C12	2	49475
3407	2020-10-01	516	26.0	NaN	C29	0	41099
3597	2020-02-01	541	27.0	NaN	C1	2	71812
4173	2020-02-01	611	32.0	NaN	C10	1	39216
4340	2019-12-01	640	26.0	NaN	C8	1	105931
4727	2019-04-01	709	31.0	NaN	C13	1	135436
5310	2020-12-01	793	26.0	NaN	C7	2	92670

	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
239	2018-07-13	2019-02-20			1	1
257	2018-05-25	NaT			1	2
263	2018-05-25	NaT			1	2
463	2015-10-18	NaT			1	3
817	2018-12-04	NaT			1	1
856	2019-11-16	NaT			3	3
1509	2020-07-13	NaT			3	3
1885	2018-06-10	NaT			1	2
2267	2018-11-30	NaT			1	1
2349	2020-02-01	NaT			2	2
2656	2016-05-17	NaT			1	2
2905	2019-11-28	NaT			1	1
2969	2020-02-01	NaT			3	3
3262	2019-10-18	NaT			1	1
3407	2019-07-04	NaT			1	1
3597	2017-02-10	NaT			1	2
4173	2020-01-13	NaT			2	2
4340	2019-12-07	NaT			3	3
4727	2018-08-04	NaT			5	5
5310	2020-10-30	NaT			3	3

	Total Business Value	Quarterly Rating	Churn
239	0	1	1.0
257	0	1	0.0
263	300300	2	0.0
463	544930	3	0.0
817	129590	1	0.0
856	0	1	0.0
1509	337020	3	0.0
1885	145670	2	0.0
2267	0	1	0.0
2349	0	1	0.0
2656	427570	4	0.0
2905	1561820	3	0.0
2969	890060	2	0.0
3262	500510	3	0.0

3407	422710	1	0.0
3597	200000	2	0.0
4173	0	1	0.0
4340	0	1	0.0
4727	0	1	0.0
5310	127800	1	0.0

```
[54]: # Lets get the driver_id for whom there are missing values in "Gender" column
gender_missing_driver_ids=list(set(df.loc[df['Gender'].isna(),'Driver_ID']))
len(gender_missing_driver_ids)
# There are total 51 driver_ids for whom there are missing values in Gender
```

```
[54]: 51
```

```
[55]: # Lets check if we have values present for "Gender" in other observations for
↳ the same driver_ids.
temp=df[df['Driver_ID'].isin(gender_missing_driver_ids)]
temp=temp.loc[~temp['Gender'].isna(),['Driver_ID','Gender']]
temp=temp[~temp.duplicated()]
temp.shape
# There are values present for all the 51 driver_ids
```

```
[55]: (51, 2)
```

```
[56]: mapping=dict(zip(temp['Driver_ID'],temp['Gender']))

# Lets replace the missing values in the dataset.
for d_id in gender_missing_driver_ids:
    df.loc[df['Driver_ID']==d_id,'Gender']=mapping[d_id]
```

```
[57]: #Checking if missing values have been removed
df['Gender'].isna().sum()
```

```
[57]: 0
```

```
[58]: # Now lets store the Gender of all drivers in the aggregate dataset
gender_drivers=df.loc[(~df[['Driver_ID','Gender']].
↳ duplicated()),['Driver_ID','Gender']]

dataset=pd.merge(left=dataset,right=gender_drivers,on='Driver_ID')
dataset.head()
#Sucessfully created a new feature which stores the Gender for each driver.
```

```
[58]:   Driver_ID  Churn  Duration  City  Education_Level  Income  Income_Flag  \
0         1      1         77   C23                2  57387.0            0
1         2      0         25   C7                 2  67016.0            0
2         4      1        142  C13                 2  65603.0            0
3         5      1         57   C9                 0  46368.0            0
```

4	6	0	123	C11	1	78728.0	0
---	---	---	-----	-----	---	---------	---

	Age	Gender
0	28.0	0.0
1	31.0	0.0
2	43.0	0.0
3	29.0	0.0
4	31.0	1.0

Note : It was asked in the question to use KNN_Imputer to impute missing values. However during my exploration process, I found out that there are original values present in the dataset that can be directly used instead of using KNN, which would be an approximation missing value imputation technique.

[]:

0.1.10 Inspecting and Aggregating “Joining_Designation” for each Driver in the new Dataset

```
[59]: # Lets check if there were multiple Joining_Designation for any of the drivers
df.groupby(by='Driver_ID').nunique()['Joining Designation'].value_counts()

# Observation - All the drivers have only 1 joining designation.
```

```
[59]: 1    2299
      Name: Joining Designation, dtype: int64
```

```
[60]: # We can store the Joining_Designation for each of the drivers in the
      ↪ aggregated dataset.
joining_designation=df.loc[~df[['Driver_ID','Joining Designation']].
      ↪ duplicated(),['Driver_ID','Joining Designation']]

dataset=pd.merge(left=dataset,right=joining_designation,on='Driver_ID')
dataset.head()
#Sucessfully created a new feature which stores the Joining_Designation for
      ↪ each driver.
```

```
[60]:
```

	Driver_ID	Churn	Duration	City	Education_Level	Income	Income_Flag	\
0	1	1	77	C23	2	57387.0	0	
1	2	0	25	C7	2	67016.0	0	
2	4	1	142	C13	2	65603.0	0	
3	5	1	57	C9	0	46368.0	0	
4	6	0	123	C11	1	78728.0	0	

	Age	Gender	Joining Designation
0	28.0	0.0	1
1	31.0	0.0	2

2	43.0	0.0	2
3	29.0	0.0	1
4	31.0	1.0	3

[]:

0.1.11 Inspecting and Aggregating “Grade” for each Driver in the new Dataset

```
[61]: # Lets check if there are multiple Grades for any of the drivers
df.groupby(by='Driver_ID').nunique()['Grade'].value_counts()

# Observation - Maximum drivers have 2 Grades,& few drivers have 1 Grade.Lets
↪check few records for drivers having 2 Grades.
```

```
[61]: 1    2255
      2     44
      Name: Grade, dtype: int64
```

```
[62]: temp=df.groupby(by='Driver_ID').nunique()['Grade']
multiple_grade_driver_ids=temp[temp==2].index
len(multiple_grade_driver_ids)

#There are 44 drivers who have 2 Grades. Lets have a look the observations for
↪one such driver.
```

```
[62]: 44
```

```
[63]: df[df['Driver_ID']==multiple_grade_driver_ids[0]]
# Perhaps with increase in time and performace, the Grade increases for drivers.
```

```
[63]:
```

	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	\
137	2019-01-01	26	41.0	0.0	C14	2	121529	
138	2019-02-01	26	41.0	0.0	C14	2	121529	
139	2019-03-01	26	41.0	0.0	C14	2	121529	
140	2019-04-01	26	41.0	0.0	C14	2	121529	
141	2019-05-01	26	41.0	0.0	C14	2	121529	
142	2019-06-01	26	41.0	0.0	C14	2	121529	
143	2019-07-01	26	41.0	0.0	C14	2	121529	
144	2019-08-01	26	41.0	0.0	C14	2	121529	
145	2019-09-01	26	42.0	0.0	C14	2	121529	
146	2019-10-01	26	42.0	0.0	C14	2	121529	
147	2019-11-01	26	42.0	0.0	C14	2	121529	
148	2019-12-01	26	42.0	0.0	C14	2	121529	
149	2020-01-01	26	42.0	0.0	C14	2	121529	
150	2020-02-01	26	42.0	0.0	C14	2	121529	
151	2020-03-01	26	42.0	0.0	C14	2	132577	
152	2020-04-01	26	42.0	0.0	C14	2	132577	

153	2020-05-01	26	42.0	0.0	C14	2	132577
154	2020-06-01	26	42.0	0.0	C14	2	132577
155	2020-07-01	26	42.0	0.0	C14	2	132577
156	2020-08-01	26	42.0	0.0	C14	2	132577
157	2020-09-01	26	43.0	0.0	C14	2	132577
158	2020-10-01	26	43.0	0.0	C14	2	132577
159	2020-11-01	26	43.0	0.0	C14	2	132577
160	2020-12-01	26	43.0	0.0	C14	2	132577

	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
137	2018-05-07	NaT	1		3	
138	2018-05-07	NaT	1		3	
139	2018-05-07	NaT	1		3	
140	2018-05-07	NaT	1		3	
141	2018-05-07	NaT	1		3	
142	2018-05-07	NaT	1		3	
143	2018-05-07	NaT	1		3	
144	2018-05-07	NaT	1		3	
145	2018-05-07	NaT	1		3	
146	2018-05-07	NaT	1		3	
147	2018-05-07	NaT	1		3	
148	2018-05-07	NaT	1		3	
149	2018-05-07	NaT	1		3	
150	2018-05-07	NaT	1		3	
151	2018-05-07	NaT	1		4	
152	2018-05-07	NaT	1		4	
153	2018-05-07	NaT	1		4	
154	2018-05-07	NaT	1		4	
155	2018-05-07	NaT	1		4	
156	2018-05-07	NaT	1		4	
157	2018-05-07	NaT	1		4	
158	2018-05-07	NaT	1		4	
159	2018-05-07	NaT	1		4	
160	2018-05-07	NaT	1		4	

	Total Business Value	Quarterly Rating	Churn
137	243110	4	0.0
138	646330	4	0.0
139	17651940	4	0.0
140	242510	4	0.0
141	1098080	4	0.0
142	1212720	4	0.0
143	2695910	4	0.0
144	494710	4	0.0
145	986540	4	0.0
146	9950710	4	0.0
147	13097320	4	0.0

148	1086370	4	0.0
149	732410	2	0.0
150	1511840	2	0.0
151	9226690	4	0.0
152	1940050	2	0.0
153	970030	2	0.0
154	0	2	0.0
155	153590	2	0.0
156	1710410	2	0.0
157	440550	2	0.0
158	744590	2	0.0
159	1578270	2	0.0
160	1453220	2	0.0

```
[64]: # Now lets store the Maximum Grade of all drivers in the aggregate dataset
grade_of_driver_ids=df.groupby(by='Driver_ID').max()['Grade']

dataset=pd.merge(left=dataset,right=grade_of_driver_ids,on='Driver_ID')
dataset.head()
#Sucessfully created a new feature which stores the Max Grade for each driver.
```

```
[64]:
```

	Driver_ID	Churn	Duration	City	Education_Level	Income	Income_Flag	\
0	1	1	77	C23	2	57387.0	0	
1	2	0	25	C7	2	67016.0	0	
2	4	1	142	C13	2	65603.0	0	
3	5	1	57	C9	0	46368.0	0	
4	6	0	123	C11	1	78728.0	0	

	Age	Gender	Joining	Designation	Grade
0	28.0	0.0		1	1
1	31.0	0.0		2	2
2	43.0	0.0		2	2
3	29.0	0.0		1	1
4	31.0	1.0		3	3

```
[ ]:
```

0.1.12 Inspecting and Aggregating “Total Business Value” for each Driver in the new Dataset

```
[65]: # For "Total Business Value", we can take the summation of the all the
↳ observations for this feature for each driver.
total_business_of_each_driver=df.groupby(by='Driver_ID').sum()['Total Business_
↳ Value']
```

```
dataset=pd.
↳merge(left=dataset,right=total_business_of_each_driver,on='Driver_ID')
dataset.head()
#Successfully created a new feature which stores the sum of total_business_value
↳for each driver in the aggregated dataset.
```

```
[65]:
```

	Driver_ID	Churn	Duration	City	Education_Level	Income	Income_Flag	\
0	1	1	77	C23	2	57387.0	0	
1	2	0	25	C7	2	67016.0	0	
2	4	1	142	C13	2	65603.0	0	
3	5	1	57	C9	0	46368.0	0	
4	6	0	123	C11	1	78728.0	0	

	Age	Gender	Joining	Designation	Grade	Total Business Value
0	28.0	0.0			1	1715580
1	31.0	0.0			2	0
2	43.0	0.0			2	350000
3	29.0	0.0			1	120360
4	31.0	1.0			3	1265000

```
[ ]:
```

0.1.13 Inspecting and Aggregating “Quarterly Rating” for each Driver in the new Dataset

```
[66]: # Lets check if there were multiple Quarterly Ratings for any of the drivers
df.groupby(by='Driver_ID').nunique()['Quarterly Rating'].value_counts()

#Observation-Drivers have got different ratings as per their quarter
↳performance.Most have got 1 kind of rating consitently.
```

```
[66]: 1    1203
      2    654
      3    352
      4     90
      Name: Quarterly Rating, dtype: int64
```

```
[67]: # Now lets store the Average Quarterly Ratings of all drivers in the aggregate
↳dataset
ratings_drivers=df.groupby(by='Driver_ID').mean()['Quarterly Rating']

dataset=pd.merge(left=dataset,right=ratings_drivers,on='Driver_ID')
dataset.head()
#Successfully created a new feature which stores the Average Quarterly Ratings
↳for each driver.
```

```
[67]:
```

	Driver_ID	Churn	Duration	City	Education_Level	Income	Income_Flag	\
0	1	1	77	C23	2	57387.0	0	
1	2	0	25	C7	2	67016.0	0	
2	4	1	142	C13	2	65603.0	0	
3	5	1	57	C9	0	46368.0	0	
4	6	0	123	C11	1	78728.0	0	

	Age	Gender	Joining	Designation	Grade	Total Business Value	\
0	28.0	0.0		1	1	1715580	
1	31.0	0.0		2	2	0	
2	43.0	0.0		2	2	350000	
3	29.0	0.0		1	1	120360	
4	31.0	1.0		3	3	1265000	

	Quarterly Rating
0	2.0
1	1.0
2	1.0
3	1.0
4	1.6

```
[ ]:
```

```
[68]: # Lets do another Feature Engineering step for Quarterly_Rating
# If the Quarterly_Rating has increased for a particular driver, then we can
    ↳ flag him as 1, or else 0.
# To achieve this, we can compare the first month Quarterly_Rating and last
    ↳ month Quarterly_Rating for every driver.

# Getting the list of driver_ids
all_driver_ids=dataset['Driver_ID']

#Getting the indexes for every driver, and storing the key-value pair in a dict
dict1={}
for id in all_driver_ids:
    dict1[id]=list(df[df['Driver_ID']==id].index)
```

```
[69]: # Creating a Flag Column
dataset['Quarterly_Rating_Flag']=-1

# Now lets compare the first and last month Quarterly_Rating for every driver
    ↳ and create the flag feature in our dataset.
for driver_id in all_driver_ids:
    first_index=min(dict1[driver_id])
    last_index=max(dict1[driver_id])
    first_income=df.iloc[first_index]['Quarterly Rating']
    last_income=df.iloc[last_index]['Quarterly Rating']
```

```

if last_income>first_income:
    dataset.loc[dataset['Driver_ID']==driver_id,'Quarterly_Rating_Flag']=1
else:
    dataset.loc[dataset['Driver_ID']==driver_id,'Quarterly_Rating_Flag']=0

```

```

[70]: # Lets have a look at the new feature that we created.
dataset.head()

```

```

[70]:   Driver_ID  Churn  Duration  City  Education_Level  Income  Income_Flag  \
0         1      1         77  C23                2  57387.0            0
1         2      0         25  C7                 2  67016.0            0
2         4      1        142  C13                2  65603.0            0
3         5      1         57  C9                 0  46368.0            0
4         6      0        123  C11                1  78728.0            0

```

```

      Age  Gender  Joining Designation  Grade  Total Business Value  \
0  28.0    0.0                1      1          1715580
1  31.0    0.0                2      2              0
2  43.0    0.0                2      2          350000
3  29.0    0.0                1      1          120360
4  31.0    1.0                3      3          1265000

```

```

      Quarterly Rating  Quarterly_Rating_Flag
0                2.0                0
1                1.0                0
2                1.0                0
3                1.0                0
4                1.6                1

```

```
[ ]:
```

0.2 Final Dataset

```

[71]: # We have our final aggregated dataset, which can be used for analysis and
      ↪model building.
dataset.head()

# Taking a look at out Final_Dataset.
dataset.head()

```

```

[71]:   Driver_ID  Churn  Duration  City  Education_Level  Income  Income_Flag  \
0         1      1         77  C23                2  57387.0            0
1         2      0         25  C7                 2  67016.0            0
2         4      1        142  C13                2  65603.0            0
3         5      1         57  C9                 0  46368.0            0
4         6      0        123  C11                1  78728.0            0

```

	Age	Gender	Joining	Designation	Grade	Total Business Value \
0	28.0	0.0		1	1	1715580
1	31.0	0.0		2	2	0
2	43.0	0.0		2	2	350000
3	29.0	0.0		1	1	120360
4	31.0	1.0		3	3	1265000

	Quarterly Rating	Quarterly_Rating_Flag
0	2.0	0
1	1.0	0
2	1.0	0
3	1.0	0
4	1.6	1

```
[72]: #Checking the shape of the dataset
dataset.shape
```

```
[72]: (2299, 14)
```

```
[73]: # Final Check for NULL values
dataset.isna().sum()
#There are no null values.
```

```
[73]: Driver_ID      0
      Churn         0
      Duration      0
      City          0
      Education_Level 0
      Income        0
      Income_Flag   0
      Age           0
      Gender        0
      Joining Designation 0
      Grade         0
      Total Business Value 0
      Quarterly Rating 0
      Quarterly_Rating_Flag 0
      dtype: int64
```

```
[74]: # Checking the Data Types of the columns
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2299 entries, 0 to 2298
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Driver_ID             2299 non-null  int64
```

```

1  Churn                2299 non-null  int64
2  Duration             2299 non-null  int64
3  City                 2299 non-null  object
4  Education_Level      2299 non-null  int64
5  Income               2299 non-null  float64
6  Income_Flag          2299 non-null  int64
7  Age                  2299 non-null  float64
8  Gender               2299 non-null  float64
9  Joining Designation  2299 non-null  int64
10 Grade                2299 non-null  int64
11 Total Business Value 2299 non-null  int64
12 Quarterly Rating     2299 non-null  float64
13 Quarterly_Rating_Flag 2299 non-null  int64
dtypes: float64(4), int64(9), object(1)
memory usage: 269.4+ KB

```

```

[75]: #Checking for duplicate values.
dataset[dataset.duplicated()]
# There are not duplicate values.

```

```

[75]: Empty DataFrame
Columns: [Driver_ID, Churn, Duration, City, Education_Level, Income,
Income_Flag, Age, Gender, Joining Designation, Grade, Total Business Value,
Quarterly Rating, Quarterly_Rating_Flag]
Index: []

```

```
[ ]:
```

0.2.1 Lets determine the categorical and continuous columns in our dataset.

```

[76]: data_types=[]
categories=[]
for column in dataset.columns:
    categories.append(dataset[column].nunique())
    data_types.append(dataset[column].dtype)

pd.
↳DataFrame(data=[categories,data_types],index=['Categories','Data_Type'],columns=dataset.
↳columns).T

```

```

[76]:
Categories Data_Type
Driver_ID      2299  int64
Churn           2    int64
Duration       879  int64
City           29   object
Education_Level  3    int64
Income        2260  float64
Income_Flag     2    int64

```

Age	36	float64
Gender	2	float64
Joining Designation	5	int64
Grade	5	int64
Total Business Value	1615	int64
Quarterly Rating	163	float64
Quarterly_Rating_Flag	2	int64

```
[77]: continuous_columns=['Duration','Income','Age','Total Business Value','Quarterly_
      ↳Rating']
      categorical_columns=['City','Education_Level','Income_Flag','Gender','Joining_
      ↳Designation','Grade','Quarterly_Rating_Flag']
```

```
[78]: # Converting the data type of categorical columns to category
      for column in categorical_columns:
          dataset[column]=dataset[column].astype('category')
```

```
[ ]:
```

0.2.2 Splitting The Data

```
[79]: # Splitting the dataset into train and test datasets
      dataset_train,dataset_test=train_test_split(dataset,test_size=0.
      ↳2,random_state=42,stratify=dataset['Churn'])
      print("Train",dataset_train.shape)
      print("Test",dataset_test.shape)
```

Train (1839, 14)

Test (460, 14)

Lets keep the Test_Data aside for now.

```
[ ]:
```

0.2.3 Descriptive Statistics

```
[80]: # Continuous Variables
      dataset_train[continuous_columns].describe()
```

```
[80]:
```

	Duration	Income	Age	Total Business Value \
count	1839.000000	1839.000000	1839.000000	1.839000e+03
mean	435.377379	59156.150565	33.764546	4.714062e+06
std	570.022273	28358.625209	6.007660	9.237594e+06
min	0.000000	10747.000000	21.000000	-1.385530e+06
25%	97.000000	39111.500000	29.000000	0.000000e+00
50%	185.000000	55108.000000	33.000000	8.643500e+05
75%	479.000000	75380.500000	37.000000	4.272055e+06
max	2801.000000	188418.000000	58.000000	9.533106e+07

```

Quarterly Rating
count      1839.000000
mean        1.580203
std         0.725855
min         1.000000
25%         1.000000
50%         1.200000
75%         2.000000
max         4.000000

```

```

[81]: plt.figure(figsize=(7,6))
      sns.heatmap(dataset_train[continuous_columns].corr(),annot=True,cmap='Blues')
      plt.show()

```



There is strong correlation between :

- “Duration” and “Total_Business_Value”
- “Quarterly_Rating” and “Total_Business_Value”
- “Duration” and “Quarterly_Rating”

```
[82]: # Categorical Variables
dataset_train[categorical_columns].describe()
```

```
[82]:      City  Education_Level  Income_Flag  Gender  Joining Designation \
count    1839             1839          1839  1839.0             1839
unique     29              3              2     2.0              5
top       C20              1              0     0.0              1
freq      115             622          1807  1074.0             798

      Grade  Quarterly_Rating_Flag
count    1839             1839
unique     5              2
top        2              0
freq      661             1562
```


```
[ ]:
```

0.3 Analysis

0.3.1 Churn

```
[83]: #Target Column - Churn
print(np.round(100*dataset_train['Churn'].value_counts(normalize=True)))
sns.countplot(data=dataset_train,x='Churn')
plt.show()
# We have imbalanced data. Need to balance the data before building the model.
```

```
1    70.0
0    30.0
Name: Churn, dtype: float64
```





[]:

0.3.2 Duration

```
[84]: # Lets look at the distribution of "Duration"
sns.histplot(data=dataset_train,x='Duration',bins=30)
plt.show()

sns.kdeplot(data=dataset_train,x='Duration')
plt.show()

# The distribution is highly right skewed.
# We can use Transformation to try to make it Normally Distributed.
```



```
[85]: #Lets check the range of "Duration" feature.  
dataset_train['Duration'].describe()
```

```

# The minimum value is 0, we cannot use Log_Transformation here, since log(0)
↳ is -inf.
# Box-Cox transformation also cannot be used, since Box-Cox transformation can
↳ only be used for positive values.
# Lets try using Yeo-Johnson Transformation

```

```

[85]: count    1839.000000
      mean      435.377379
      std       570.022273
      min        0.000000
      25%       97.000000
      50%      185.000000
      75%      479.000000
      max     2801.000000
      Name: Duration, dtype: float64

```

```

[86]: yeo_johnson_transformer=PowerTransformer(method='yeo-johnson')
      Duration_transformed=yeo_johnson_transformer.
      ↳fit_transform(dataset_train[['Duration']])

```



```

[87]: # Lets look at the distribution of "Duration"
      sns.histplot(Duration_transformed,bins=30)
      plt.show()

      sns.kdeplot(Duration_transformed.reshape(-1))
      plt.show()


      # The distribution is not exactly normal, but it is better than the
      ↳ original distribution.
      # We can use Yeo-Johnson transformation on "Duration" feature.

```



```
[88]: # Transforming the "Duration" column
dataset_train['Duration']=Duration_transformed.reshape(-1)
dataset_test['Duration']=yeo_johnson_transformer.
    ↳transform(dataset_test[['Duration']]).reshape(-1)
```

```
[89]: # Relationship of Duration with Churn
sns.boxplot(data=dataset_train,x='Churn',y='Duration')
plt.show()
# For people who Churn, the median Duration is lower.
```



```
[ ]:
```

0.3.3 City

```
[90]: # Lets look at the unique number of categories for "City" Feature
print(dataset_train['City'].unique(),"Categories") # There are 29 cities.
print(dataset_train['City'].unique())
```

29 Categories

['C4', 'C29', 'C6', 'C9', 'C7', ..., 'C5', 'C12', 'C13', 'C11', 'C18']

Length: 29

Categories (29, object): ['C1', 'C10', 'C11', 'C12', ..., 'C6', 'C7', 'C8', 'C9']

```
[91]: # Lets look at the frequencies of the difference categories.
round(dataset_train['City'].value_counts(normalize=True)*100,2)
```

```
[91]: C20    6.25
      C15    4.19
      C26    4.19
```

```

C27    4.02
C29    3.92
C16    3.70
C1     3.64
C14    3.64
C10    3.59
C28    3.48
C3     3.43
C8     3.37
C4     3.37
C22    3.26
C21    3.26
C2     3.26
C18    3.26
C5     3.21
C9     3.21
C25    3.15
C12    3.10
C6     3.05
C23    2.99
C24    2.99
C7     2.99
C11    2.94
C19    2.88
C17    2.88
C13    2.77
Name: City, dtype: float64

```

```
[92]: # We can One-Hot-Encode this categorical column later.
```

```
[ ]:
```

0.3.4 Education_Level

```

[93]: # Lets check the frequency counts of the different categories of
      ↪ "Education_Level" feature.
print(round(100*dataset_train['Education_Level'].
      ↪ value_counts(normalize=True),2))
sns.countplot(dataset_train['Education_Level'])
plt.show()


# The percentages of the different categories appearing are almost same.

```

```

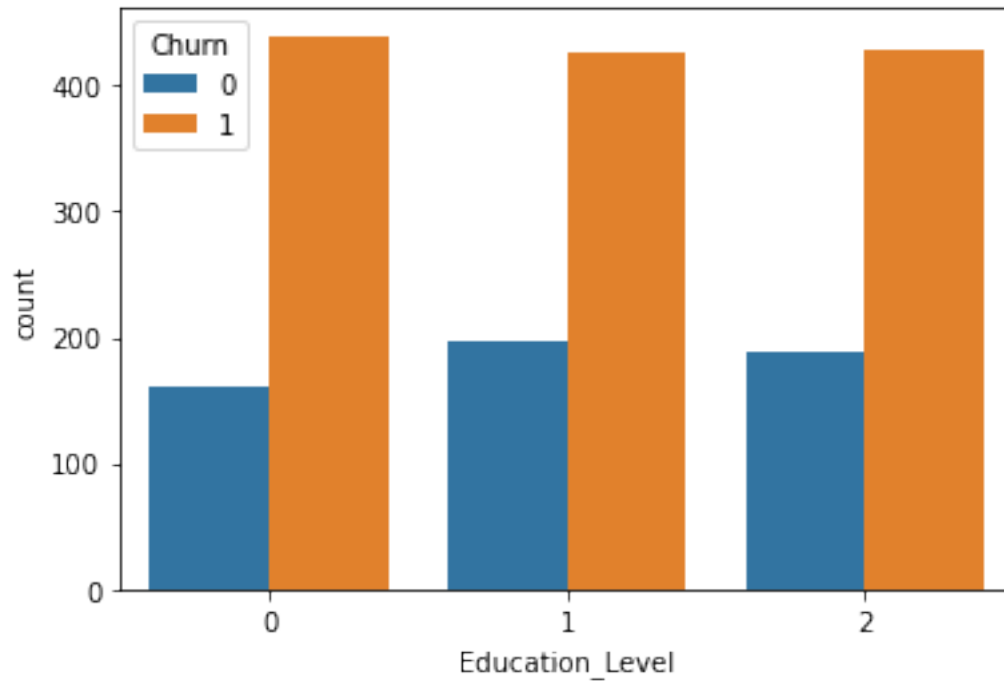
1     33.82
2     33.55
0     32.63
Name: Education_Level, dtype: float64

```



```
[94]: # Relationship of Churn with Education_Level
sns.countplot(data=dataset_train,hue='Churn',x='Education_Level')
plt.show()

# The ratio of churn to non-churn is little different for Education_Level-0
```



[]:

0.3.5 Income

```
[95]: # Lets look at the distribution of "Income"
sns.histplot(data=dataset_train,x='Income',bins=30)
plt.show()

sns.kdeplot(data=dataset_train,x='Income')
plt.show()

# The distribution is right skewed.
# We can use Transformation to try to make it Normally Distributed.
```



```
[96]: #Lets check the range of "Income" feature.
dataset_train['Income'].describe()


# We can use Log Transformation to try to make it Normally Distributed.
```


```
[96]: count      1839.000000
      mean      59156.150565
      std       28358.625209
      min       10747.000000
      25%       39111.500000
      50%       55108.000000
      75%       75380.500000
      max       188418.000000
      Name: Income, dtype: float64
```

```
[97]: sns.histplot(np.log(dataset_train['Income']),bins=30)
      plt.show()

      sns.kdeplot(np.log(dataset_train['Income']))
      plt.show()


# After doing Log Transformation, we get a distribution which is almost
↪ Normally Distributed.
# We can use Log transformation on "Income" feature.
```





```
[98]: # Transforming the "Income" column
dataset_train['Income']=np.log(dataset_train['Income'])
dataset_test['Income']=np.log(dataset_test['Income'])
```

```
[99]: # Relationship of Income with Churn
sns.boxplot(data=dataset_train,x='Churn',y='Income')
plt.show()
# For people who Churn, the median Income is lower.
```



[]:

0.3.6 Income_Flag

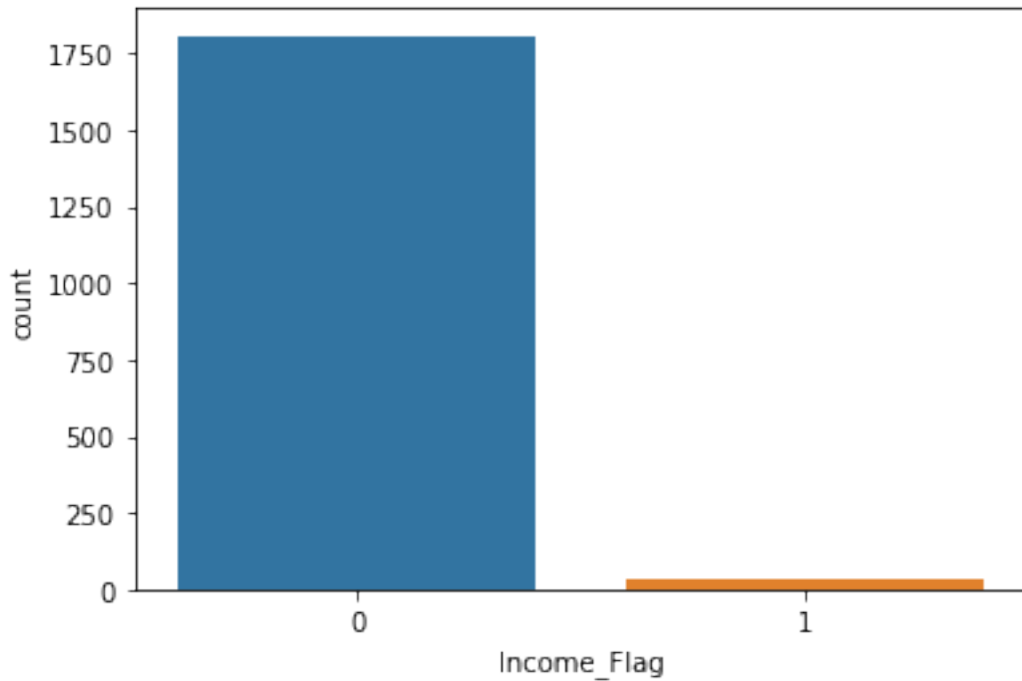
```
[100]: # Lets look at the frequencies of the difference categories of "Income_Flag".
print(round(dataset_train['Income_Flag'].value_counts(normalize=True)*100,2))
sns.countplot(dataset_train['Income_Flag'])
plt.show()

# Income Flag-1 signifies that the income of a particular driver has increased_
→over time.
# For very few people, the income has increased.
```

0 98.26

1 1.74


Name: Income_Flag, dtype: float64



```
[101]: # Relationship of Churn with Income_Flag
sns.countplot(data=dataset_train,hue='Churn',x='Income_Flag')
plt.show()

# The ratio of churn to non-churn is different for different Income_Flag.

pd.
↳ crosstab(dataset_train['Churn'],dataset_train['Income_Flag'],normalize='columns')
```



```
[101]: Income_Flag      0      1
      Churn
      0      0.28611  0.90625
      1      0.71389  0.09375
```



```
[ ]:
```

0.3.7 Age

```
[102]: # Lets look at the distribution of "Age"
sns.histplot(data=dataset_train,x='Age',bins=30)
plt.show()

sns.kdeplot(data=dataset_train,x='Age')
plt.show()

# The distribution is little right skewed.
# We can use Transformation to try to make it Normally Distributed.
```




```
[103]: #Lets check the range of "Age" feature.
dataset_train['Age'].describe()


# We can use Log Transformation to try to make it Normally Distributed.
```


```
[103]: count      1839.000000
mean         33.764546
std           6.007660
min          21.000000
25%          29.000000
50%          33.000000
75%          37.000000
max          58.000000
Name: Age, dtype: float64
```

```
[104]: sns.histplot(np.log(dataset_train['Age']),bins=30)
plt.show()

sns.kdeplot(np.log(dataset_train['Age']))
plt.show()


# After doing Log Transformation, we get a distribution which is almost
↳ Normally Distributed.
# We can use Log transformation on "Age" feature.
```





```
[105]: # Transforming the "Income" column  
dataset_train['Age']=np.log(dataset_train['Age'])  
dataset_test['Age']=np.log(dataset_test['Age'])
```

```
[106]: # Relationship of Income with Churn  
sns.boxplot(data=dataset_train,x='Churn',y='Age')  
plt.show()  
# For people who Churn, the median Age is lower.
```



[]:

0.3.8 Gender


```
[107]: # Lets look at the frequencies of the difference categories of "Gender".
print(round(dataset_train['Gender'].value_counts(normalize=True)*100,2))
sns.countplot(dataset_train['Gender'])
plt.show()
```

Gender-1 signifies that the driver is a female. Female:Male = 60:40

0.0 58.4

1.0 41.6

Name: Gender, dtype: float64




```
[108]: # Relationship of Gender with Churn

sns.countplot(data=dataset_train,hue='Churn',x='Gender')
plt.show()

# The ratio of churn to non-churn is little different for different Genders.

pd.crosstab(dataset_train['Churn'],dataset_train['Gender'],normalize='columns')
```



```
[108]: Gender      0.0      1.0
      Churn
      0      0.308194  0.281046
      1      0.691806  0.718954
```

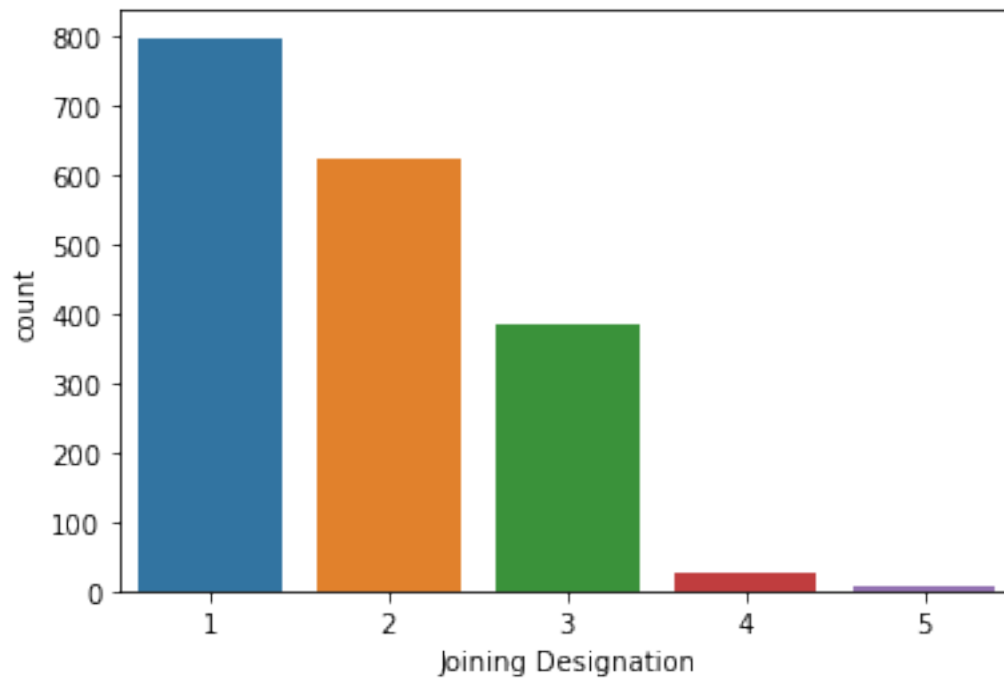
```
[ ]:
```

0.3.9 Joining Designation

```
[109]: # Lets check the frequency counts of the different categories of "Joining_
      ↪ Designation" feature.
      print(round(100*dataset_train['Joining Designation'].
      ↪ value_counts(normalize=True),2))
      sns.countplot(dataset_train['Joining Designation'])
      plt.show()

      # Most drivers had joined at designations - 1 and 2.
```

```
1      43.39
2      33.82
3      20.83
4       1.52
5       0.44
Name: Joining Designation, dtype: float64
```

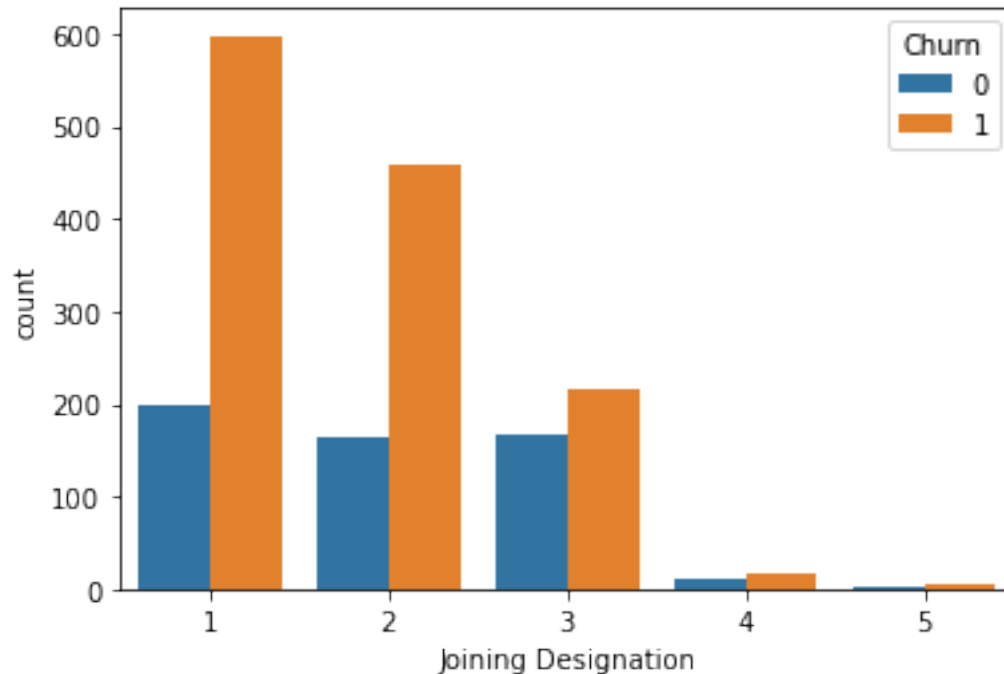


```
[110]: # Relationship of Joining_Designation with Churn

sns.countplot(data=dataset_train,hue='Churn',x='Joining Designation')
plt.show()

# The ratio of churn to non-churn is different for different Joining_
↳Designations.

pd.crosstab(dataset_train['Churn'],dataset_train['Joining_
↳Designation'],normalize='columns')
```



```
[110]: Joining Designation      1      2      3      4      5
Churn
0          0.250627  0.263666  0.438642  0.428571  0.25
1          0.749373  0.736334  0.561358  0.571429  0.75
```


```
[ ]:
```

0.3.10 Grade

```
[111]: # Lets check the frequency counts of the different categories of "Grade"
        ↳ feature.
print(round(100*dataset_train['Grade'].value_counts(normalize=True),2))
sns.countplot(dataset_train['Grade'])
plt.show()

# Most drivers have current Grade as - 1,2 and 3.
```

```
2    35.94
1    30.94
3    26.54
4     5.49
5     1.09
Name: Grade, dtype: float64
```




```
[112]: # Relationship of Grade with Churn

sns.countplot(data=dataset_train,hue='Churn',x='Grade')
plt.show()

# The ratio of churn to non-churn is different for different Grades.

pd.crosstab(dataset_train['Churn'],dataset_train['Grade'],normalize='columns')
```

```
[112]: Grade      1      2      3      4      5
      Churn
      0      0.165202  0.255673  0.456967  0.49505  0.5
      1      0.834798  0.744327  0.543033  0.50495  0.5
```

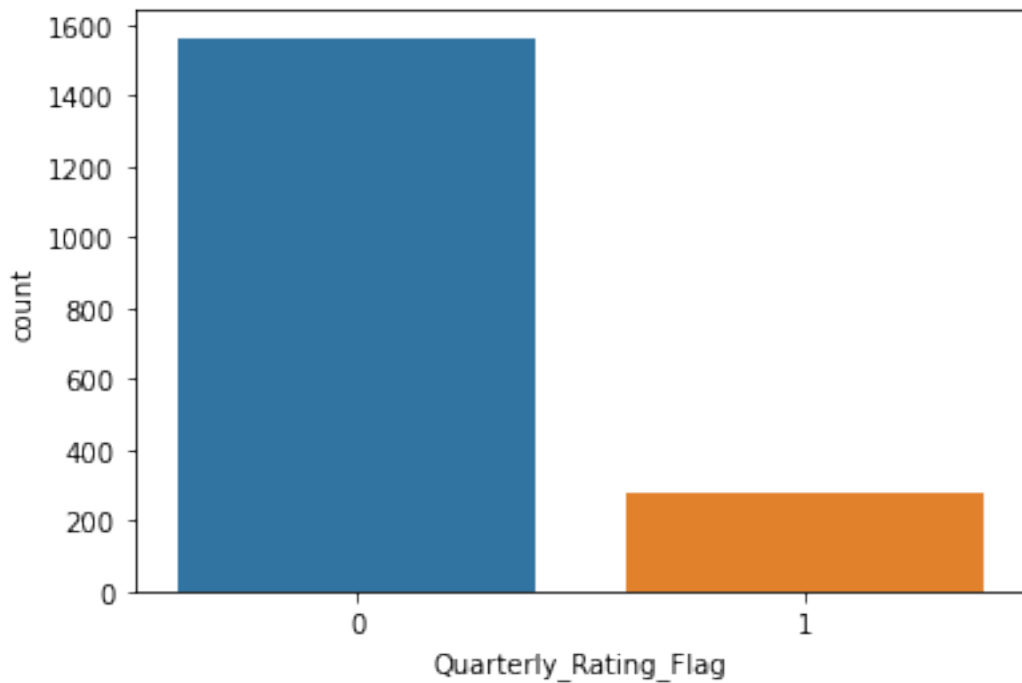
```
[ ]:
```

0.3.11 Quarterly_Rating_Flag

```
[113]: # Lets check the frequency counts of the different categories of "Quarterly_
      ↳Rating" feature.
      print(round(100*dataset_train['Quarterly_Rating_Flag'].
      ↳value_counts(normalize=True),2))
      sns.countplot(dataset_train['Quarterly_Rating_Flag'])
      plt.show()

      # Most drivers have current Grade as - 1,2 and 3.
```

```
0      84.94
1      15.06
Name: Quarterly_Rating_Flag, dtype: float64
```

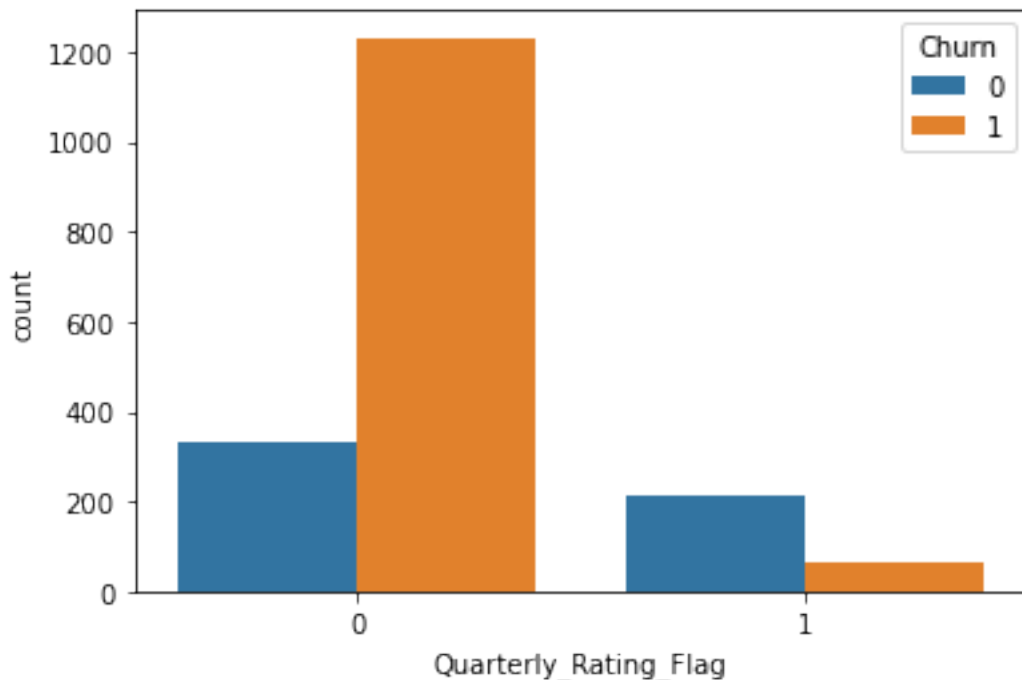


```
[114]: # Relationship of Quarterly_Rating_Flag with Churn

sns.countplot(data=dataset_train,hue='Churn',x='Quarterly_Rating_Flag')
plt.show()

# The ratio of churn to non-churn is different for different_
↳ Quarterly_Rating_Flags.

pd.
↳ crosstab(dataset_train['Churn'],dataset_train['Quarterly_Rating_Flag'],normalize='columns')
```



```
[114]: Quarterly_Rating_Flag      0      1
      Churn
      0      0.211908  0.776173
      1      0.788092  0.223827
```

```
[ ]:
```

0.4 Splitting into X and Y

```
[115]: # Firstly lets drop the "Driver_ID" column
dataset_train.drop(columns=['Driver_ID'],inplace=True)
dataset_test.drop(columns=['Driver_ID'],inplace=True)
```

```
[116]: # Splitting into X and Y
X_train=dataset_train.drop(columns=['Churn'])
Y_train=dataset_train['Churn']
X_test=dataset_test.drop(columns=['Churn'])
Y_test=dataset_test['Churn']
```

```
[ ]:
```

0.5 Encoding

```
[117]: # Lets decide on which features to do Encoding.  
X_train[categorical_columns].head()
```

```
# Only "City" feature need One Hot Encoding.
```

```
[117]:
```

	City	Education_Level	Income_Flag	Gender	Joining	Designation	Grade	\
731	C4	1	0	0.0			1	1
1793	C29	2	0	0.0			3	3
177	C6	2	0	1.0			3	3
28	C29	2	0	0.0			4	4
272	C9	0	0	1.0			2	2

	Quarterly_Rating_Flag
731	0
1793	1
177	0
28	1
272	0

```
[ ]:
```

0.6 Creating Transformer

```
[118]: # Lets use a Transformer, for OneHotEncoding categorical columns and for  
↳Scaling Numerical Columns.
```

```
transformer=ColumnTransformer(transformers=[  
    ('scaling',StandardScaler(),continuous_columns),  
    ↳  
    ↳('one_hot_encoding',OneHotEncoder(drop='first',sparse=False,handle_unknown='ignore'),['City'  
    ],remainder="passthrough")
```

```
[119]: # Transformed Data  
X_train_transformed=transformer.fit_transform(X_train)  
X_test_transformed=transformer.transform(X_test)
```

```
[ ]:
```

0.7 Handling Imabalanced Data using SMOTE

```
[120]: # Getting Balanced Data using SMOTE.  
smote_model=SMOTE(sampling_strategy='minority',random_state=42)  
X_train_transformed_SMOTE,Y_train_SMOTE=smote_model.  
↳fit_resample(X_train_transformed,Y_train)
```

```
[ ]:
```

1 Model 1 - Random Forests

```
[121]: # Creating a Random_Forest Model.
rf_model=RandomForestClassifier()

# Creating a parameters_list to do Hyperparameter Tuning
parameters = {
    'n_estimators': [15,20,25,30,35,40,45,50],
    'max_depth': [9,10,11,12,13,14]}

# Creating a Grid_Search Object
model=GridSearchCV(rf_model,parameters,scoring='accuracy',cv=5)
model.fit(X_train_transformed_SMOTE,Y_train_SMOTE)

print(model.best_params_)
print(model.best_score_)

{'max_depth': 14, 'n_estimators': 45}
0.887490198128496
```

```
[122]: # Using Grid_Search, the best accuracy that we are getting is 88.6 percent.
```

```
[ ]:
```

2 Model 2 - XGBoost

```
[123]: # Creating a XGBoost Model.
xgb_model=XGBClassifier()

# Creating a parameters_list to do Hyperparameter Tuning
parameters = {
    'n_estimators': [15,20,25,30,35,40,45,50],
    'max_depth': [9,10,11,12,13,14],
    'learning_rate': [0.5,0.75,1]}

# Creating a Grid_Search Object
model=GridSearchCV(xgb_model,parameters,scoring='accuracy',cv=5)
model.fit(X_train_transformed_SMOTE,Y_train_SMOTE)

print(model.best_params_)
print(model.best_score_)

{'learning_rate': 0.5, 'max_depth': 14, 'n_estimators': 50}
0.8932891720125763
```

```
[124]: # Using Grid_Search, the best accuracy that we are getting is 89.3 percent.
```

```
[125]: #Since the XGBoost model has a better validation score, we can therefore use  
→this model for determining the test_set accuracy
```

```
# Creating the final XGBoost model to find test dataset accuracy.  
model=XGBClassifier(learning_rate= 0.5, max_depth= 14, n_estimators= 50)
```

```
#Fitting the model  
model.fit(X_train_transformed_SMOTE,Y_train_SMOTE)
```

```
# Predicting Test Dataset Accuracy  
print("Accuracy:",model.score(X_test_transformed,Y_test))  
# The test dataset accuracy is 84.3 percent, which is a good score.
```

```
# Getting the Predictions  
Y_pred=model.predict(X_test_transformed)
```


```
# Getting the probabilities.  
Y_pred_prob=model.predict_proba(X_test_transformed)
```

Accuracy: 0.8434782608695652

```
[126]: # Getting the ROC_AUC Score.  
roc_auc_score(Y_test,Y_pred_prob[:,1])  
  
# A ROC_AUC score 0.88 is very good.
```


[126]: 0.8788049987570903

```
[127]: # ROC_AUC_Curve.  
fpr,tpr,thres=roc_curve(Y_test,Y_pred_prob[:,1])  
plt.plot(fpr,tpr)  
plt.show()  
  
# The curve is close to the ideal model curve
```



[]:

```
[128]: # Confusion Matrix
sns.heatmap(pd.crosstab(Y_pred,Y_test),annot=True,fmt='g',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel("Actual")
plt.show()
```



[]:

2.1 Insights & Recommendations

- Our XGBoost model has an accuracy of almost 85% and ROC_AUC score of almost 88%. Doing better preprocessing and feature transformation, engineering might help increase the model's performance even further.
- Duration is an important feature to determine Churn. Drivers who stay longer have a lower chance of leaving Ola. Ola can give more salary or other perks to make drivers stay longer. Income is a very important feature to determine Churn. Drivers who have a lower income are more probable to Churn.
- Almost 90 percent of drivers who had got a raise from Ola did not Churn. This can be an important point where Ola can give incentives and lucrative hikes to retain their drivers.
- Most drivers who Churn are relatively younger than those who don't. Ola can try to give certain goodies or hampers which can cater to the young drivers.
- Drivers who either join as Amateurs or as Highly Experienced have more chance of leaving. Ola can pay more to experienced drivers to earn their loyalty and can give lucrative incentives to young drivers to make them stay longer.
- Drivers with lower grades are more probable to Churn. Ola can increase their Grades after satisfactory performance to make drivers continue with Ola.
- Drivers whose Quarterly ratings had an increase had a 75 percent chance of not Churning. Therefore better hikes and ratings can also be an important factor to motivate drivers to continue with Ola.

[]: