

Project_10_Scaler

October 15, 2022

0.1 Problem Statement

Scaler is an online tech-versity offering intensive computer science & Data Science courses through live classes delivered by tech leaders and subject matter experts. The meticulously structured program enhances the skills of software professionals by offering a modern curriculum with exposure to the latest technologies. It is a product by InterviewBit.

You are working as a data scientist with the analytics vertical of Scaler, focused on profiling the best companies and job positions to work for from the Scaler database. You are provided with the information for a segment of learners and tasked to cluster them on the basis of their job profile, company, and other features. Ideally, these clusters should have similar characteristics.

0.2 Data Dictionary:

- 'Unnamed 0'- Index of the dataset
- Email_hash- Anonymised Personal Identifiable Information (PII)
- Company- Current employer of the learner
- orgyear- Employment start date
- CTC- Current CTC
- Job_position- Job profile in the company
- CTC_updated_year: Year in which CTC got updated (Yearly increments, Promotions)

```
[3]: #Importing the required libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder
import warnings
warnings.filterwarnings('ignore')
import re
from sklearn.cluster import AgglomerativeClustering, KMeans, DBSCAN
from sklearn.mixture import GaussianMixture
from yellowbrick.cluster import SilhouetteVisualizer
from sklearn.cluster import AgglomerativeClustering
from scipy.cluster.hierarchy import dendrogram, linkage
```

```
[4]: # Getting the Dataset.
df=pd.read_csv('https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/
↳002/856/original/scaler_clustering.csv')
```

```
[5]: # Checking the first 5 observations.
df.head()
```

```
[5]:      Unnamed: 0      company_hash \
0          0      atrgxmnt xzaxv
1          1  qtrxvzwt xzegwgb rxbxnta
2          2      ojzwnvwnxw vx
3          3      ngpgutaxv
4          4      qxen sqghu

      email_hash  orgyear      ctc \
0  6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...  2016.0  1100000
1  b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...  2018.0   449999
2  4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...  2015.0  2000000
3  effdede7a2e7c2af664c8a31d9346385016128d66bbc58...  2017.0   700000
4  6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...  2017.0  1400000

      job_position  ctc_updated_year
0          Other          2020.0
1  FullStack Engineer          2019.0
2   Backend Engineer          2020.0
3   Backend Engineer          2019.0
4  FullStack Engineer          2019.0
```

```
[6]: # Checking the last 5 observations.
df.tail()
```

```
[6]:      Unnamed: 0      company_hash \
205838      206918      vuurt xzw
205839      206919      husqvawgb
205840      206920      vwwgrxnt
205841      206921  zgn vuurxwvmrt
205842      206922  bgqsvz onvzrtj

      email_hash  orgyear      ctc \
205838  70027b728c8ee901fe979533ed94ffda97be08fc23f33b...  2008.0   220000
205839  7f7292ffad724ebbe9ca860f515245368d714c84705b42...  2017.0   500000
205840  cb25cc7304e9a24facda7f5567c7922ffc48e3d5d6018c...  2021.0   700000
205841  fb46a1a2752f5f652ce634f6178d0578ef6995ee59f6c8...  2019.0  5100000
205842  0bcfc1d05f2e8dc4147743a1313aa70a119b41b30d4a1f...  2014.0  1240000

      job_position  ctc_updated_year
205838          NaN          2019.0
```

205839	NaN	2020.0
205840	NaN	2021.0
205841	NaN	2019.0
205842	NaN	2016.0

```
[7]: # Checking the shape of the dataset.
df.shape

# Need to check if there are 205843 number of unique learners, otherwise we
→ need to remove/aggregate the duplicate records.
```

```
[7]: (205843, 7)
```

```
[8]: # Checking for number of duplicate values in "Unnamed: 0" column
print((df['Unnamed: 0'].duplicated()).sum())
# There are no duplicate values. All the values are unique.

# Dropping the "Unnamed: 0" column
df.drop(columns=['Unnamed: 0'],inplace=True)
```

```
0
```

```
[9]: # Checking for number of duplicate records.
print((df.duplicated()).sum())

# There are 33 duplicate records. We can drop these observations.
df=df[~df.duplicated()]
```

```
33
```

```
[10]: #Resetting the Index.
df.reset_index(drop=True,inplace=True)
```

```
[11]: # Checking the shape of the dataset.
df.shape
```

```
[11]: (205810, 6)
```

```
[12]: # Checking for total number of unique learners using the email-Id column.
df['email_hash'].nunique()

# Total Records:205810,Total Unique Learners:153443.Need to inspect the
→ remaining duplicate records, and then aggregate.
```

```
[12]: 153443
```

```
[13]: # Finding the list of email_ids for whom we have more than 1 record.
temp=df['email_hash'].value_counts()
duplicated_learners=temp[temp>1].index
```

```
# Finding the list of email_ids who have just 1 unique record.
non_duplicated_learners=temp[temp==1].index

print("Total Duplicated Learners",len(duplicated_learners))
print("Total Non-Duplicated Learners",len(non_duplicated_learners))
```

Total Duplicated Learners 41191
Total Non-Duplicated Learners 112252

```
[14]: # Checking for missing values for every column
df.isna().sum()

# There are no missing values for "email_hash", "ctc" and "ctc_updated_year".
# Need to handle missing values for "company_hash", "orgyear" and
↳ "job_position".
```

```
[14]: company_hash      44
email_hash             0
orgyear               86
ctc                   0
job_position          52547
ctc_updated_year       0
dtype: int64
```

```
[15]: # Checking the data types of the columns
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205810 entries, 0 to 205809
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   company_hash          205766 non-null  object
1   email_hash            205810 non-null  object
2   orgyear               205724 non-null  float64
3   ctc                   205810 non-null  int64
4   job_position          153263 non-null  object
5   ctc_updated_year      205810 non-null  float64
dtypes: float64(2), int64(1), object(3)
memory usage: 9.4+ MB
```

```
[15]:
```

Lets check the duplicated records for few learners to understand how to do the aggregation.

```
[16]: # Checking records of Learner 1
display(df[df['email_hash']==duplicated_learners[7]])
```

```
# Observation : The 8 records look almost the same except the ctc and
→ job_position.
```

	company_hash	email_hash \
37734	bvi ogenfvqt b4d5afa09bec8689017d8b29701b80d664ca37b83cb883...	
45982	bvi ogenfvqt b4d5afa09bec8689017d8b29701b80d664ca37b83cb883...	
144744	bvi ogenfvqt b4d5afa09bec8689017d8b29701b80d664ca37b83cb883...	
151696	bvi ogenfvqt b4d5afa09bec8689017d8b29701b80d664ca37b83cb883...	
153848	bvi ogenfvqt b4d5afa09bec8689017d8b29701b80d664ca37b83cb883...	
154626	bvi ogenfvqt b4d5afa09bec8689017d8b29701b80d664ca37b83cb883...	
197115	bvi ogenfvqt b4d5afa09bec8689017d8b29701b80d664ca37b83cb883...	
203140	bvi ogenfvqt b4d5afa09bec8689017d8b29701b80d664ca37b83cb883...	

	orgyear	ctc	job_position	ctc_updated_year
37734	2020.0	900000	Engineering Leadership	2021.0
45982	2020.0	900000	Engineering Intern	2021.0
144744	2020.0	900000	Data Analyst	2021.0
151696	2020.0	900000	Data Scientist	2021.0
153848	2020.0	900000	NaN	2021.0
154626	2020.0	900000	Software Engineer 1	2021.0
197115	2020.0	2000000	Engineering Intern	2021.0
203140	2020.0	2000000	Data Analyst	2021.0

```
[17]: # Checking records of Learner 2
display(df[df['email_hash']==duplicated_learners[15]])
```

```
# Observation : The 7 records look almost the same
→ organisation_year, ctc and job_position.
```

	company_hash \	email_hash	orgyear	ctc \
2396	ouxqg ogrhnxgz uqxcvnt rxbxnta			
8292	ouxqg ogrhnxgz uqxcvnt rxbxnta			
10927	ouxqg ogrhnxgz uqxcvnt rxbxnta			
35627	ouxqg ogrhnxgz uqxcvnt rxbxnta			
88098	ouxqg ogrhnxgz uqxcvnt rxbxnta			
88137	ouxqg ogrhnxgz uqxcvnt rxbxnta			
137341	ouxqg ogrhnxgz uqxcvnt rxbxnta			
2396	021ea9c97b6b287336e9345f39f9308c33ccbd15ac366d...	2020.0	300000	
8292	021ea9c97b6b287336e9345f39f9308c33ccbd15ac366d...	2020.0	300000	
10927	021ea9c97b6b287336e9345f39f9308c33ccbd15ac366d...	2020.0	300000	
35627	021ea9c97b6b287336e9345f39f9308c33ccbd15ac366d...	2020.0	300000	
88098	021ea9c97b6b287336e9345f39f9308c33ccbd15ac366d...	2022.0	420000	
88137	021ea9c97b6b287336e9345f39f9308c33ccbd15ac366d...	2022.0	420000	
137341	021ea9c97b6b287336e9345f39f9308c33ccbd15ac366d...	2022.0	420000	

	job_position	ctc_updated_year
2396	Engineering Intern	2021.0
8292	Other	2021.0
10927	Backend Engineer	2021.0
35627	Non Coder	2021.0
88098	Backend Engineer	2021.0
88137	Engineering Intern	2021.0
137341	Other	2021.0

```
[18]: # Checking records of Learner 3
display(df[df['email_hash']==duplicated_learners[36]])

# Observation : The 6 records look almost the same except the job_position and
↳ ctc_updated_year.
```

	company_hash	email_hash \
8336	atigat	8b711b2bb77300ab4bfb6f8bf54d9b28cdd214d8c66e9e...
8645	atigat	8b711b2bb77300ab4bfb6f8bf54d9b28cdd214d8c66e9e...
22481	atigat	8b711b2bb77300ab4bfb6f8bf54d9b28cdd214d8c66e9e...
194677	atigat	8b711b2bb77300ab4bfb6f8bf54d9b28cdd214d8c66e9e...
196770	atigat	8b711b2bb77300ab4bfb6f8bf54d9b28cdd214d8c66e9e...
198817	atigat	8b711b2bb77300ab4bfb6f8bf54d9b28cdd214d8c66e9e...

	orgyear	ctc	job_position	ctc_updated_year
8336	2018.0	3800000	FullStack Engineer	2019.0
8645	2018.0	3800000	Other	2019.0
22481	2018.0	3800000	Backend Engineer	2019.0
194677	2018.0	3800000	Backend Engineer	2020.0
196770	2018.0	3800000	FullStack Engineer	2020.0
198817	2018.0	3800000	Other	2020.0

```
[19]: # Checking records of Learner 4
display(df[df['email_hash']==duplicated_learners[41]])

# Observation : The 6 records look almost the same except the company_hash, ctc
↳ and job_position.
```

	company_hash \	email_hash	orgyear	ctc \
52127	yxuug cxatg			
58375	yxuug cxatg			
143785	nvvn wgzohrnvwj otqcxwto			
144053	nvvn wgzohrnvwj otqcxwto			
161826	nvvn wgzohrnvwj otqcxwto			
165732	nvvn wgzohrnvwj otqcxwto			
52127	8ffe41ee60bc738df2cb50dbf5c248300b4ddf830093c7...	2019.0	750000	
58375	8ffe41ee60bc738df2cb50dbf5c248300b4ddf830093c7...	2019.0	750000	
143785	8ffe41ee60bc738df2cb50dbf5c248300b4ddf830093c7...	2019.0	396000	

144053	8ffe41ee60bc738df2cb50dbf5c248300b4ddf830093c7...	2019.0	396000
161826	8ffe41ee60bc738df2cb50dbf5c248300b4ddf830093c7...	2019.0	396000
165732	8ffe41ee60bc738df2cb50dbf5c248300b4ddf830093c7...	2019.0	396000

	job_position	ctc_updated_year
52127	FullStack Engineer	2020.0
58375	NaN	2020.0
143785	Support Engineer	2020.0
144053	Backend Engineer	2020.0
161826	Other	2020.0
165732	NaN	2020.0

[19]:

- To aggregate the duplicate records, let's first sort the dataset according to, "email_hash", decreasing "ctc" and decreasing "ctc_updated_year".

- We can then extract the first record for every learner.

[20]: `df=df.
→sort_values(by=['email_hash','ctc','ctc_updated_year'],ascending=[True,False,False])`

[20]:

0.2.1 Creating Aggregated Dataset

[21]: `# Getting the first record for every learner.
df=df.groupby(by='email_hash').first().reset_index()

We can drop the "email_hash" column from the dataset since it's of no use to
→us anymore.
df.drop(columns='email_hash',inplace=True)

df.head()`

[21]:

	company_hash	orgyear	ctc	job_position \
0	bxwqgogen	2012.0	3500000	Backend Engineer
1	nqsn axsnvr	2013.0	250000	Backend Engineer
2	gunhb	2021.0	1300000	FullStack Engineer
3	bxwqgotbx wgqugqvnvgz	2004.0	2000000	FullStack Engineer
4	fvrbvqn rvmo	2009.0	3400000	None

	ctc_updated_year
0	2019.0
1	2020.0
2	2019.0
3	2021.0
4	2018.0

```
[22]: # Checking for missing values in Percentage.
round(100*df.isna().mean(),2)

# "company_hash" has almost 0.02% missing values.
# "orgyear" has almost 0.05% missing values.
# "job_position" has almost 13.18% missing values.
```

```
[22]: company_hash      0.02
orgyear      0.05
ctc          0.00
job_position 13.18
ctc_updated_year 0.00
dtype: float64
```

```
[23]: # Lets try to first deal with the missing values of "job_position".

# Lets look at the number of unique values of "job_position".
df['job_position'].nunique()

#There are a lot of categories for this column.Lets try doing a few data_
↳cleaning techniques to reduce the categories.
```

```
[23]: 813
```

```
[24]: # First using Regex to try to reduce duplicated categories with special_
↳characters.

def function(string):
    if string!=None:
        return re.sub('[^A-Za-z0-9 ]+', '', string)

df['job_position']=df['job_position'].apply(function)
```

```
[25]: # Lets convert all the "job_position" into lower-case and also removing leading_
↳and trailing spaces.
df['job_position']=(df['job_position'].str.lower()).str.strip()
```

```
[26]: # Lets also remove if there are spaces between the words.
def remove_spaces(job):
    if job !=None:
        job=job.split()
        job=' '.join(job)
        return job

df['job_position']=df['job_position'].apply(remove_spaces)
```



```
[27]: # Lets check how many categories we have now.  
df['job_position'].nunique()
```

```
# Great! We were able to remove almost 90 duplicate categories.
```

```
[27]: 721
```

```
[28]: # Since "job_position" is a categorical variable, lets look at its value_counts.  
df['job_position'].value_counts(normalize=True, dropna=False)*100
```

```
# Observation : We do not have a clear majority mode here.  
# Lets see try to look at the top 20 occurring job profiles.  
(df['job_position'].value_counts(normalize=True)*100)[:20]
```

```
[28]: backend engineer          27.981744  
fullstack engineer          16.064525  
other                       11.854916  
frontend engineer           6.957716  
engineering leadership      4.659245  
qa engineer                 4.616459  
data scientist              3.678154  
android engineer            3.672149  
sdet                       3.392159  
devops engineer             3.229269  
support engineer            2.372785  
ios engineer                1.926152  
data analyst                1.909638  
engineering intern          1.580105  
product designer            0.957821  
product manager             0.801687  
backend architect           0.794932  
research engineers          0.728124  
program manager             0.571990  
non coder                   0.399342  
Name: job_position, dtype: float64
```

```
[29]: # We can create 1 new category "Missing" to impute the missing values.  
# And we can create another new category "Combined" to combine all the  
→ categories whose occurrence is less than 1 percent.
```

```
df['job_position'].fillna('Missing', inplace=True)  
  
temp=df['job_position'].value_counts(normalize=True)*100  
jobs_to_replace=temp[temp<1].index  
df.loc[df['job_position'].isin(jobs_to_replace), 'job_position']='Combined'
```

```
[30]: # Lets check if the missing value handling was successful.
df['job_position'].isna().sum()
```

```
[30]: 0
```

```
[31]: # Lets check how many values are missing for "company_hash"
print('Total Missing Values',df['company_hash'].isna().sum())

# Lets check for the total number of categories in "company_hash"
print('Total Categories',df['company_hash'].nunique())

# Since there are so many companies and the number of missing values are very
↳ less, we can drop these rows.
df=df[~df['company_hash'].isna()]

# Lets try using Regex to try to reduce duplicated categories with special
↳ characters if present.

def function(string):
    if string!=None:
        return re.sub('[^A-Za-z0-9 ]+', '', string)
df['job_position']=df['job_position'].apply(function)

# Lets check for the total number of categories in "company_hash"
print('Total Categories after Regex',df['company_hash'].nunique())

# Resetting the Index
df.reset_index(drop=True,inplace=True)
```

Total Missing Values 32

Total Categories 36323

Total Categories after Regex 36323

```
[32]: # Lets check how many values are missing for "orgyear"
print('Total Missing Values:',df['orgyear'].isna().sum())
print()

# Lets check for the total number of categories in "orgyear"
print('Total Categories:',df['orgyear'].nunique())
print()

# Lets look at the possible values of "orgyear"
print(sorted(df['orgyear'].unique().tolist()))
print()

# We can see lot of illogical values such as 0,1,2,208,209,2204,20165. We can
↳ drop such rows.
```

```

# Lets check the percentage of records that lie between 2005 and 2022. Almost
↳ 97.5%, which is a good proportion.
print(df['orgyear'].between(2005,2022).mean()*100,'%')
print()
# We can drop the other observations.

df=df[df['orgyear'].between(2005,2022)]

# Lets check whether the missing values are still present.
print('Total Missing Values:',df['orgyear'].isna().sum())

# Hurrah! The missing values were dropped as part of the previous process.
# There was no need of using KNN imputation to fill the missing values.

# Resetting the Index
df.reset_index(drop=True,inplace=True)

```

Total Missing Values: 78

Total Categories: 77

```

[1.0, 1986.0, 1989.0, 1990.0, 1991.0, 1992.0, 1993.0, 1994.0, 1995.0, 1996.0,
1997.0, 1998.0, 1999.0, 2000.0, 2001.0, 2002.0, 2003.0, 2004.0, 2005.0, 2006.0,
2007.0, 2008.0, 2009.0, 2010.0, 2011.0, 2012.0, 2013.0, 2014.0, 2015.0, 2016.0,
2017.0, 2018.0, 2019.0, 2020.0, 2021.0, 2022.0, 2023.0, nan, 0.0, 2.0, 3.0, 4.0,
5.0, 6.0, 38.0, 83.0, 91.0, 200.0, 201.0, 206.0, 208.0, 209.0, 1900.0, 1970.0,
1971.0, 1972.0, 1973.0, 1976.0, 1977.0, 1979.0, 1981.0, 1982.0, 1984.0, 1985.0,
1987.0, 1988.0, 2024.0, 2025.0, 2026.0, 2027.0, 2028.0, 2029.0, 2031.0, 2101.0,
2106.0, 2107.0, 2204.0, 20165.0]

```

96.47352536649915 %

Total Missing Values: 0

```

[33]: # Final Check for presence of missing values.
df.isna().sum()

```

```

[33]: company_hash      0
      orgyyear         0
      ctc              0
      job_position     0
      ctc_updated_year  0
      dtype: int64

```


[33]:

0.3 New Features

0.3.1 1) “Years of Experience”

```
[34]: # We can get this new feature by subtracting "ctc_updated_year" from "orgyear".  
df['Years_of_Experience']=df['ctc_updated_year']-df['orgyear']
```

```
[35]: # Lets check if we have only non-negative values of "Years of Experience".  
plt.figure(figsize=(12,5))  
sns.countplot(data=df,x='Years_of_Experience')  
plt.xticks(rotation=90)  
plt.show()
```



```
[36]: # We observed that there are negative values for "Years of Experience", which  
      ↪ is not possible in the real world.  
      # We can therefore drop those records for which "Years of Experience" has  
      ↪ negative values.
```

```
df=df[df['Years_of_Experience']>=0]  
  
# Resetting the Index  
df.reset_index(drop=True,inplace=True)
```

```
[37]: # Now we can drop the columns "orgyear" and "ctc_updated_year"  
df.drop(columns=['orgyear','ctc_updated_year'],inplace=True)
```

```
[37]:
```

```
[38]: # Lets take a look at our dataset.  
df.head(3)
```

```
[38]:    company_hash      ctc      job_position  Years_of_Experience
      0      bxwqgogen  3500000  backend engineer              7.0
      1      nqsn axsxnv  250000  backend engineer              7.0
      2      fvrbvqn rvmo  3400000              Missing              9.0
```

```
[39]: df.shape
```

```
[39]: (141111, 4)
```

```
[40]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 141111 entries, 0 to 141110
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   company_hash          141111 non-null  object
1   ctc                   141111 non-null  int64
2   job_position          141111 non-null  object
3   Years_of_Experience   141111 non-null  float64
dtypes: float64(1), int64(1), object(2)
memory usage: 4.3+ MB
```

```
[41]: # Descriptive Statistics for Continuous Columns
display(df[['ctc', 'Years_of_Experience']].describe().T)

# Observation - The max value of "ctc" is very large. And the min value is
→strangely very small. Need to inspect further.
```

	count	mean	std	min	25%	\
ctc	141111.0	2.425988e+06	1.283431e+07	2.0	550000.0	
Years_of_Experience	141111.0	4.471267e+00	3.266160e+00	0.0	2.0	

	50%	75%	max
ctc	950000.0	1680000.0	1.000150e+09
Years_of_Experience	4.0	6.0	1.600000e+01

```
[42]: # Descriptive Statistics for Categorical Columns
display(df[['company_hash', 'job_position']].describe())
```

	company_hash	job_position
count	141111	141111
unique	33953	16
top	nvvn wgzhornvzwj otqcxwto	backend engineer
freq	4927	35709

```
[42]:
```

0.4 Analysis

0.4.1 1) "ctc"

```
[43]: df[['ctc']].describe().T
# As observed earlier, the max value is very large, and the min value is very
↪small. This could be a potential outlier.
```

```
[43]:          count          mean          std  min      25%      50%      75%  \
ctc  141111.0  2.425988e+06  1.283431e+07  2.0  550000.0  950000.0  1680000.0

          max
ctc  1.000150e+09
```

```
[44]: # Lets first look at smaller values of "ctc" columns.
# Lets check how many values are smaller than 1 lakh.

print((df['ctc']<100000).mean()*100)

# Almost 2.26 percent data is well below a logical minimum "ctc". Lets assume
↪these are errors and proceed to drop them.

df=df[df['ctc']>=100000]
```

2.262049025235453

```
[45]: # Now lets look at the larger values of "ctc" column.
# Lets check how many values are greater than 1 crore.

print((df['ctc']>10000000).mean()*100)

# Almost 1.28 percent data is greater than 1 crore. These are not very common
↪"ctc", so lets go ahead and drop them.


df=df[df['ctc']<=10000000]
```

1.2833619733321733

```
[46]: # Resetting the Index.
df.reset_index(drop=True,inplace=True)
```

```
[47]: # Plotting the boxplot.
plt.figure(figsize=(7,7))
sns.boxplot(data=df,x='ctc')
plt.show()


# There are very still big values in the "ctc" column. Lets use Log
↪Transformation on "ctc" column.
df['Log_ctc']=np.log(df['ctc'])
```



```
[48]: # Lets check the distribution of "ctc" column.
sns.kdeplot(df['Log_ctc'])

# The distribution looks much better. We can use this column instead to find
→ the clusters.
```

```
[48]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd61a0a4d90>
```



```
[49]: df.shape
```


```
[49]: (136149, 5)
```

0.4.2 2) “job_position”

```
[50]: # Lets look at the occurences of each of the "job_position" categories.

plt.figure(figsize=(15,4))
sns.countplot(data=df,x='job_position')
plt.xticks(rotation=90)
plt.show()

# "Backend-Engineer" and "Fullstack-Engineer" appear the most times.
```

[50]:


0.4.3 3) “ctc” with “job_position”

[51]: *# Lets look at the average salary of each of the job categories.*

```
plt.figure(figsize=(15,4))
sns.barplot(data=df,x='job_position',y='ctc')
plt.xticks(rotation=90)
plt.show()
```

"Engineering Leadership" earns the highest average "ctc" among all job categories.

"Support Engineer" has the minimum average "ctc" among all job categories.




[51]:

0.4.4 4) “Years_of_Experience” with “job_position”

```
[52]: # Lets look at the average salary of each of the job categories.
plt.figure(figsize=(15,4))
sns.barplot(data=df,x='job_position',y='Years_of_Experience')
plt.xticks(rotation=90)
plt.show()

# "Engineering Leadership" has the highest average "Years_of_Experience" among
↳ all job categories.
# "Engineering Intern" has the minimum average "Years_of_Experience" among all
↳ job categories.
```




```
[52]:
```

0.4.5 5) “company_hash”

```
[53]: # Lets look at the top 10 companies having most number of jobs.

plt.figure(figsize=(15,4))
temp=df['company_hash'].value_counts().sort_values(ascending=False)[:10]
sns.barplot(x=temp.index,y=temp.values)
plt.xticks(rotation=90)
plt.show()
```




[53]:

0.4.6 6) “company_hash” with “ctc”

[54]: *# Lets look at the top 10 companies having average "ctc".*

```
plt.figure(figsize=(15,4))
temp=df.groupby(by='company_hash').mean()['ctc'].sort_values(ascending=False)[:
    ↪10]
sns.barplot(x=temp.index,y=temp.values)
plt.xticks(rotation=90)
plt.show()
```

There are many companies who pay a very big "ctc". All the top 10 "ctc" are ↪ identical.



[54]:

0.5 Manual Clustering

0.5.1 1) “Designation Flag”

Creating “Designation_Flag” which shows learners with CTC greater/lower than the Average of their Company’s department having same Years of Experience.

- Before we create the flag, lets specify the rules based on which the flag would be created.
- If the “ctc” is less than 75% of “average_ctc”, we can flag it as 1.
- If the “ctc” is between 75% and 125% of “average_ctc”, we can flag it as 2.
- If the “ctc” is greater than 125% of “average_ctc”, we can flag it as 3.

```
[55]: # Lets find the "average_ctc" as per the "company_hash", "job_position" and
↳ "Years_of_Experience". We can name it "Average_1"
temp=df.groupby(by=["company_hash", "job_position", "Years_of_Experience"]).
↳agg({'ctc': 'mean',}).astype('int').reset_index().rename(columns={'ctc':
↳ "Average_1"})

# Merging the average_ctc with the dataset.
df=pd.
↳merge(left=df, right=temp, on=["company_hash", "job_position", "Years_of_Experience"])
```

```
[56]: # Lets have a look at our updated dataset.
df.head(3)
```

```
[56]:  company_hash      ctc      job_position  Years_of_Experience  Log_ctc  \
0    bxwqgogen  3500000  backend engineer           7.0  15.068274
1    bxwqgogen  3560000  backend engineer           7.0  15.085271
2    bxwqgogen  5500000  backend engineer           7.0  15.520259

      Average_1
0    3286052
1    3286052
2    3286052
```

```
[57]: # Lets create a function to create the "Designation_Flag" column.
def function(arr):
    if arr[0]>(1.25*arr[1]):
        return 3
    elif arr[0]<(0.75*arr[1]):
        return 1
    else:
        return 2

# Lets get the values of "Designation_Flag"
df['Designation_Flag']=df[['ctc', 'Average_1']].apply(function,axis=1)
```

```
# We can now drop the "Average_1" column.
df.drop(columns=['Average_1'],inplace=True)
```


```
[58]: # Lets have a look at our updated dataset.
df.head(3)
```

```
[58]:  company_hash      ctc      job_position  Years_of_Experience  Log_ctc  \
0    bxwqgogen  3500000  backend engineer           7.0  15.068274
1    bxwqgogen  3560000  backend engineer           7.0  15.085271
2    bxwqgogen  5500000  backend engineer           7.0  15.520259

      Designation_Flag
0                    2
1                    2
2                    3
```


```
[59]: # Lets look at the percentage of people belonging to each type of
      ↪ "Designation_Flag".
plt.figure(figsize=(15,4))
sns.countplot(data=df,x='Designation_Flag')
plt.xticks(rotation=90)
plt.show()

# There are more people who earn almost equal to the average salary as per
      ↪ their "company_hash","job_position","Years_of_Experience".
# There is a small section of learners who earn either very large or very large
      ↪ ctc.
```




```
[60]: # Lets look at the relationship of "Designation_Flag" with "ctc".
plt.figure(figsize=(15,4))
sns.barplot(data=df,x='Designation_Flag',y='ctc')
plt.xticks(rotation=90)
plt.show()
```

```
# We can see a good difference between the relative ctcs between flag_1, flag_2
→ and flag_3.
```



```
[61]: # Lets look at the relationship of "Designation_Flag" with
→ "Years_of_Experience".
plt.figure(figsize=(15,4))
sns.barplot(data=df,x='Designation_Flag',y='Years_of_Experience')
plt.xticks(rotation=90)
plt.show()
```

```
# This is an interesting discovery. Learners who earn the highest salaries do
→ not have the highest "Years_of_Experience".
```



```
[61]:
```

0.5.2 2) “Class Flag”

Creating “Class_Flag” which shows learners with CTC greater/lower than the Average of their Company’s department.

- Before we create the flag, let's specify the rules based on which the flag would be created.
- If the “ctc” is less than 75% of “average_ctc”, we can flag it as 1.
- If the “ctc” is between 75% and 125% of “average_ctc”, we can flag it as 2.

- If the “ctc” is greater than 125% of “average_ctc”, we can flag it as 3.

```
[62]: # Lets find the "average_ctc" as per the "company_hash" and "job_position". We
      ↪ can name it "Average_2"
temp=df.groupby(by=["company_hash","job_position"]).agg({'ctc':'mean',}).
      ↪ astype('int').reset_index().rename(columns={'ctc':"Average_2"})

# Merging the average_ctc with the dataset.
df=pd.merge(left=df,right=temp,on=["company_hash","job_position"])
```

```
[63]: # Lets have a look at our updated dataset.
df.head(3)
```

```
[63]:  company_hash      ctc      job_position  Years_of_Experience  Log_ctc  \
0    bxwqgogen  3500000  backend engineer             7.0  15.068274
1    bxwqgogen  3560000  backend engineer             7.0  15.085271
2    bxwqgogen  5500000  backend engineer             7.0  15.520259

      Designation_Flag  Average_2
0                    2    2873402
1                    2    2873402
2                    3    2873402
```

```
[64]: # Lets get the values of "Class_Flag"
df['Class_Flag']=df[['ctc','Average_2']].apply(function,axis=1)

# We can now drop the "Average_1" column.
df.drop(columns=['Average_2'],inplace=True)
```


```
[65]: # Lets have a look at our updated dataset.
df.head(3)
```

```
[65]:  company_hash      ctc      job_position  Years_of_Experience  Log_ctc  \
0    bxwqgogen  3500000  backend engineer             7.0  15.068274
1    bxwqgogen  3560000  backend engineer             7.0  15.085271
2    bxwqgogen  5500000  backend engineer             7.0  15.520259

      Designation_Flag  Class_Flag
0                    2           2
1                    2           2
2                    3           3
```

```
[66]: # Lets look at the percentage of people belonging to each type of "Class_Flag".
plt.figure(figsize=(15,4))
sns.countplot(data=df,x='Class_Flag')
plt.xticks(rotation=90)
plt.show()
```


```
# There are more people who earn almost equal to the average salary as per
→ their "company_hash" and "job_position"
# There is a small section of learners who earn either very large or very large
→ ctc.
```



[67]: # Lets look at the relationship of "Class_Flag" with "ctc".

```
plt.figure(figsize=(15,4))
sns.barplot(data=df,x='Class_Flag',y='ctc')
plt.xticks(rotation=90)
plt.show()
```


```
# We can see a good difference between the relative ctcs between flag_1, flag_2
→ and flag_3.
```



[68]: # Lets look at the relationship of "Class_Flag" with "Years_of_Experience".

```
plt.figure(figsize=(15,4))
sns.barplot(data=df,x='Class_Flag',y='Years_of_Experience')
plt.xticks(rotation=90)
plt.show()
```

```
# Learners who earn the highest salaries have the highest "Years_of_Experience".
```

[68]:

0.5.3 3) “Tier Flag”

Creating “Tier_Flag” which shows learners with CTC greater/lower than the Average of their Company.

- Before we create the flag, lets specify the rules based on which the flag would be created.
- If the “ctc” is less than 75% of “average_ctc”, we can flag it as 1.
- If the “ctc” is between 75% and 125% of “average_ctc”, we can flag it as 2.
- If the “ctc” is grater than 125% of “average_ctc”, we can flag it as 3.

```
[69]: # Lets find the "average_ctc" as per the "company_hash". We can name it
      ↪ "Average_3"
temp=df.groupby(by=["company_hash"]).agg({'ctc':'mean',}).astype('int').
      ↪reset_index().rename(columns={'ctc':"Average_3"})

# Merging the average_ctc with the dataset.
df=pd.merge(left=df,right=temp,on=["company_hash"])
```

```
[70]: # Lets have a look at our updated dataset.
df.head(3)
```

```
[70]:  company_hash      ctc      job_position  Years_of_Experience  Log_ctc  \
0    bxwqgogen  3500000  backend engineer           7.0  15.068274
1    bxwqgogen  3560000  backend engineer           7.0  15.085271
2    bxwqgogen  5500000  backend engineer           7.0  15.520259

      Designation_Flag  Class_Flag  Average_3
0                    2           2    2792920
1                    2           2    2792920
2                    3           3    2792920
```

```
[71]: # Lets get the values of "Tier_Flag"
df['Tier_Flag']=df[['ctc', 'Average_3']].apply(function,axis=1)
```

```
# We can now drop the "Average_1" column.
df.drop(columns=['Average_3'],inplace=True)
```


```
[72]: # Lets have a look at our updated dataset.
df.head(3)
```

```
[72]:  company_hash      ctc      job_position  Years_of_Experience  Log_ctc  \
0    bxwqgogen  3500000  backend engineer             7.0  15.068274
1    bxwqgogen  3560000  backend engineer             7.0  15.085271
2    bxwqgogen  5500000  backend engineer             7.0  15.520259

   Designation_Flag  Class_Flag  Tier_Flag
0                 2           2           3
1                 2           2           3
2                 3           3           3
```


```
[73]: # Lets look at the percentage of people belonging to each type of "Tier_Flag".
plt.figure(figsize=(15,4))
sns.countplot(data=df,x='Tier_Flag')
plt.xticks(rotation=90)
plt.show()

# There are more people who earn almost equal to the average salary as per
→ their "company_hash".
# There is a small section of learners who earn either very large or very large
→ ctc.
```




```
[74]: # Lets look at the relationship of "Tier_Flag" with "ctc".
plt.figure(figsize=(15,4))
sns.barplot(data=df,x='Tier_Flag',y='ctc')
plt.xticks(rotation=90)
plt.show()

# We can see a good difference between the relative ctcs between flag_1, flag_2
→ and flag_3.
```



```
[75]: # Lets look at the relationship of "Tier_Flag" with "Years_of_Experience".
plt.figure(figsize=(15,4))
sns.barplot(data=df,x='Tier_Flag',y='Years_of_Experience')
plt.xticks(rotation=90)
plt.show()

# Learners who earn the highest salaries have the highest "Years_of_Experience".
```



```
[75]:
```

0.6 Encoding

```
[76]: # Lets check which columns we need to encode.
display(df.head(2))

# Since "company_hash" has a lot of categories, we can drop it for doing
↳ clustering.
# For doing clustering, we can also drop "Designation_Flag", "Class_Flag" and
↳ "Tier_Flag".
# Also, we can drop "ctc" column since we have a new column - "Log_ctc".
# We need to One-Hot-Encode "job_position"
```

```
df=df[['job_position', 'Years_of_Experience', 'Log_ctc']]
encoder=OneHotEncoder(sparse=False)
temp=pd.DataFrame(data=encoder.
    ↪fit_transform(df[['job_position']]),columns=encoder.get_feature_names())
df=pd.concat((df.iloc[:,1:],temp),axis=1)
```

	company_hash	ctc	job_position	Years_of_Experience	Log_ctc	\
0	bxwqgogen	3500000	backend engineer	7.0	15.068274	
1	bxwqgogen	3560000	backend engineer	7.0	15.085271	

	Designation_Flag	Class_Flag	Tier_Flag
0	2	2	3
1	2	2	3

[76]:

0.7 Scaling

[77]: *# Lets check which columns we need to encode.*

```
display(df.head(2))

# We need to scale "Years_of_Experience" and "Log_ctc".
scaler=StandardScaler()
temp=pd.DataFrame(data=scaler.fit_transform(df.iloc[:,[0,1]]),columns=scaler.
    ↪get_feature_names_out())
df=pd.concat((temp,df.iloc[:,2:]),axis=1)
```

	Years_of_Experience	Log_ctc	x0_Combined	x0_Missing	\
0	7.0	15.068274	0.0	0.0	
1	7.0	15.085271	0.0	0.0	

	x0_android engineer	x0_backend engineer	x0_data analyst	\
0	0.0	1.0	0.0	
1	0.0	1.0	0.0	

	x0_data scientist	x0_devops engineer	x0_engineering intern	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	

	x0_engineering leadership	x0_frontend engineer	x0_fullstack engineer	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	

	x0_ios engineer	x0_other	x0_qa engineer	x0_sdet	x0_support engineer
0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0

```
[78]: # Final Look at our encoded and scaled dataset.
df.head(2)
```

```
[78]:   Years_of_Experience   Log_ctc  x0_Combined  x0_Missing \
0           0.774817   1.631987           0.0           0.0
1           0.774817   1.653532           0.0           0.0

   x0_android engineer  x0_backend engineer  x0_data analyst \
0                0.0                1.0                0.0
1                0.0                1.0                0.0

   x0_data scientist  x0_devops engineer  x0_engineering intern \
0                0.0                0.0                0.0
1                0.0                0.0                0.0

   x0_engineering leadership  x0_frontend engineer  x0_fullstack engineer \
0                0.0                0.0                0.0
1                0.0                0.0                0.0

   x0_ios engineer  x0_other  x0_qa engineer  x0_sdet  x0_support engineer
0                0.0      0.0            0.0      0.0                0.0
1                0.0      0.0            0.0      0.0                0.0
```

```
[ ]:
```

1 Unsupervised Learning - Clustering

1.1 1) K-Means Clustering


```
[ ]: # Lets use KMeans clustering to find the optimal number of clusters. Lets try
      ↪ for clusters 2,3,4,5 and 6.
inertias=[]
silhouette_scores=[]
kmeans_per_k = [KMeans(n_clusters=k, random_state=42).fit(df) for k in range(2,
      ↪ 7)]


# Inertia and Silhouette_Scores for deciding the optimal number of clusters
inertias = [model.inertia_ for model in kmeans_per_k]
silhouette_scores = [silhouette_score(df, model.labels_) for model in
      ↪ kmeans_per_k]
```

```
[ ]: # Elbow method to determine the optimal number of clusters
plt.figure(figsize=(10, 6))
plt.plot(range(2,7), inertias, "ro-")
plt.xlabel("$k$", fontsize=14)
plt.ylabel("Inertia", fontsize=14)
```

```
plt.annotate('Elbow', xy=(3, inertias[1]), xytext=(0.55, 0.55),
    ↳textcoords='figure fraction',
    fontsize=16, arrowprops=dict(facecolor='black', shrink=0.1))
plt.show()



# From the elbow method, we can see a sharp bend for k=3. Lets look at the
↳segmenation for 3 clusters.
plt.scatter(x=df['Years_of_Experience'], y=df['Log_ctc'], c=kmeans_per_k[1].
    ↳labels_, cmap='Set2')
plt.xlabel("Years of Experience", fontsize=14)
plt.ylabel("Log-CTC", fontsize=14)
plt.show()
```





```
[ ]: # Silhouette_scores to determine the optimal number of clusters.
plt.figure(figsize=(10, 6))
plt.plot(range(2,7), silhouette_scores, "bo-")
plt.xlabel("$k$", fontsize=14)
plt.ylabel("Inertia", fontsize=14)
plt.show()

# From this plot, k=2 or k=3 is the right number of clusters. Lets look at the
↳ segmenation for 2 clusters.
plt.scatter(x=df['Years_of_Experience'],y=df['Log_ctc'],c=kmeans_per_k[0].
↳ labels_,cmap='Set2')
plt.xlabel("Years of Experience", fontsize=14)
plt.ylabel("Log-CTC", fontsize=14)
plt.show()
```




```
[ ]: # Observation - 3 Clusters are more suitable according to the business sense.
# - The first cluster belongs to those young learners who are probably highly
    ↳ skilled and have joined high paying startups and MNCs.
# - The second cluster belongs to those learner who are probably working in
    ↳ service based MNCs and their salary increases with experience.
# - The third cluster belongs to those learner who are probably working in top
    ↳ product based companies and are very skillful and experienced.
```


```
[ ]:
```

1.2 2) Hierarchical Clustering

```
[81]: # Lets see how many clusters "Hierarchical Clustering" on its own.
# Since it will take a lot of time to run the algorithm on the entire dataset,
    ↳ therefore lets take only 5,000 observations.

linkage_data = linkage(df.sample(4000), method='ward', metric='euclidean')
dendrogram(linkage_data)
plt.show()

# From the above plot we can see that the number of clusters is 3, which is
    ↳ also verified by the Elbow method from K_Means.
```




```
[86]: # Lets look at the segmentation of the 3 clusters by using
      ↪ "Hierarchical_Clustering".
      # Since it will take a lot of time to run the algorithm on the entire dataset,
      ↪ therefore lets take only 15,000 observations.
      data=df.sample(15000)

      hierarchical_cluster = AgglomerativeClustering(n_clusters=3,
      ↪ affinity='euclidean', linkage='ward')
      labels = hierarchical_cluster.fit_predict(data)

      plt.scatter(data['Years_of_Experience'], data['Log_ctc'], c=labels,cmap='Set1')
      plt.show()

      # The segmentation of clusters by using "Hierarchical_Clustering" is almost the
      ↪ same as by using "KMeans".
```



2 Actionable Insights & Recommendations

- From the clustering analysis and also by business intuition, we finally decided that there are 3 meaningful clusters.
- The first cluster belongs to those young learners who are probably highly skilled and have joined high paying startups and MNCs.

- The second cluster belongs to those learner who are probably working in service based MNCs and their salary increases with experience.
- The third cluster belongs to those learner who are probably working in top product based companies and are very skillful and experienced.
- For very high paying start-ups for less experienced learners, Scaler can identify, train and make them ready to crack those jobs.
- For stable MNCs, learners who are not highly skilled or who are looking to do a career transition can be placed in such companies.
- For highly experienced jobs, some learners can be identified who are both skillful and highly experieced to get those jobs.
- There are a lot of “Backend-Engineer” and “Fullstack-Engineer” jobs available, so getting learned placed for this job role would be relatively easy.
- “Engineering Leadership” earns the highest average “ctc” among all job categories, and these learners have the highest average experience. A highly experienced learner might be better suited for “Engineering Leadership” job.
- “Support Engineer” has the minimum average “ctc” among all job categories. This can be a good marketing strategy to attract learners who are looking to do a career transition.
- As expected, “Engineering Intern” has the minimum average “Years_of_Experience” among all job categories.
- “Data Scientist” and “Backend Engineer” are one of the most highly paid jobs, and Scaler’s program these program for both these job profiles. So Scaler can get learner placed in such jobs easily.
- There are a few masked companies who hire a lot of people, so Scaler can also placed suitable learners in these companies based on the job fit requirement.
- There are a few companies who offer the best salaries to its employees. Scaler can shortlist such learners from its cohort and try to get learners placed in them to increase its credibility.

[]: